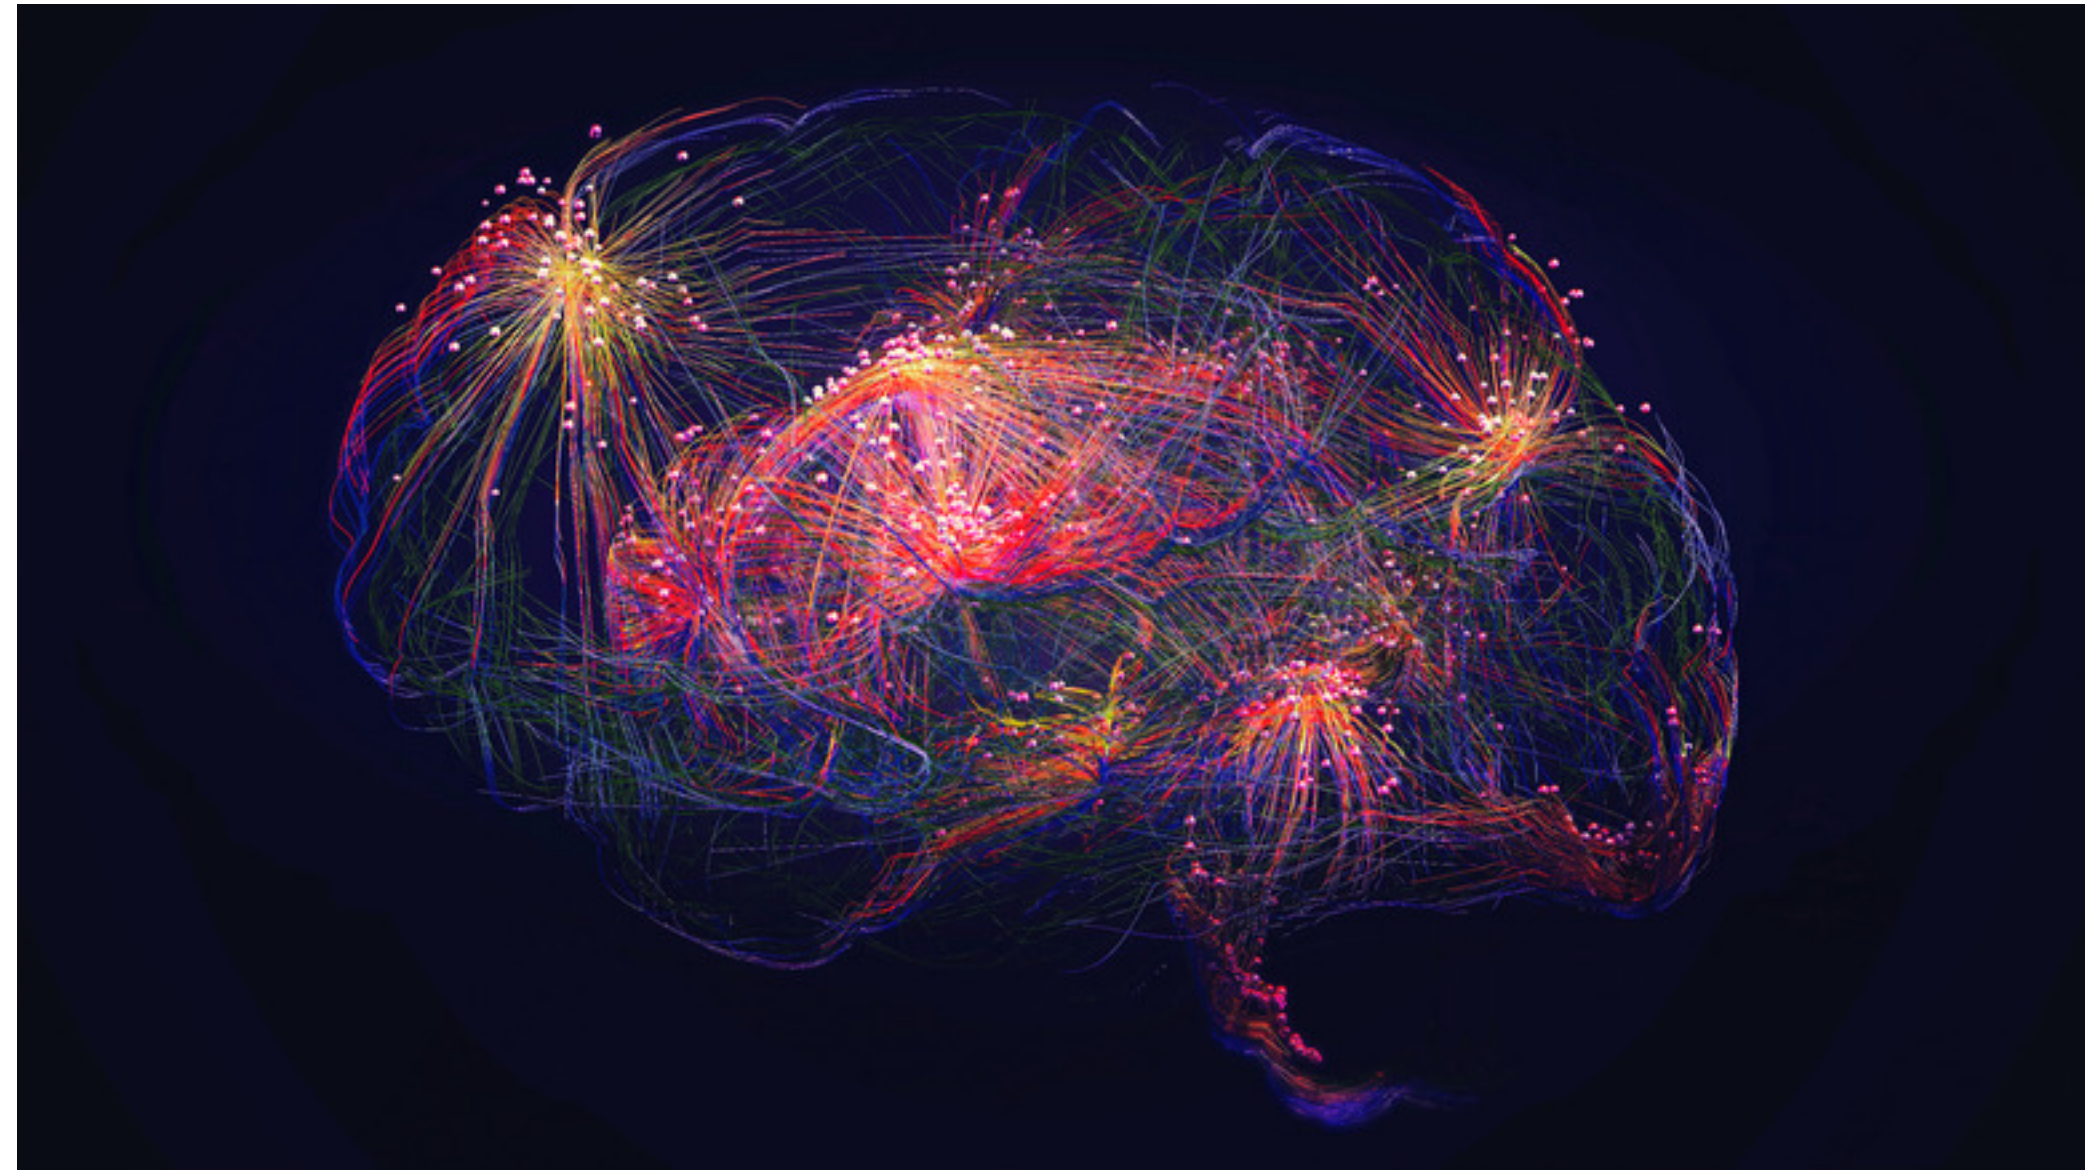


Artificial Neural Networks



Artificial Neural Networks inspired by brain research

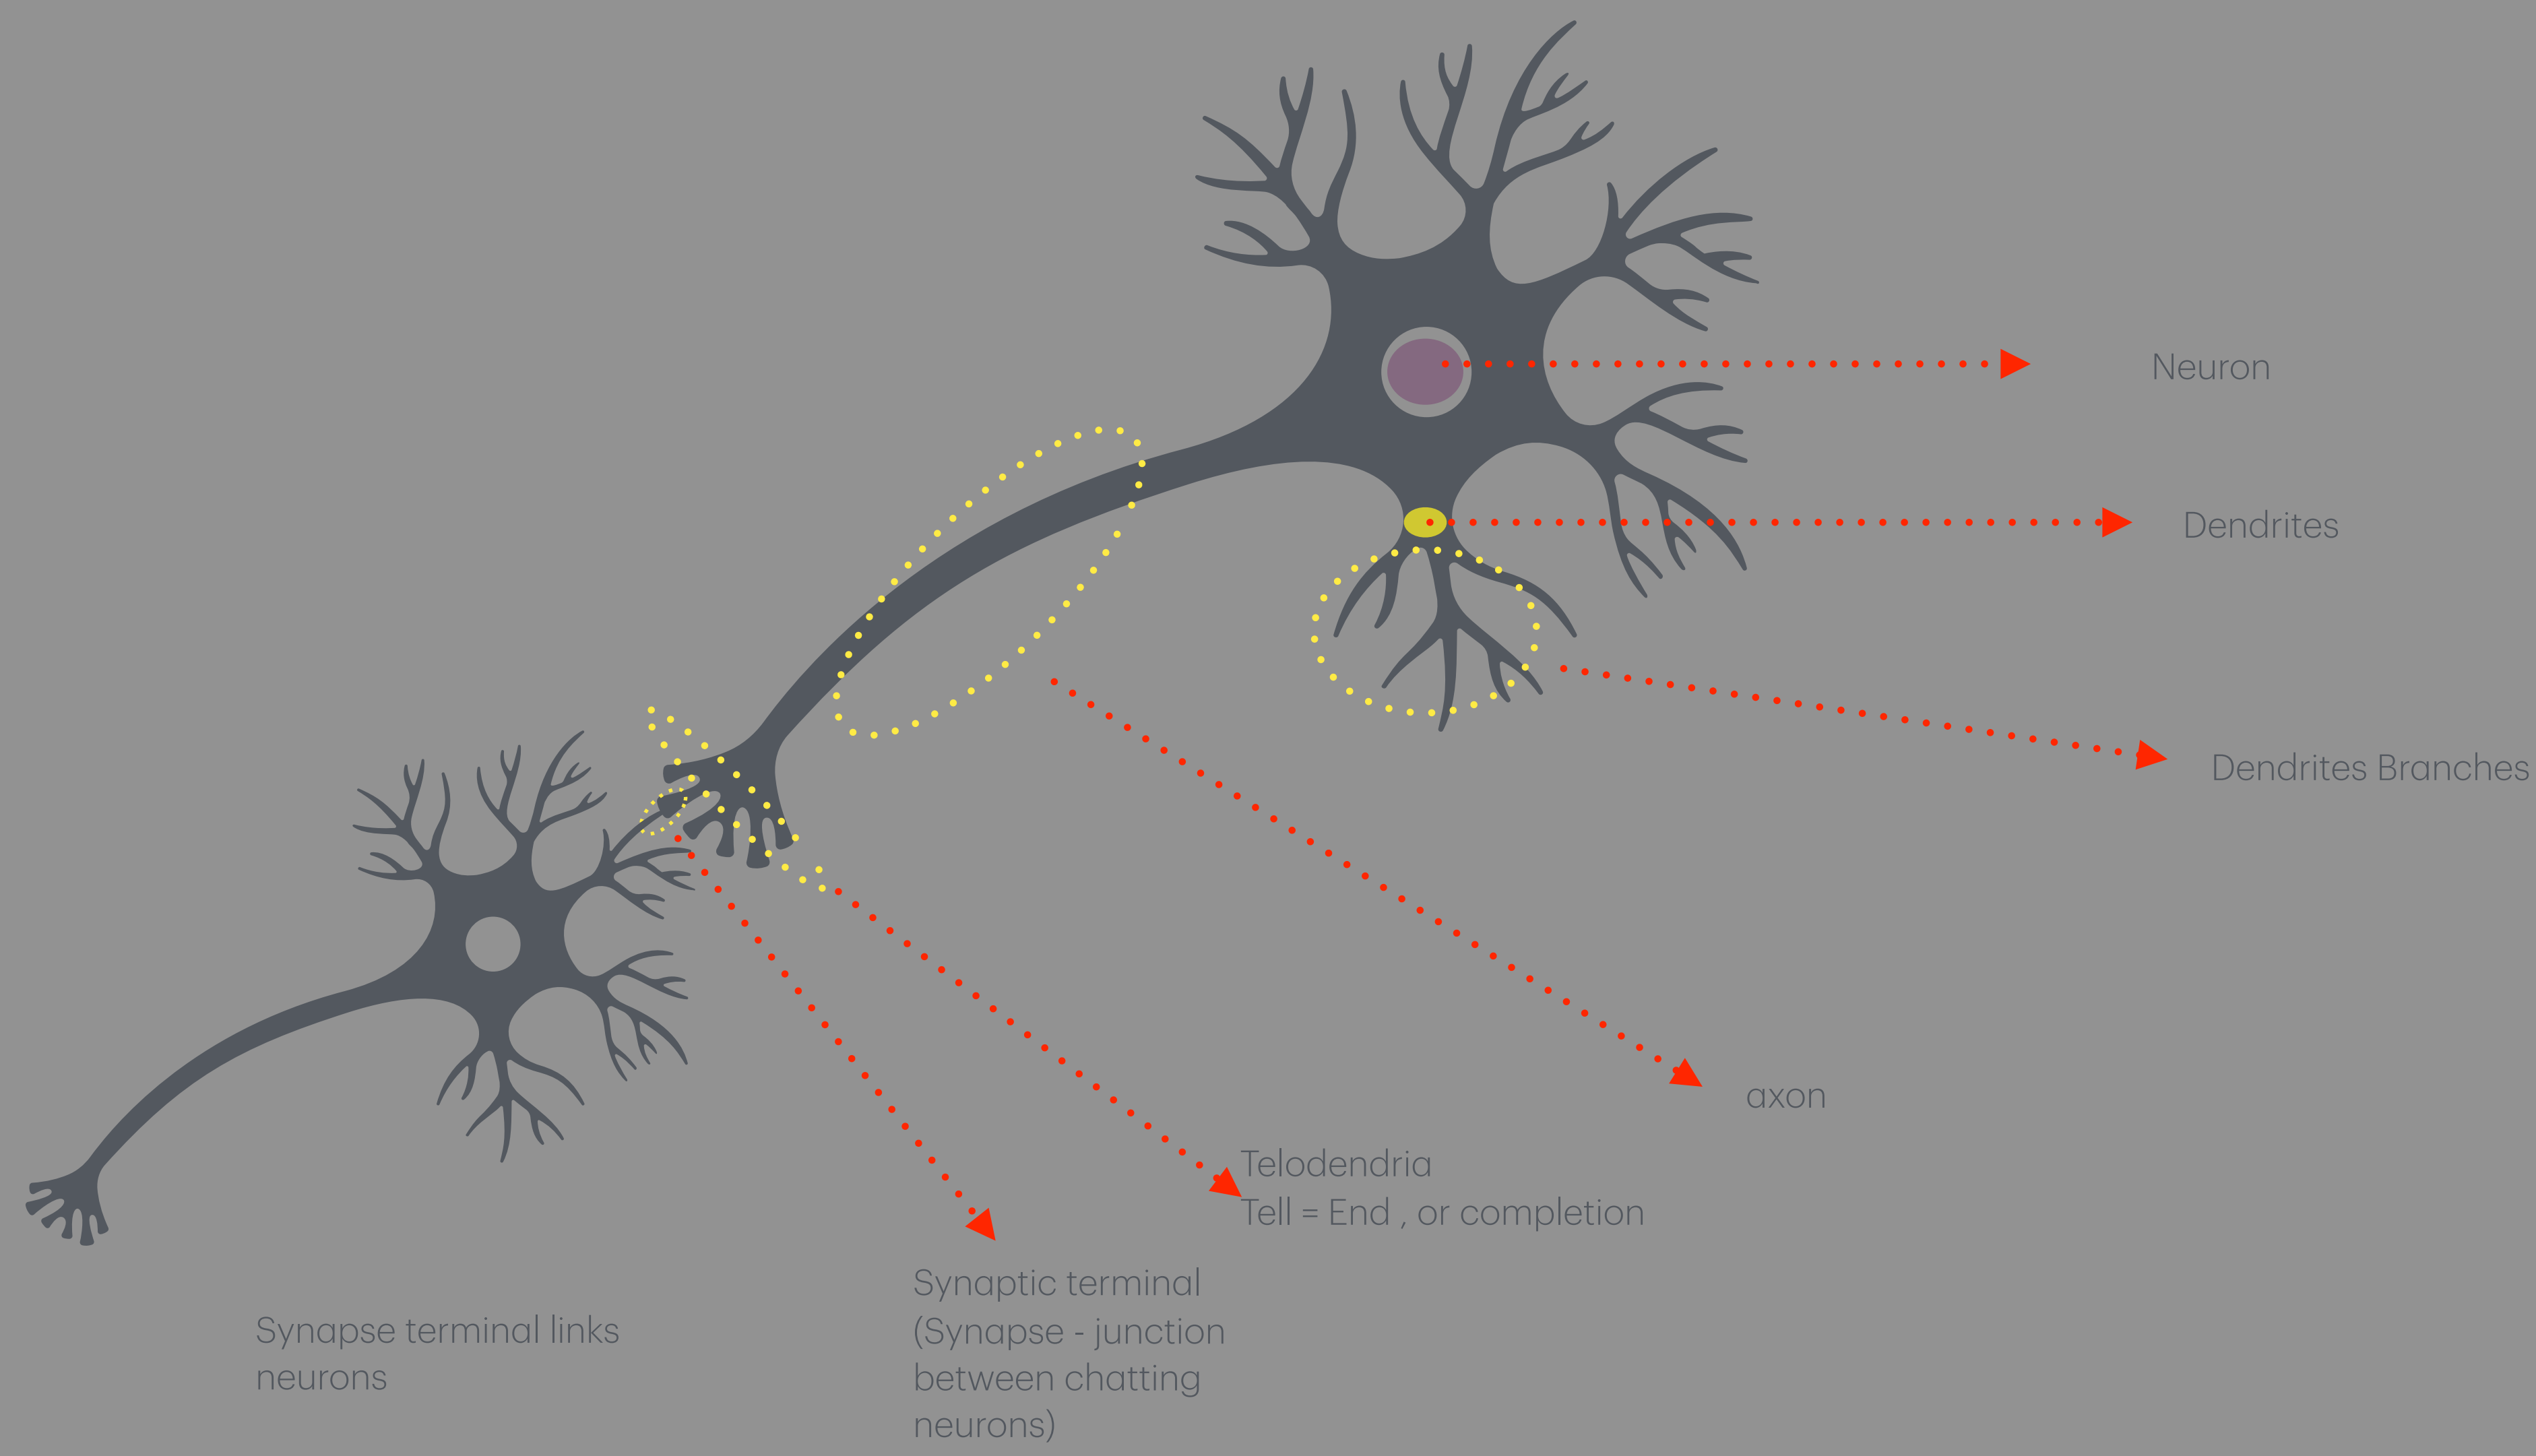
Speech Recognition

Google Images (Classification)


Go (DeepMind)

Recommendation Systems
(Youtube)


Biological Neuron



Biological Neuron



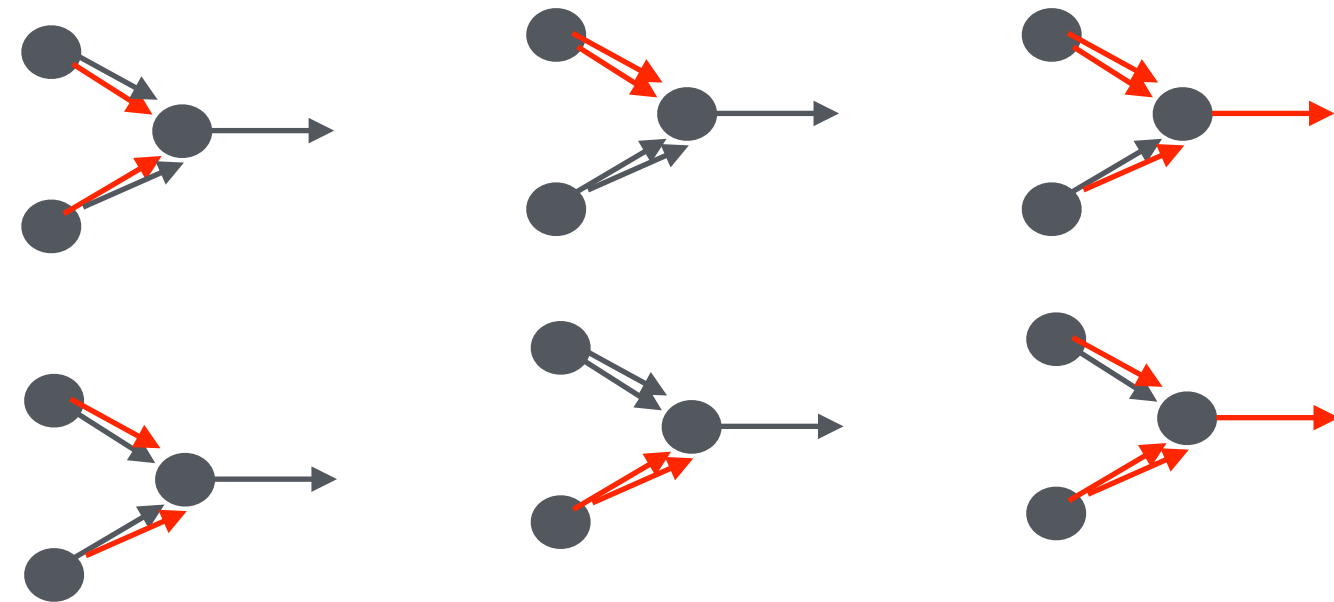
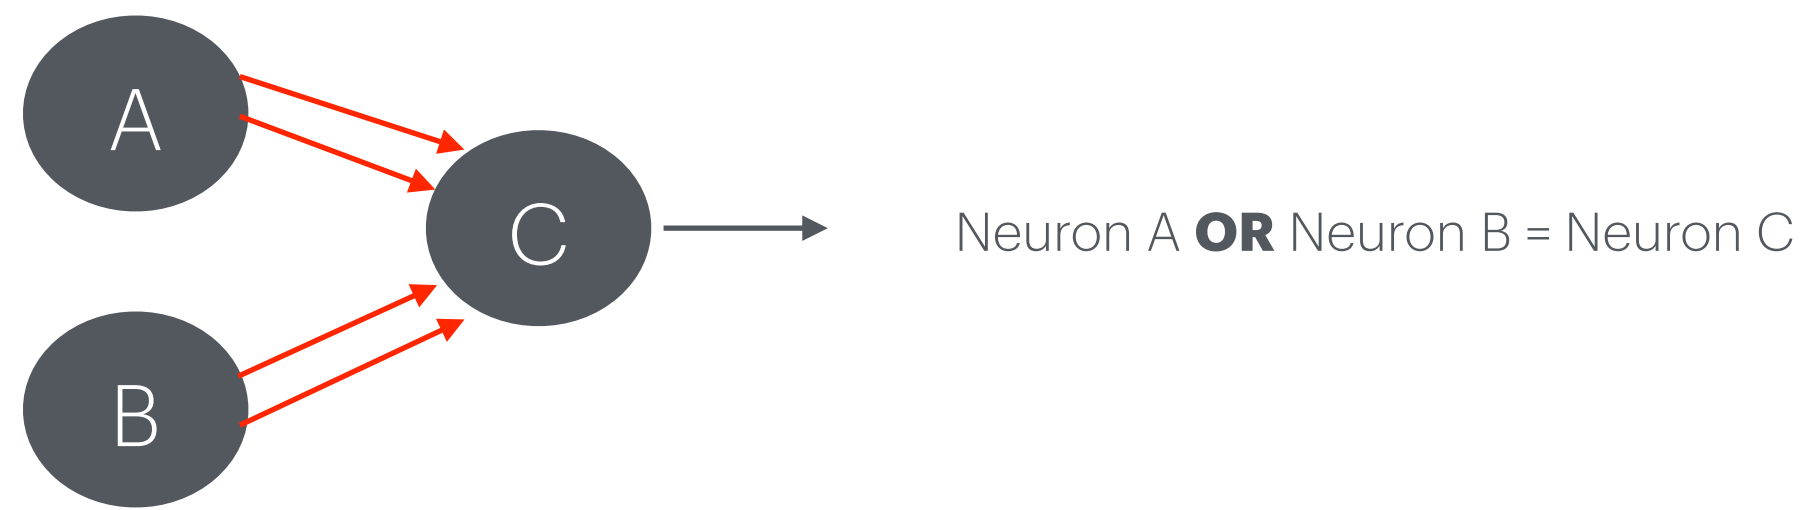
Logical Computation: Artificial Neuron








Note: Neuron activated
when at least two inputs
are active

Note: consumer neuron
C sums all the inputs

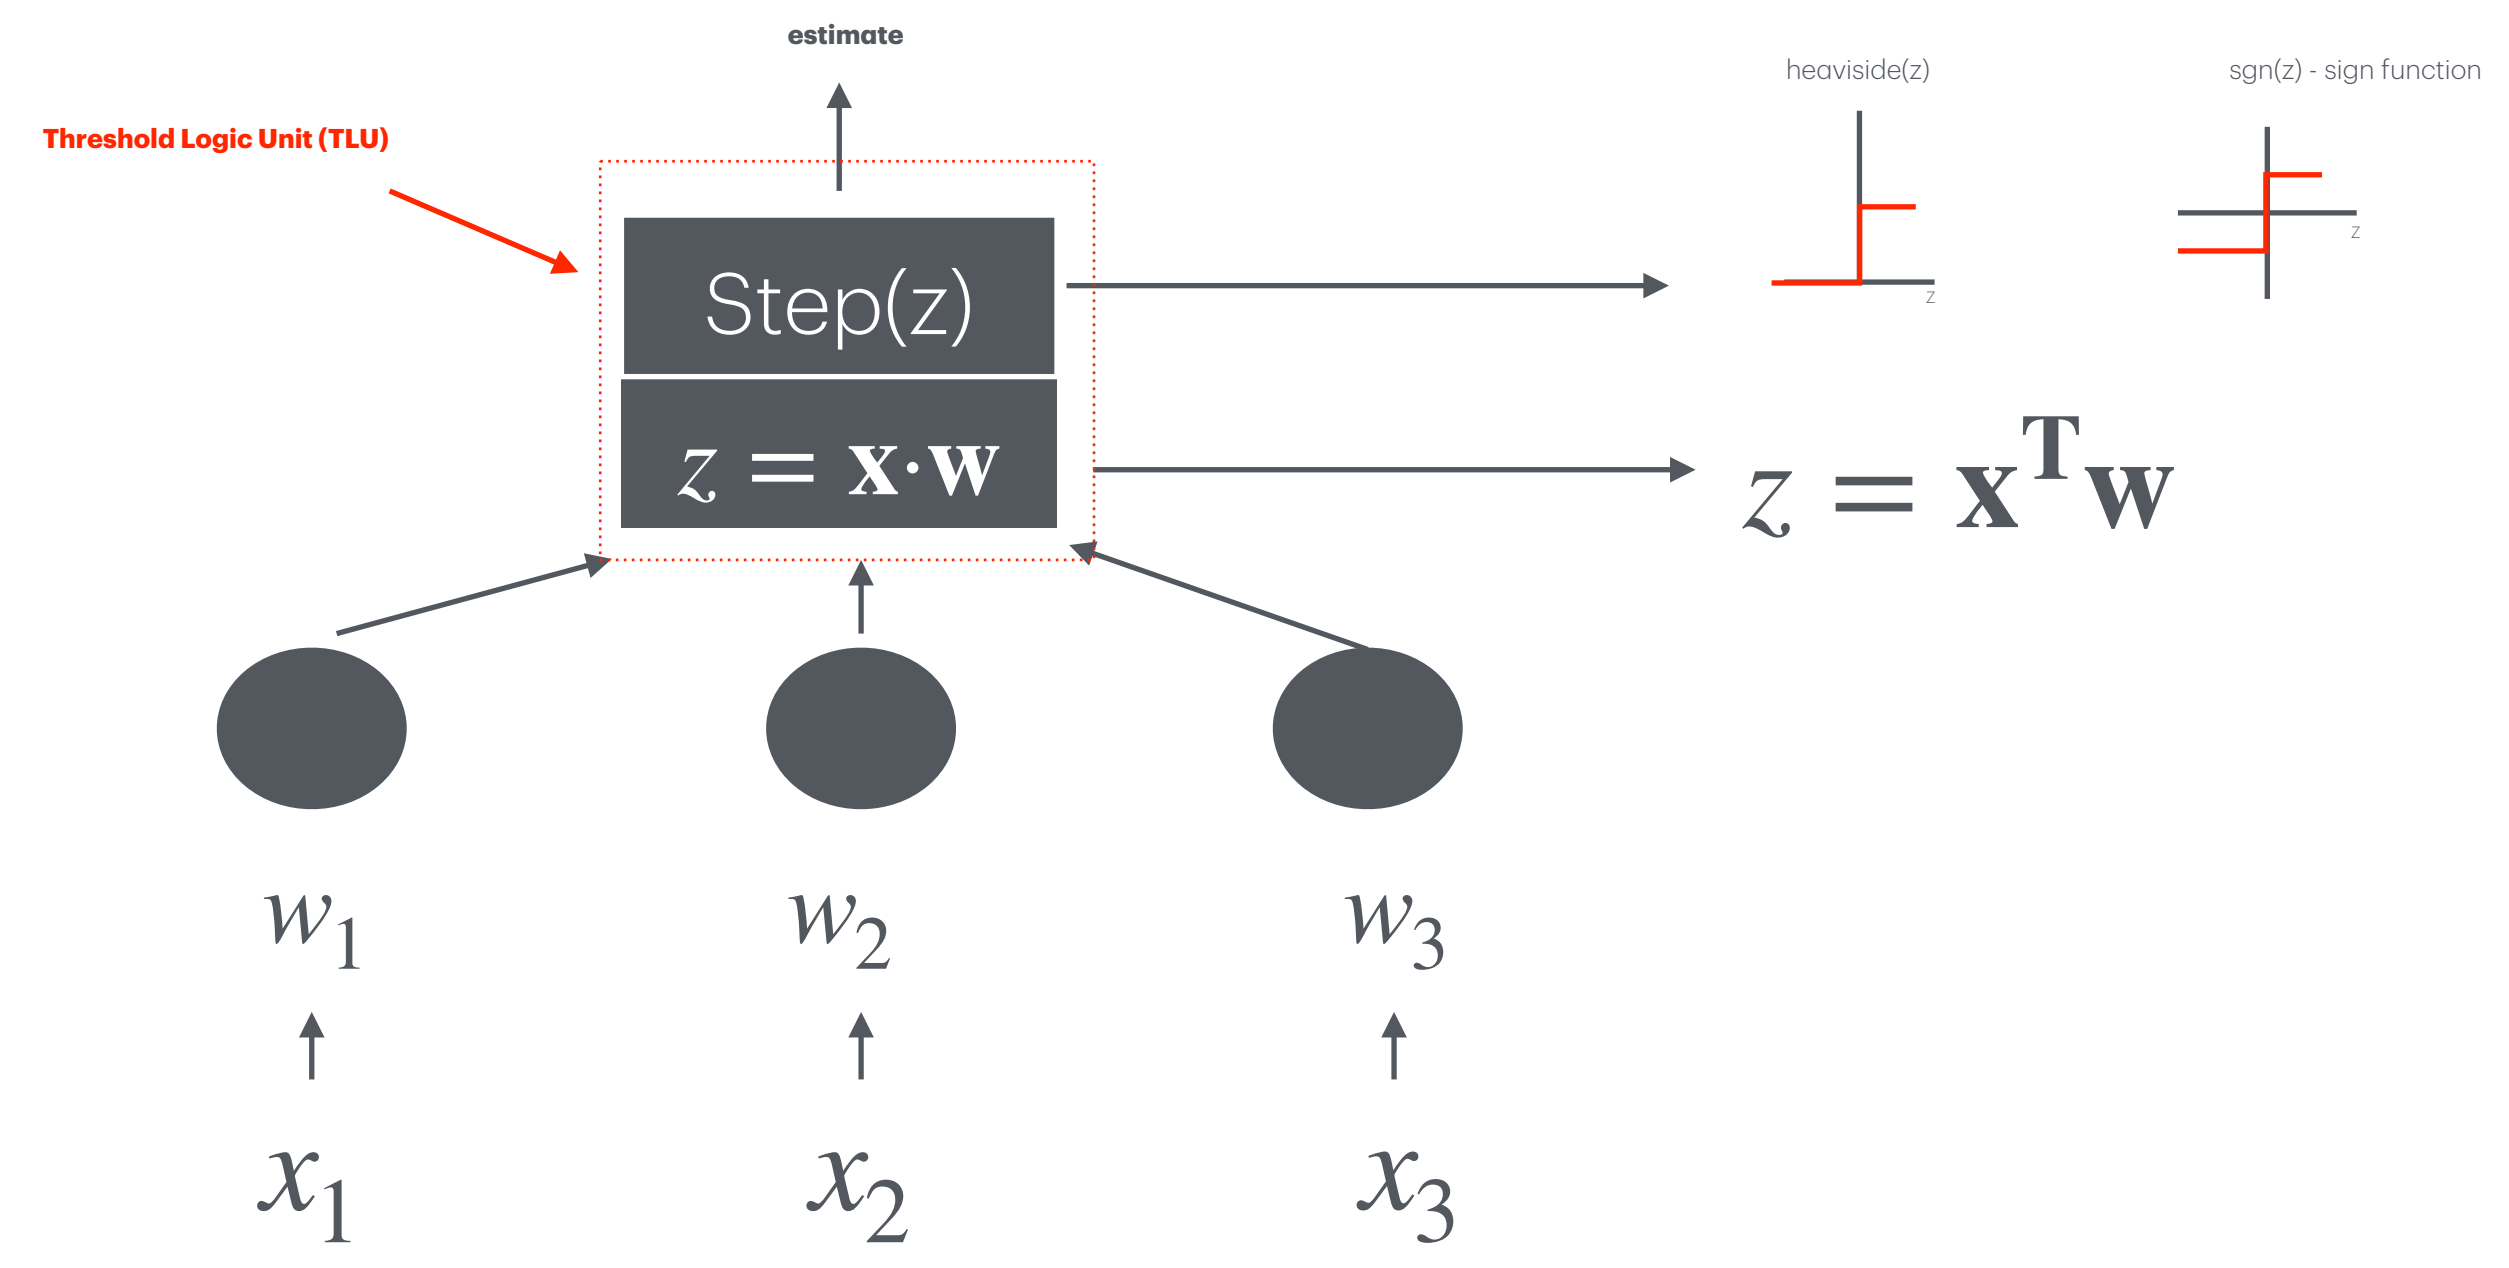
Logical Computation: Artificial Neuron



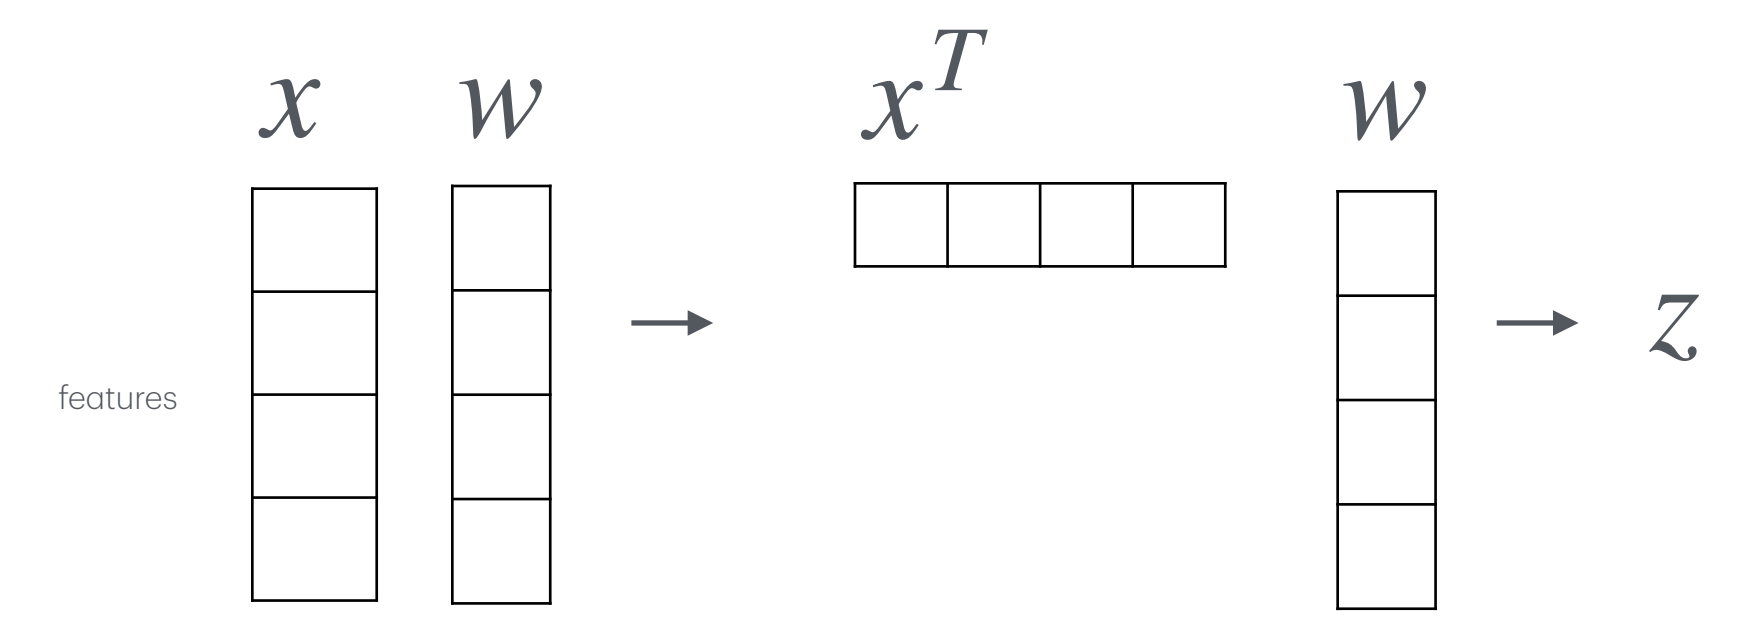
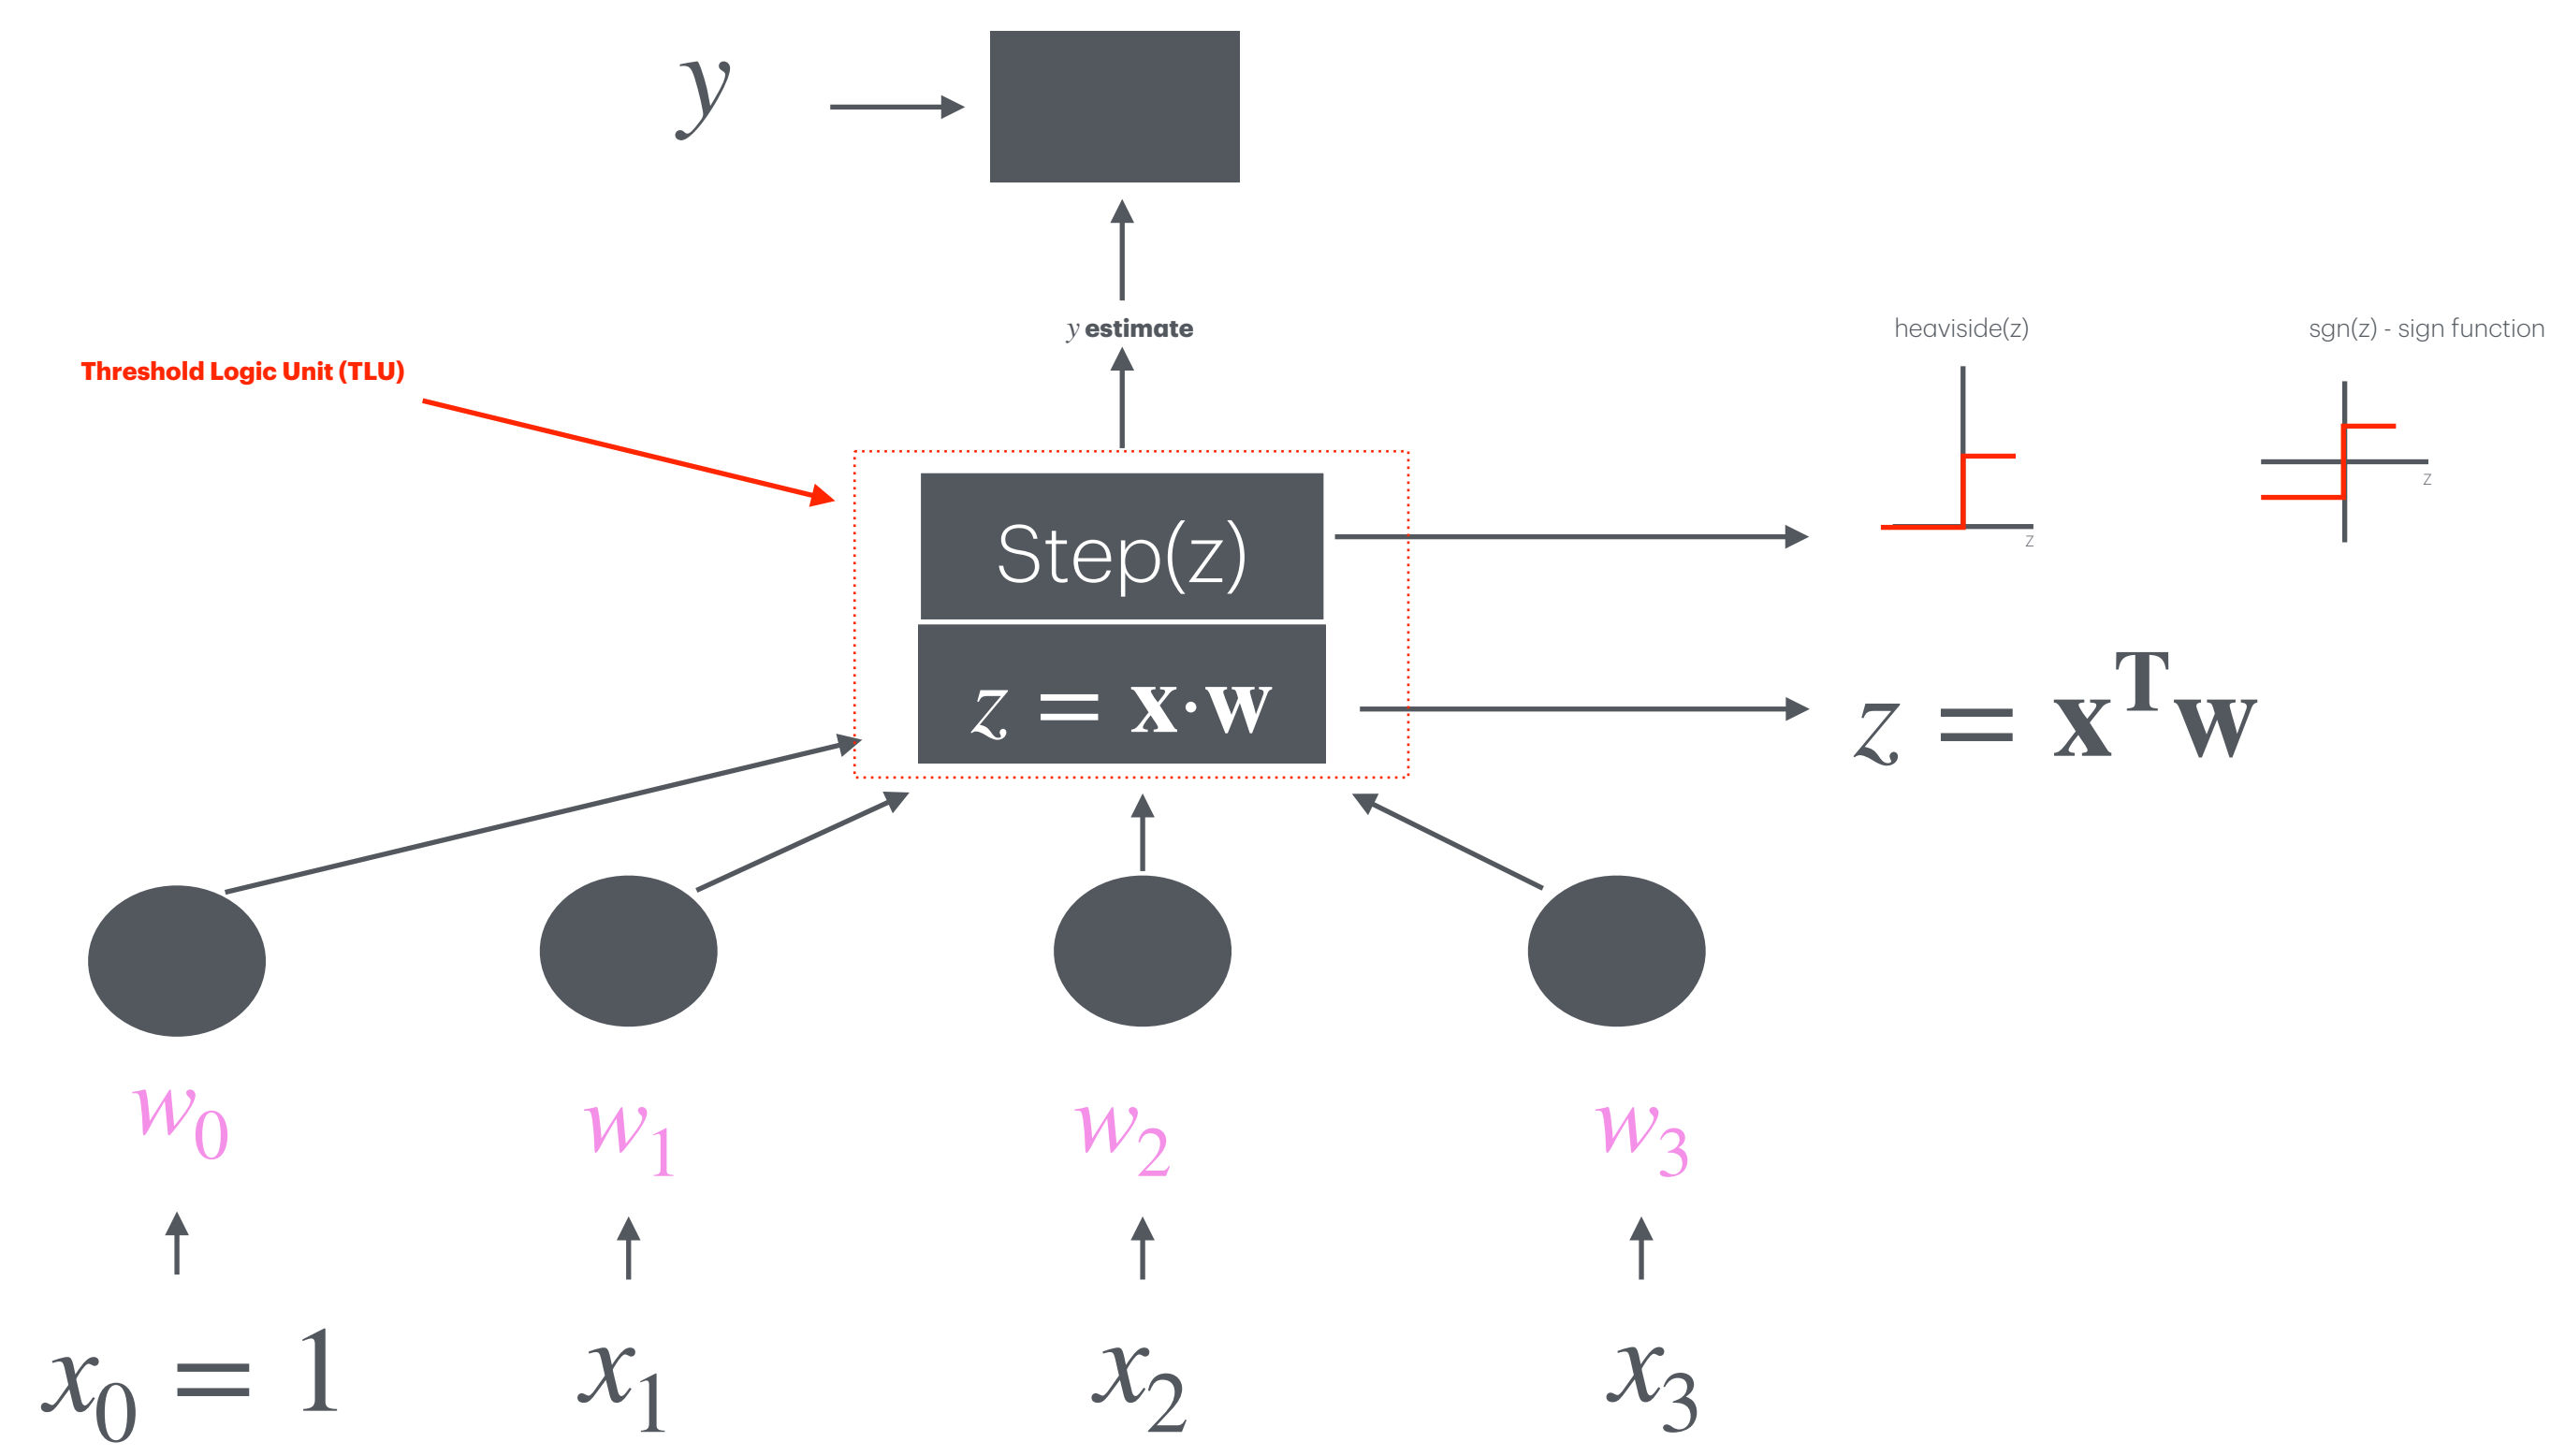
Logical Computation: Artificial Neuron



Logical Computation: Perception



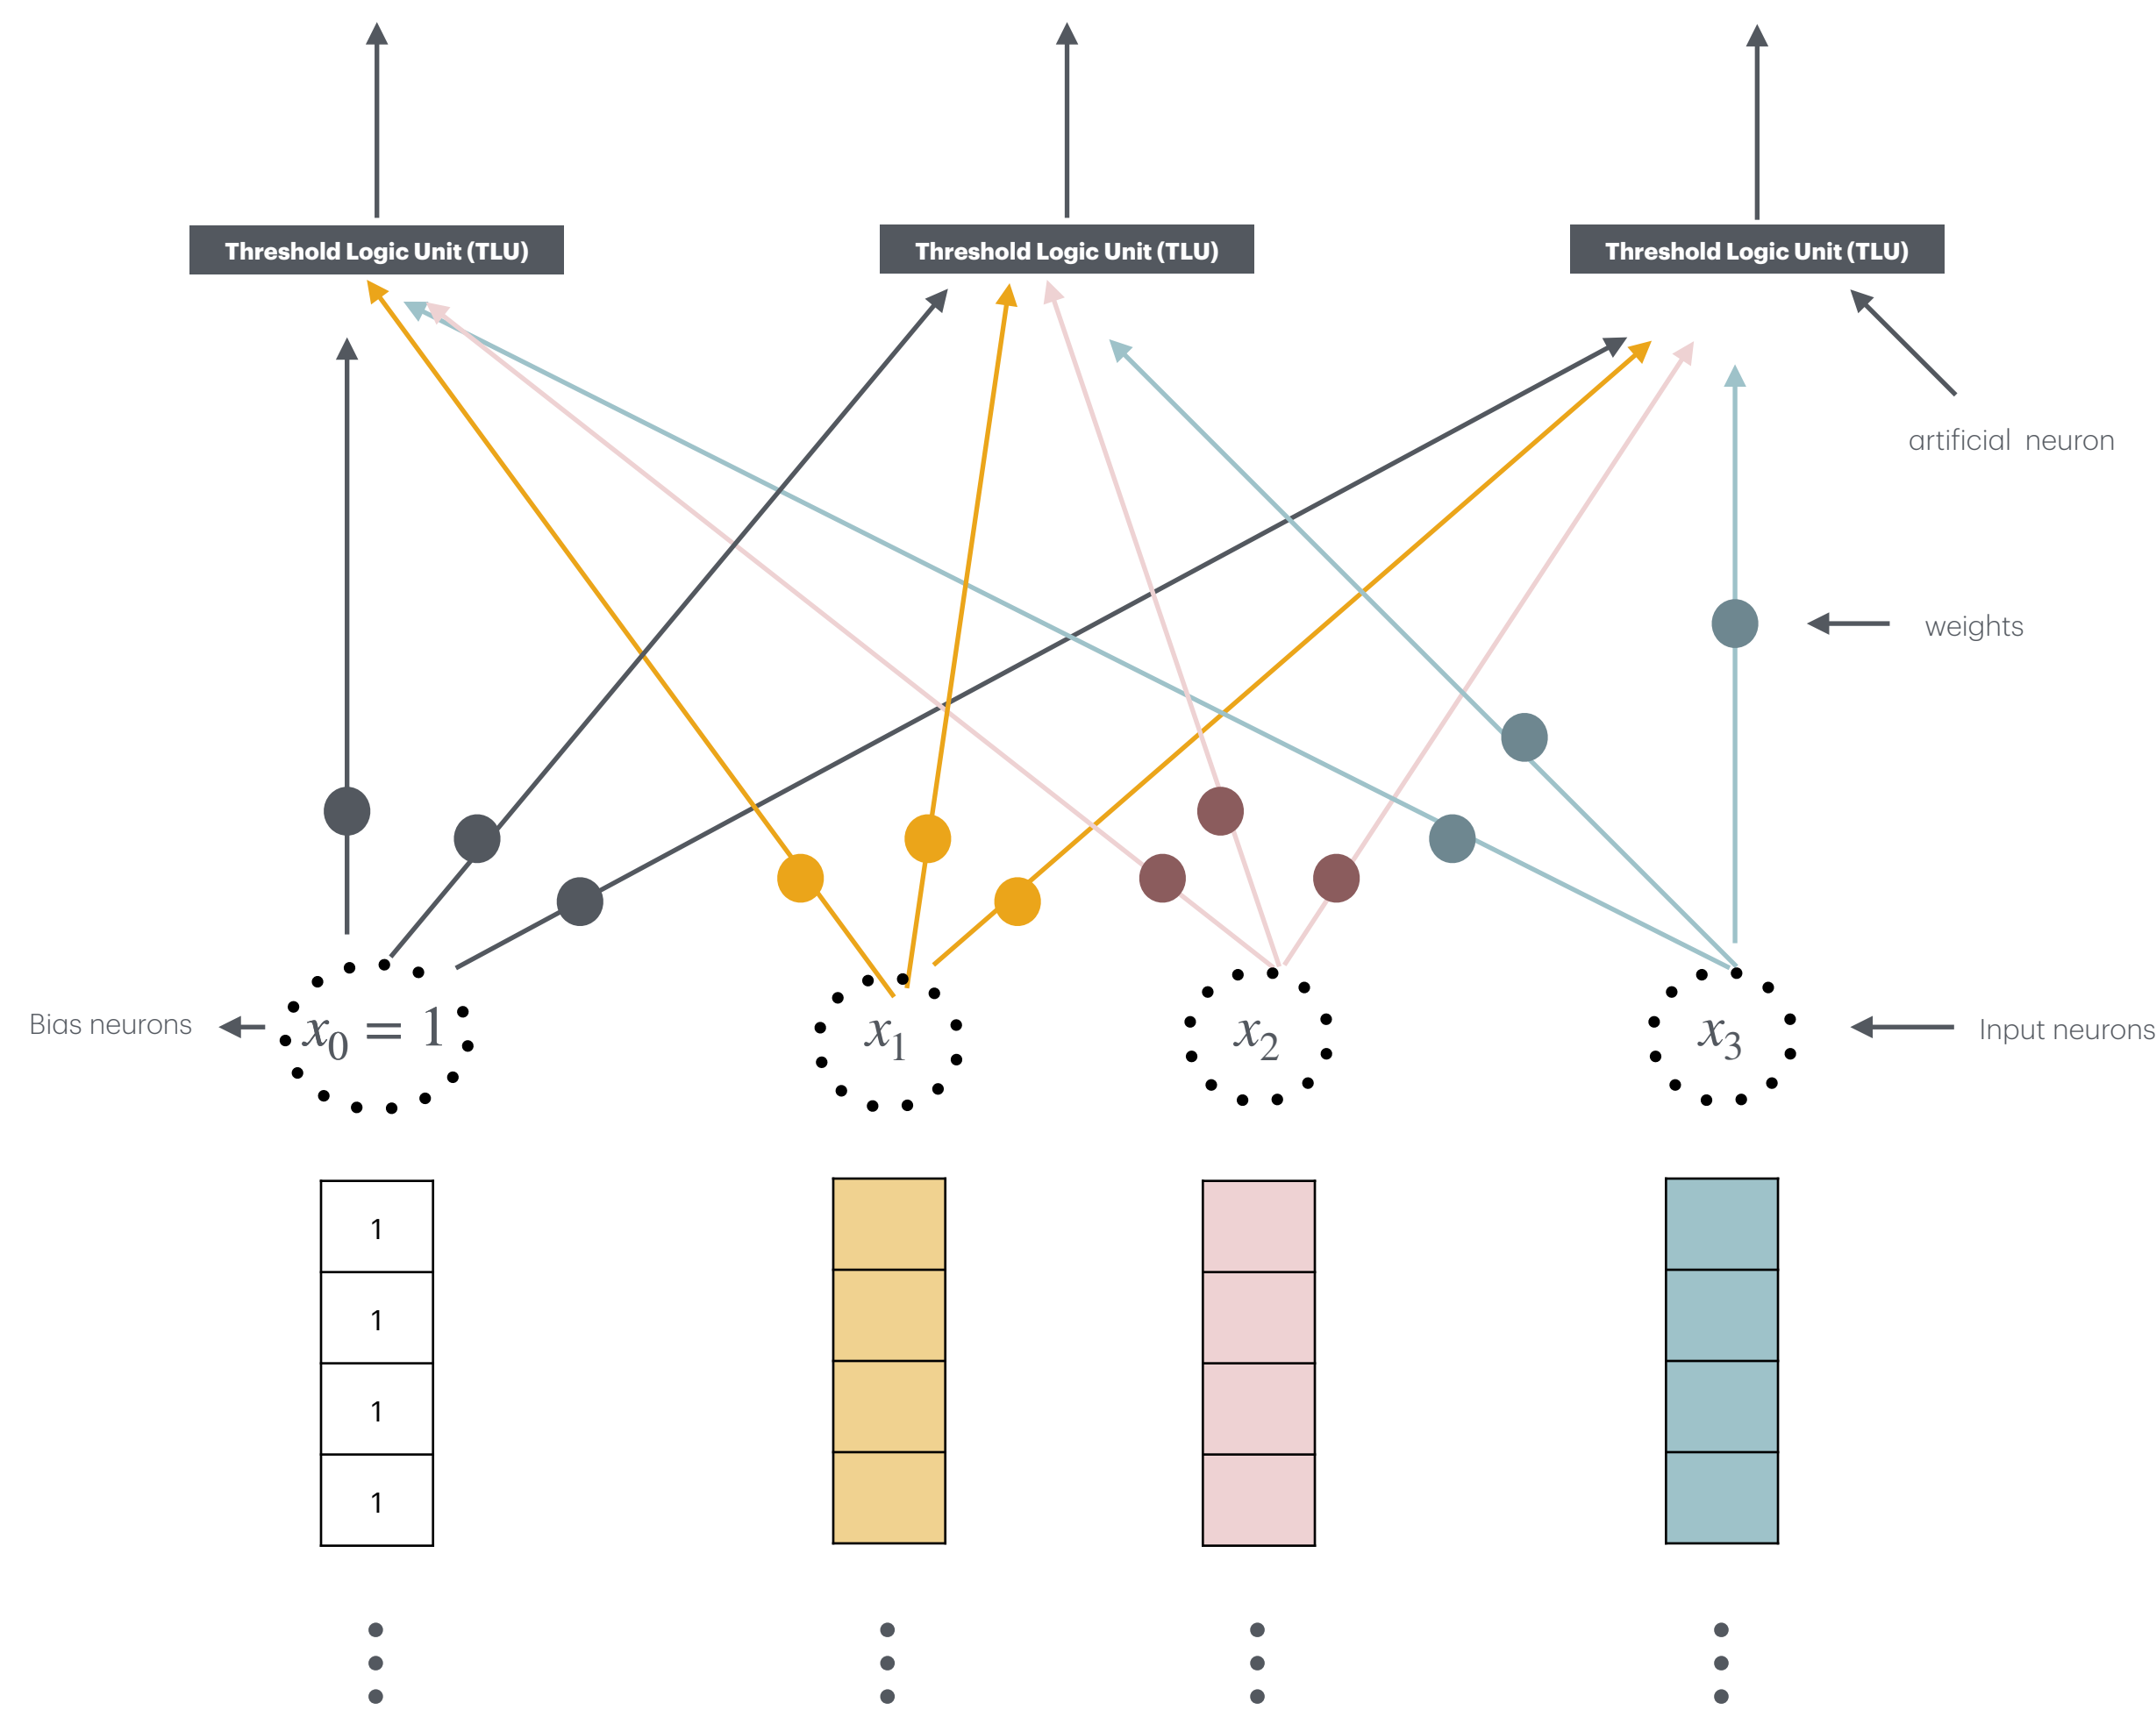
Logical Computation: Perception



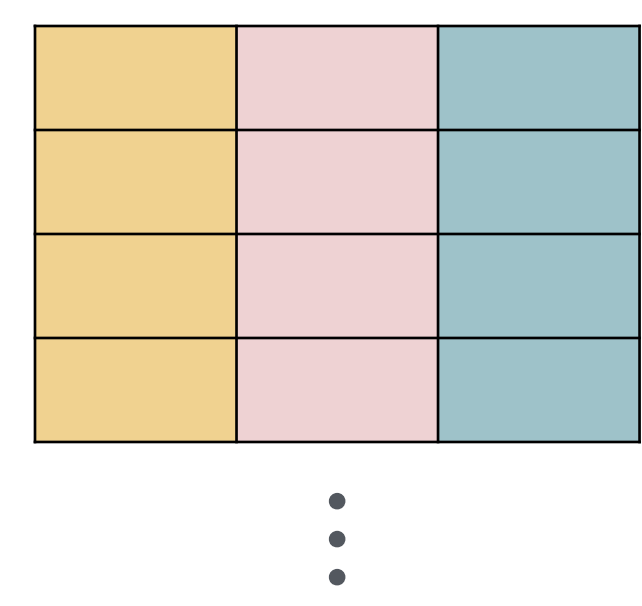
Find right values for weights that lower MSE with expectations

Used for simple linear classification (1D, 2D features)

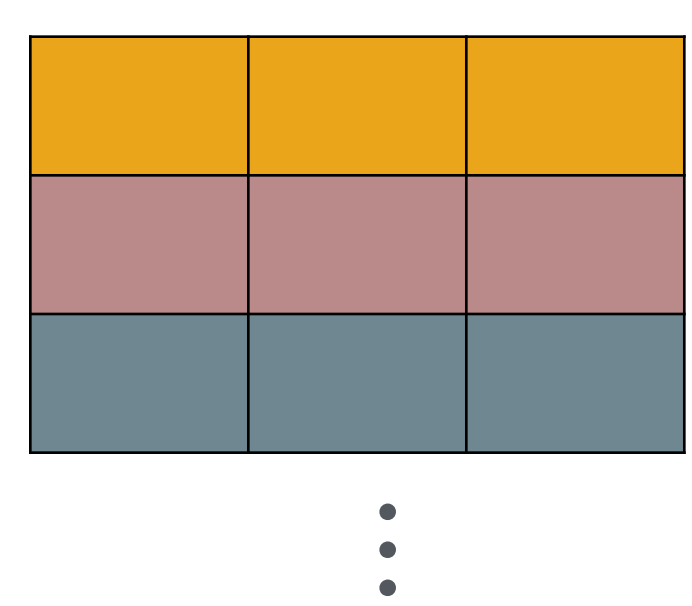
Logical Computation: Multi-Perception



X




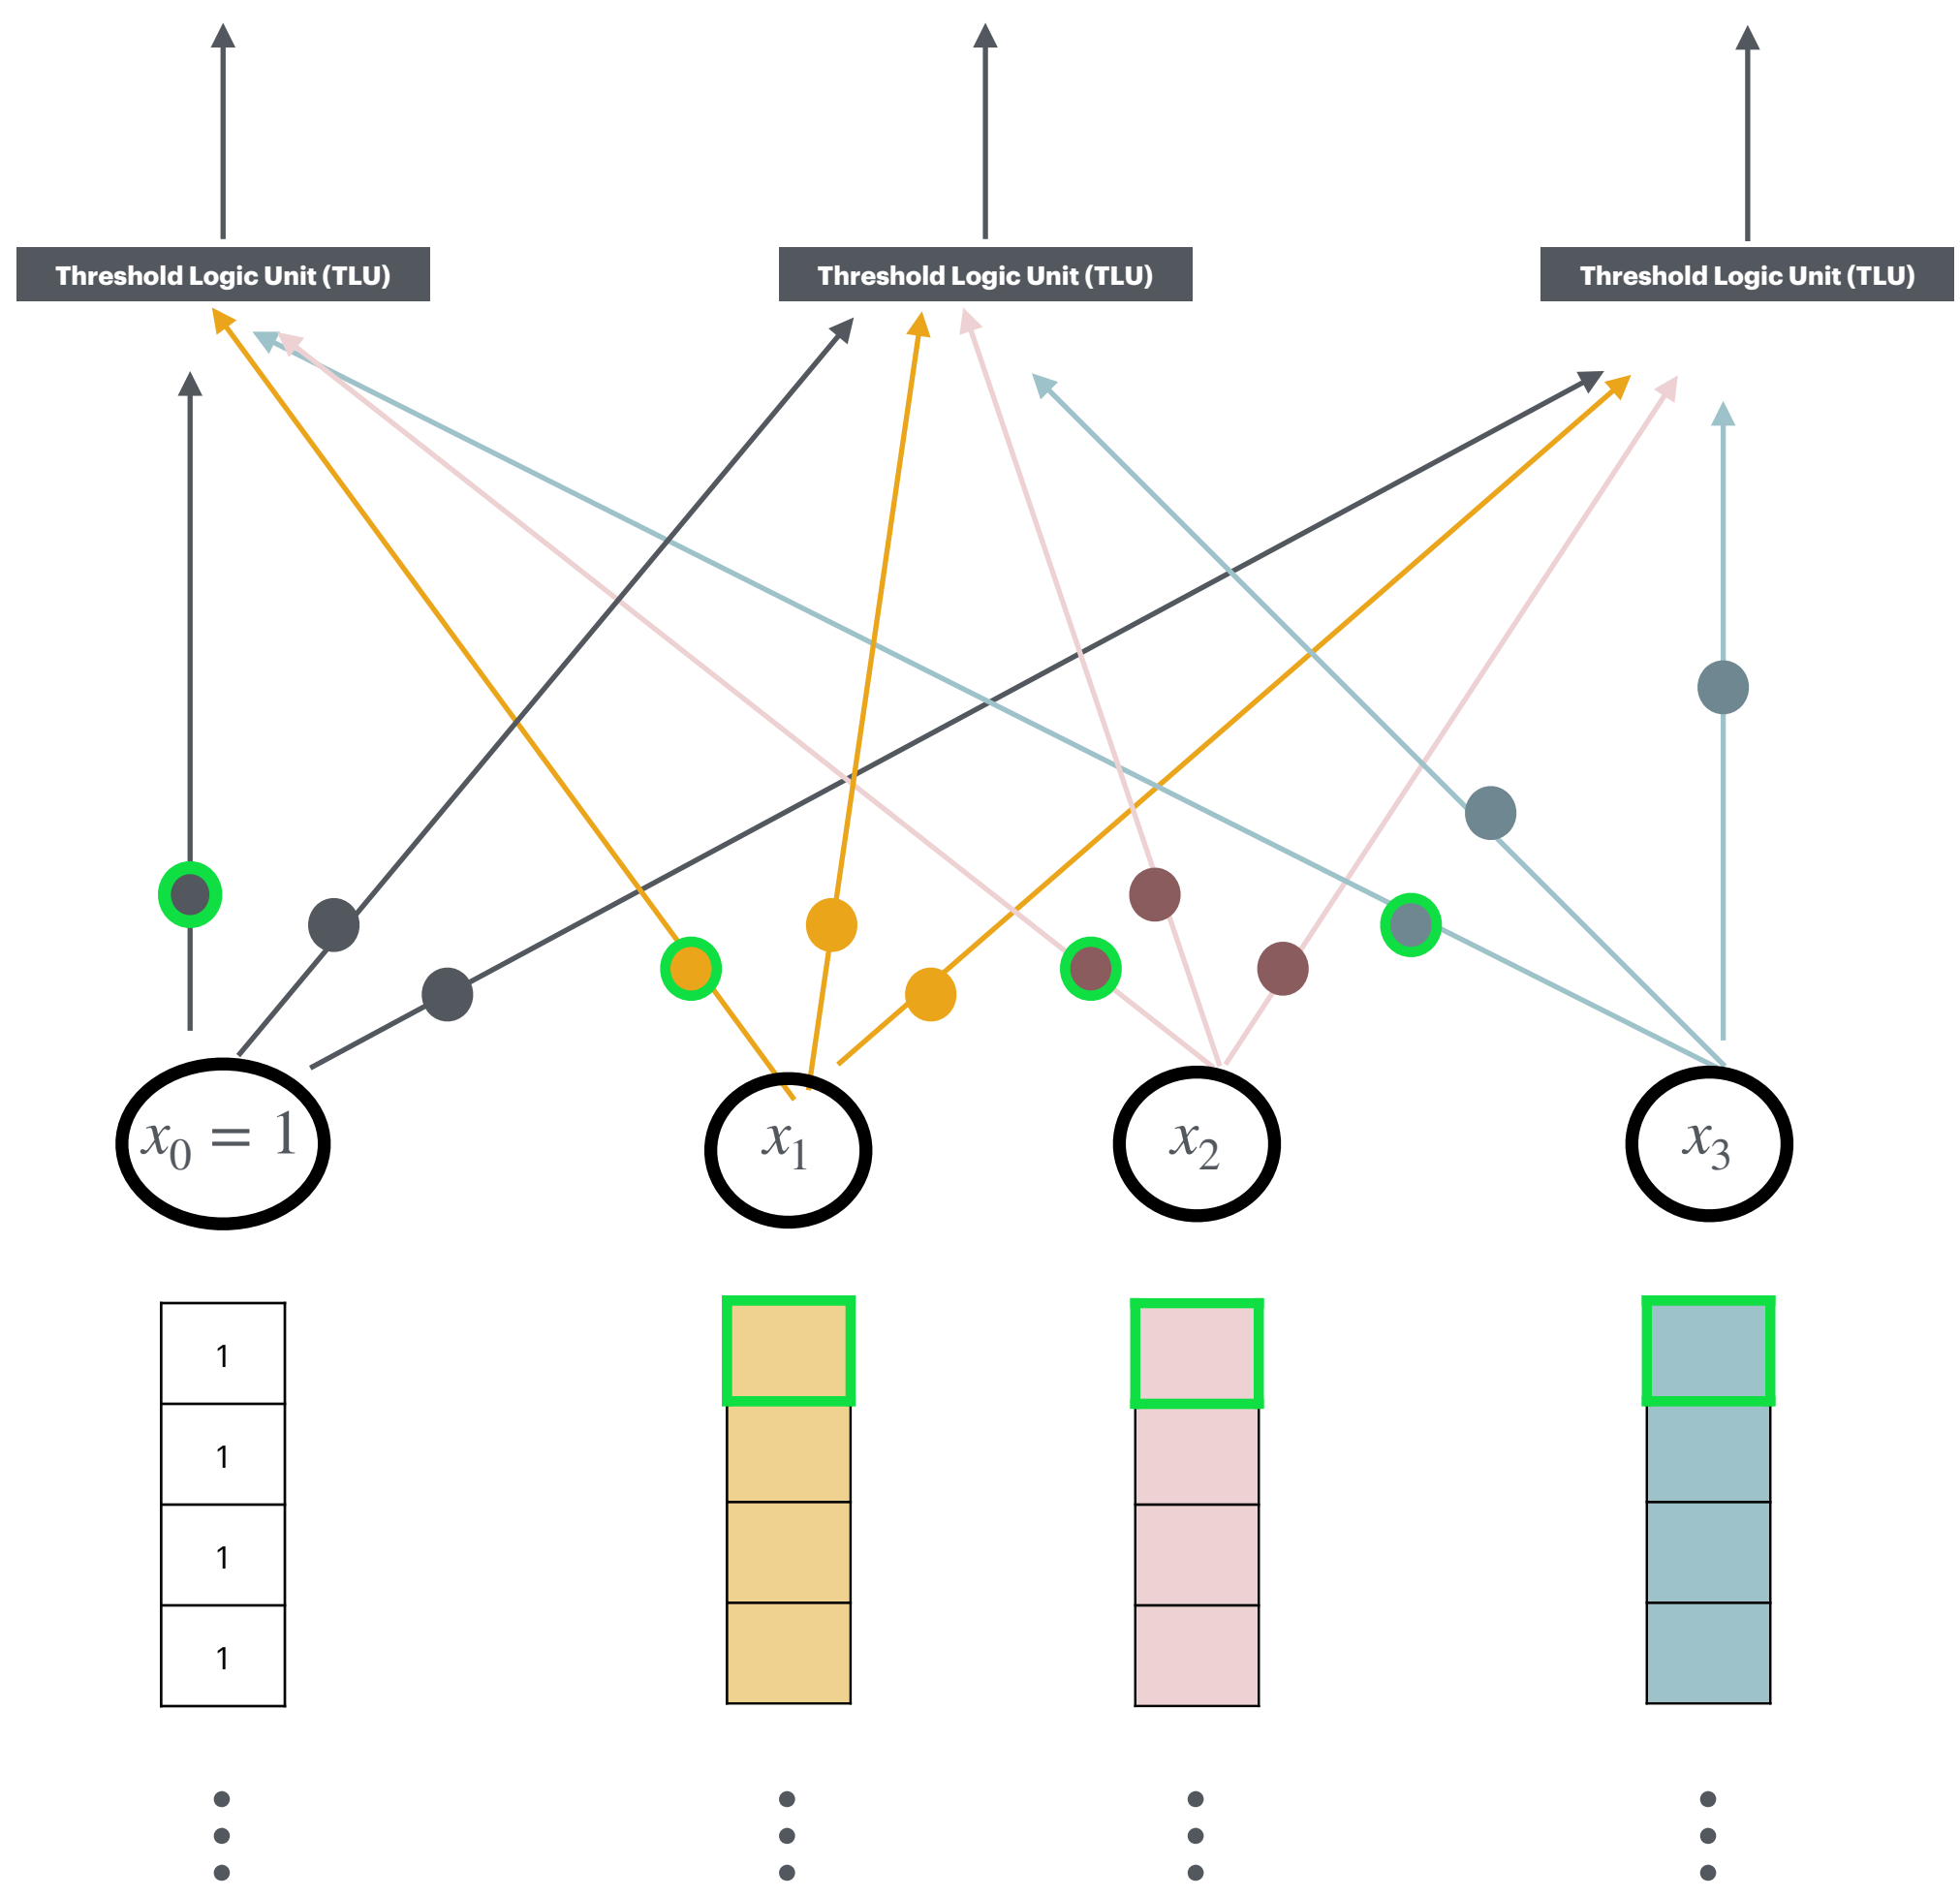
W





B



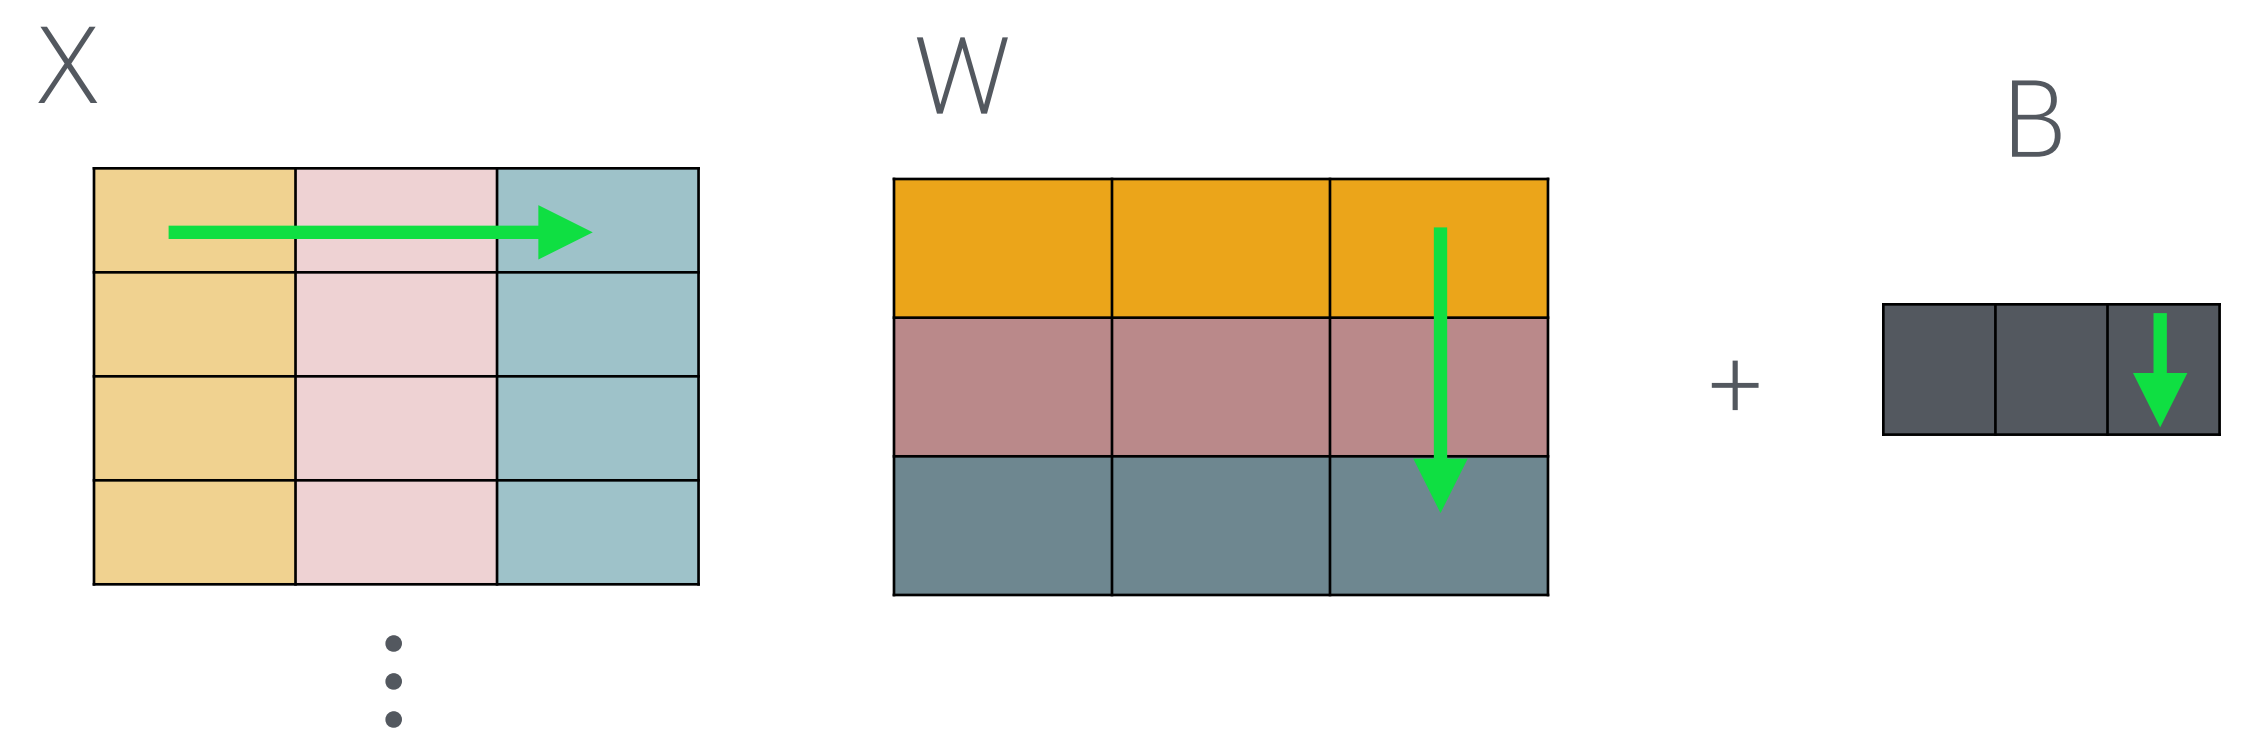
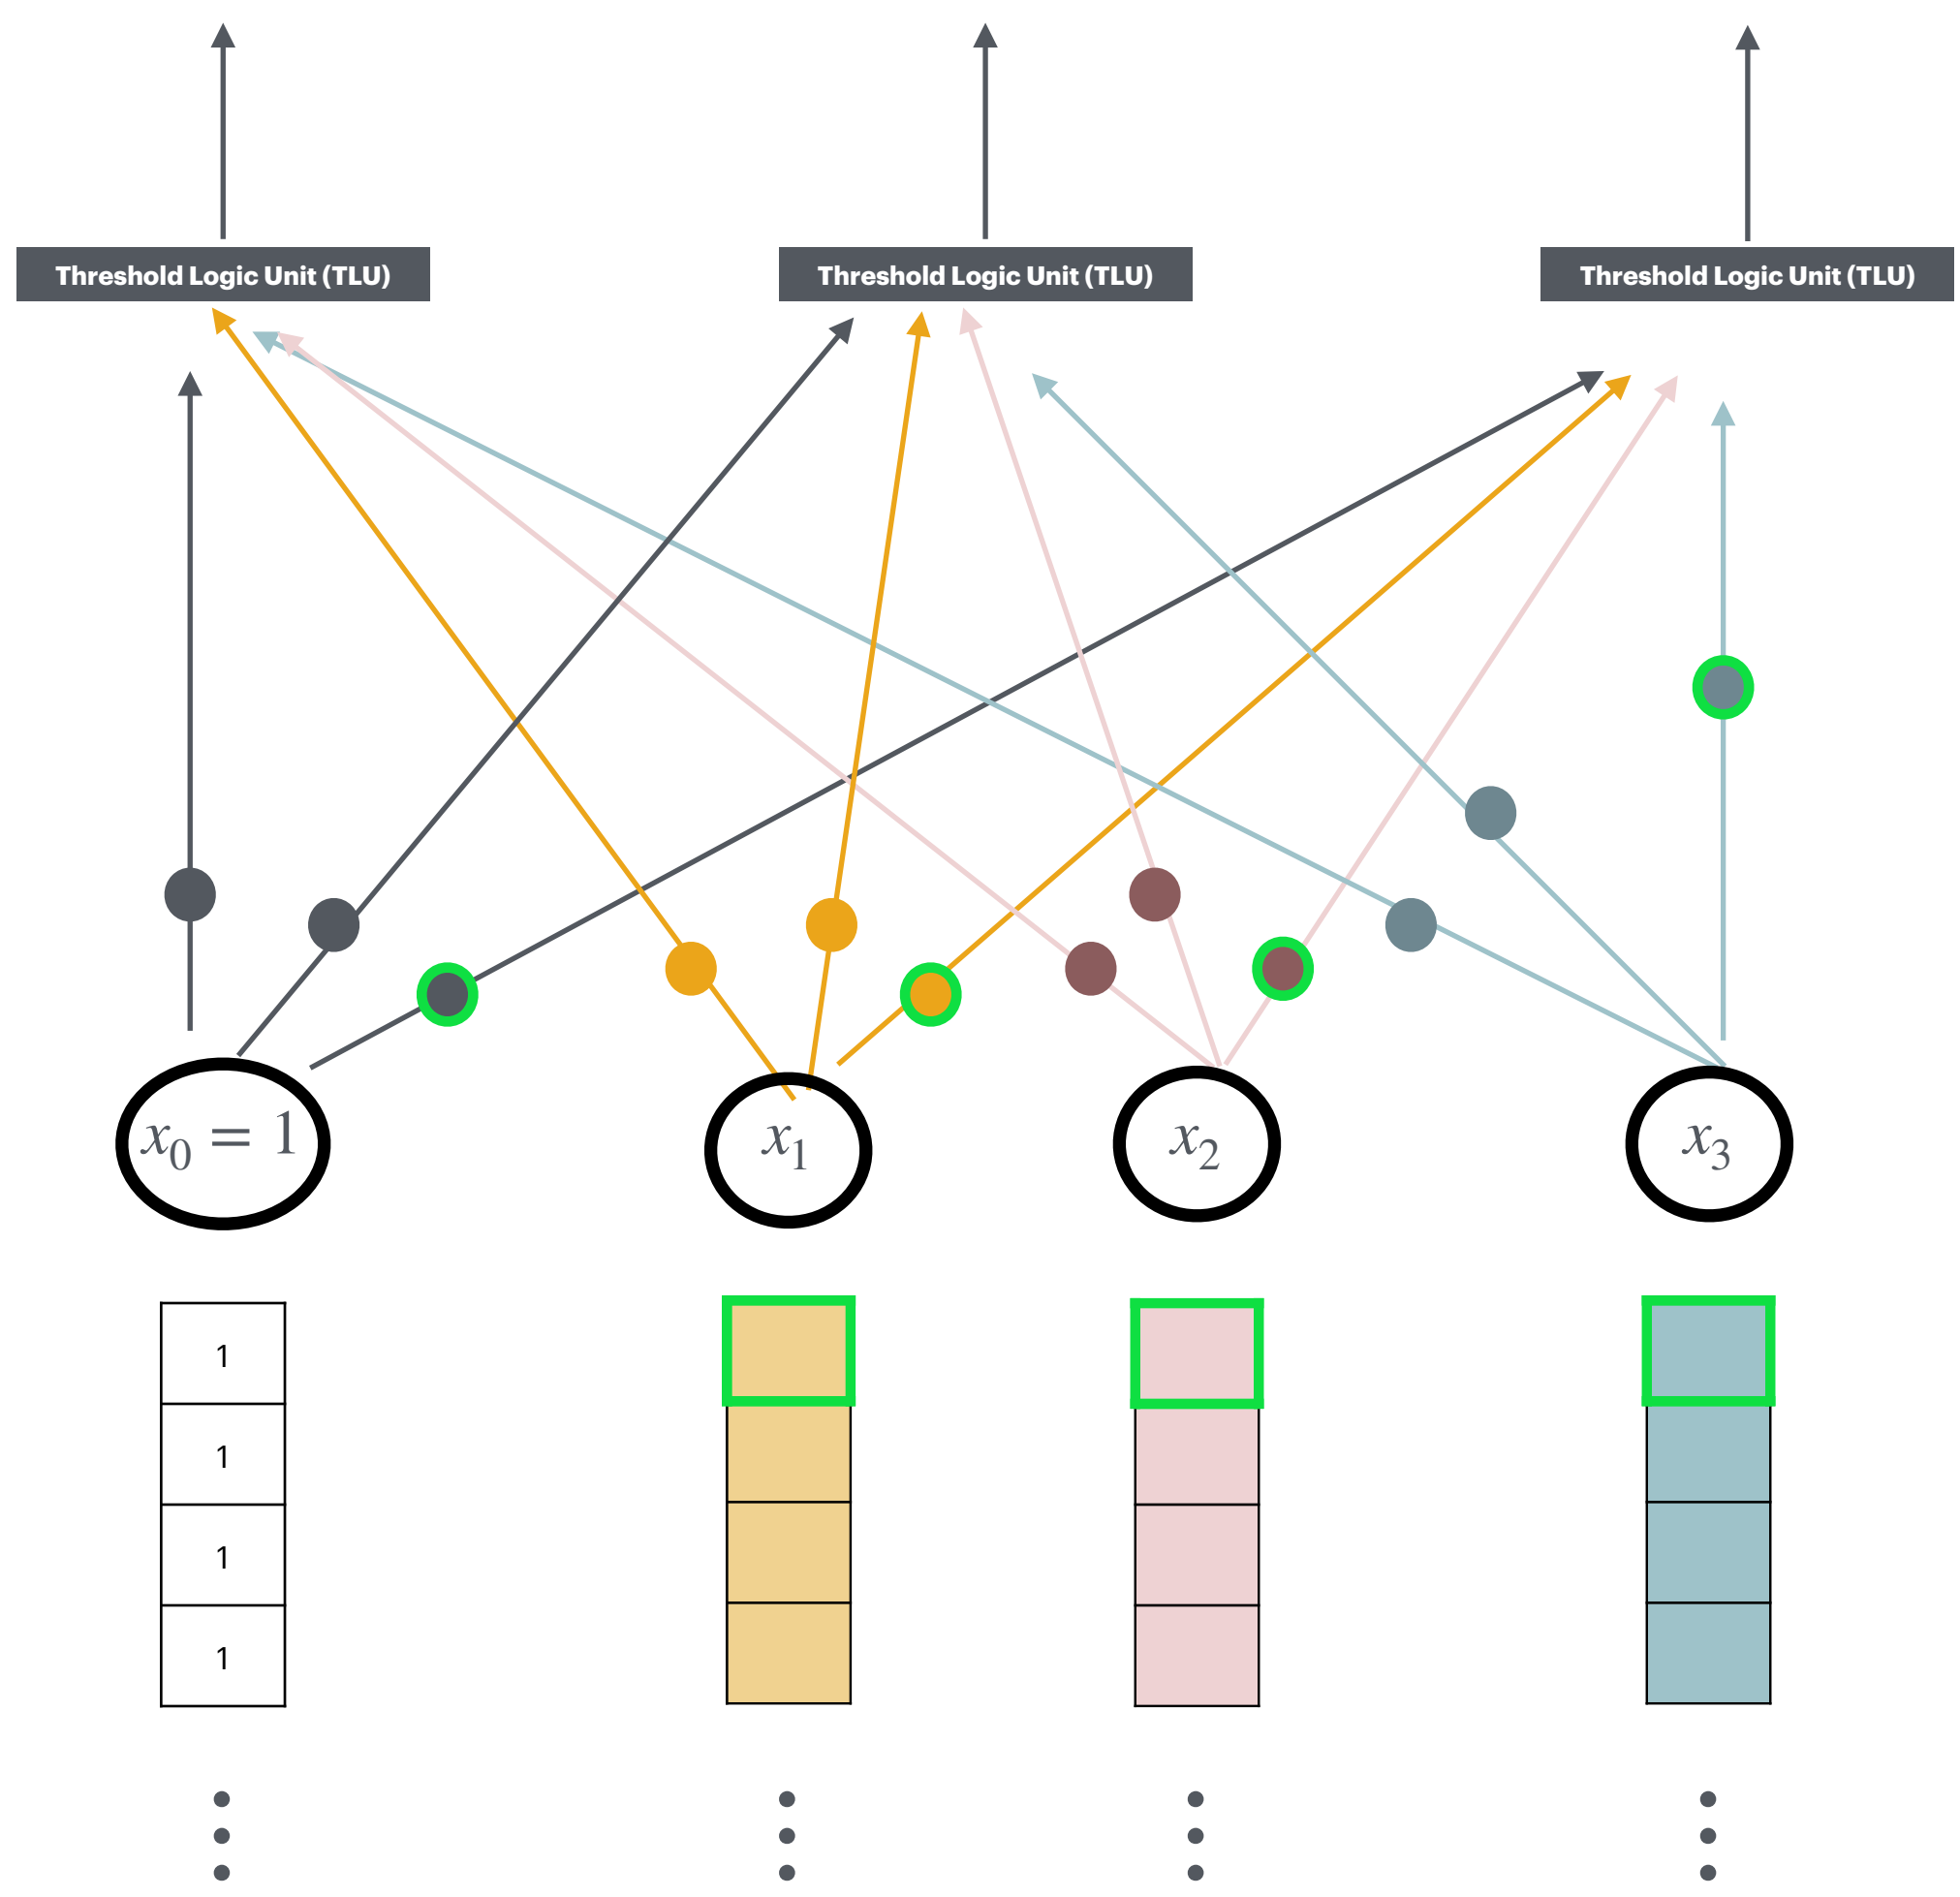
Logical Computation: Perception



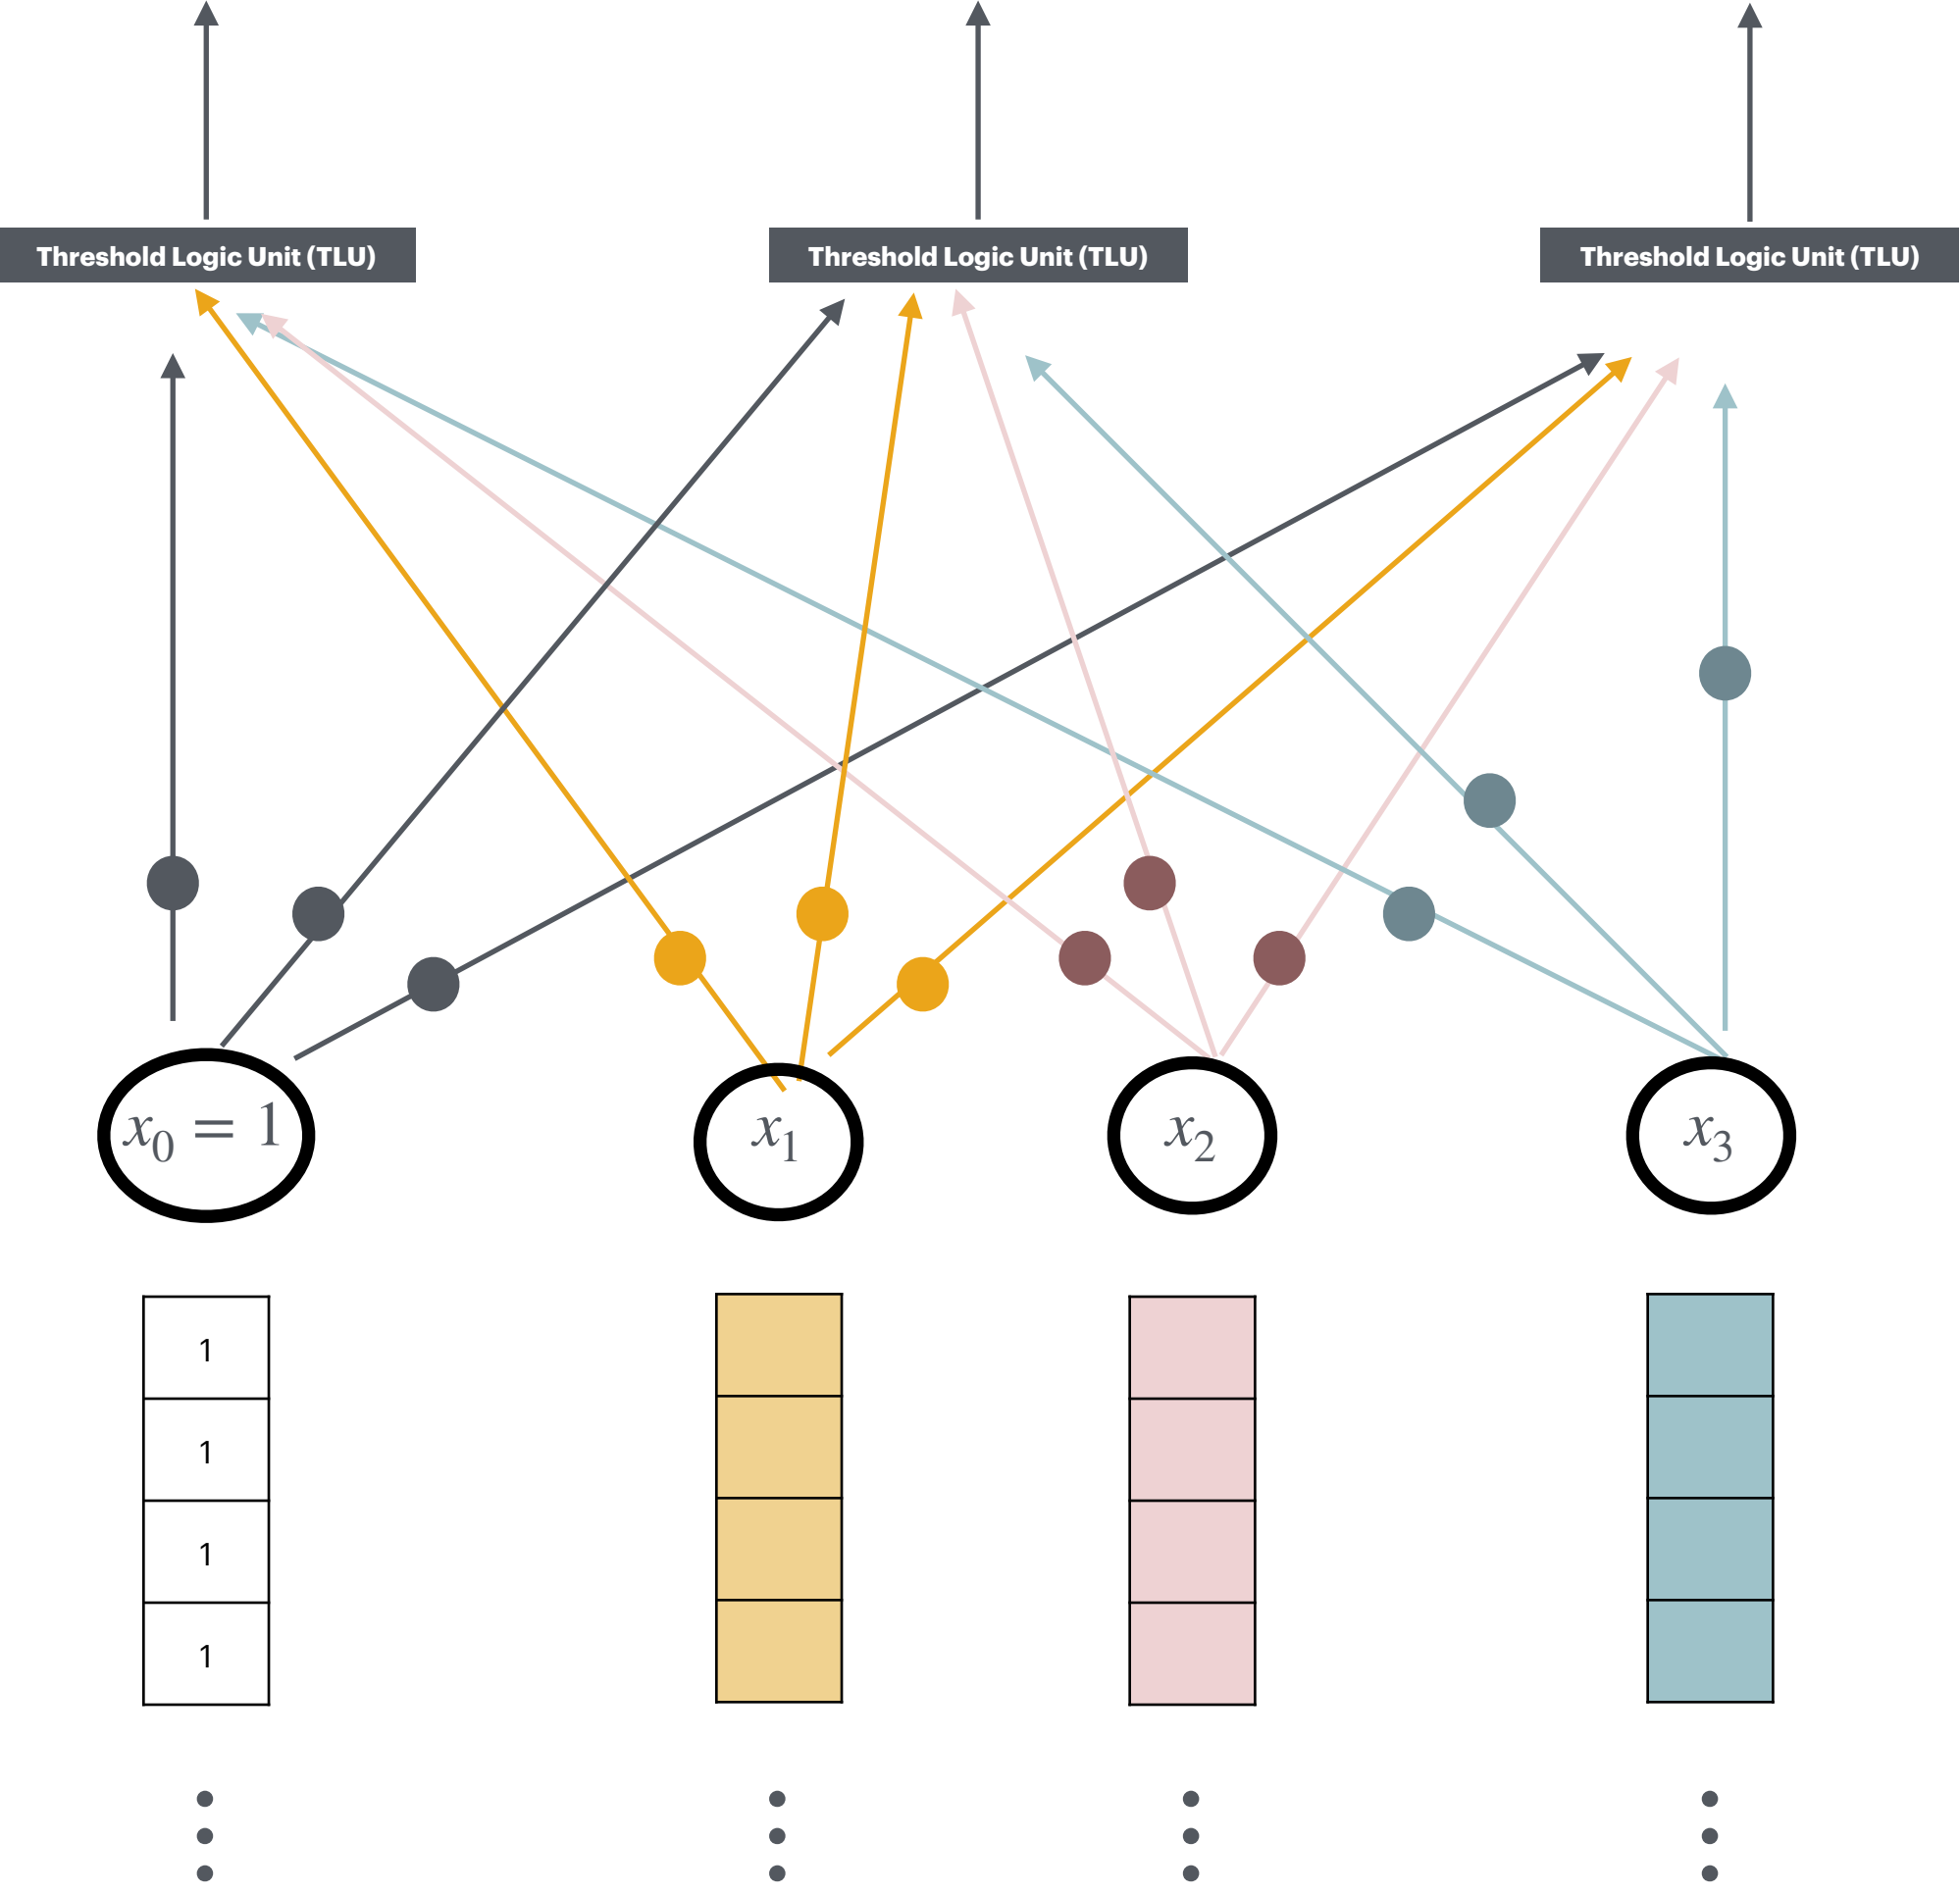
Logical Computation: Perception



Logical Computation: Perception



Logical Computation: Perception


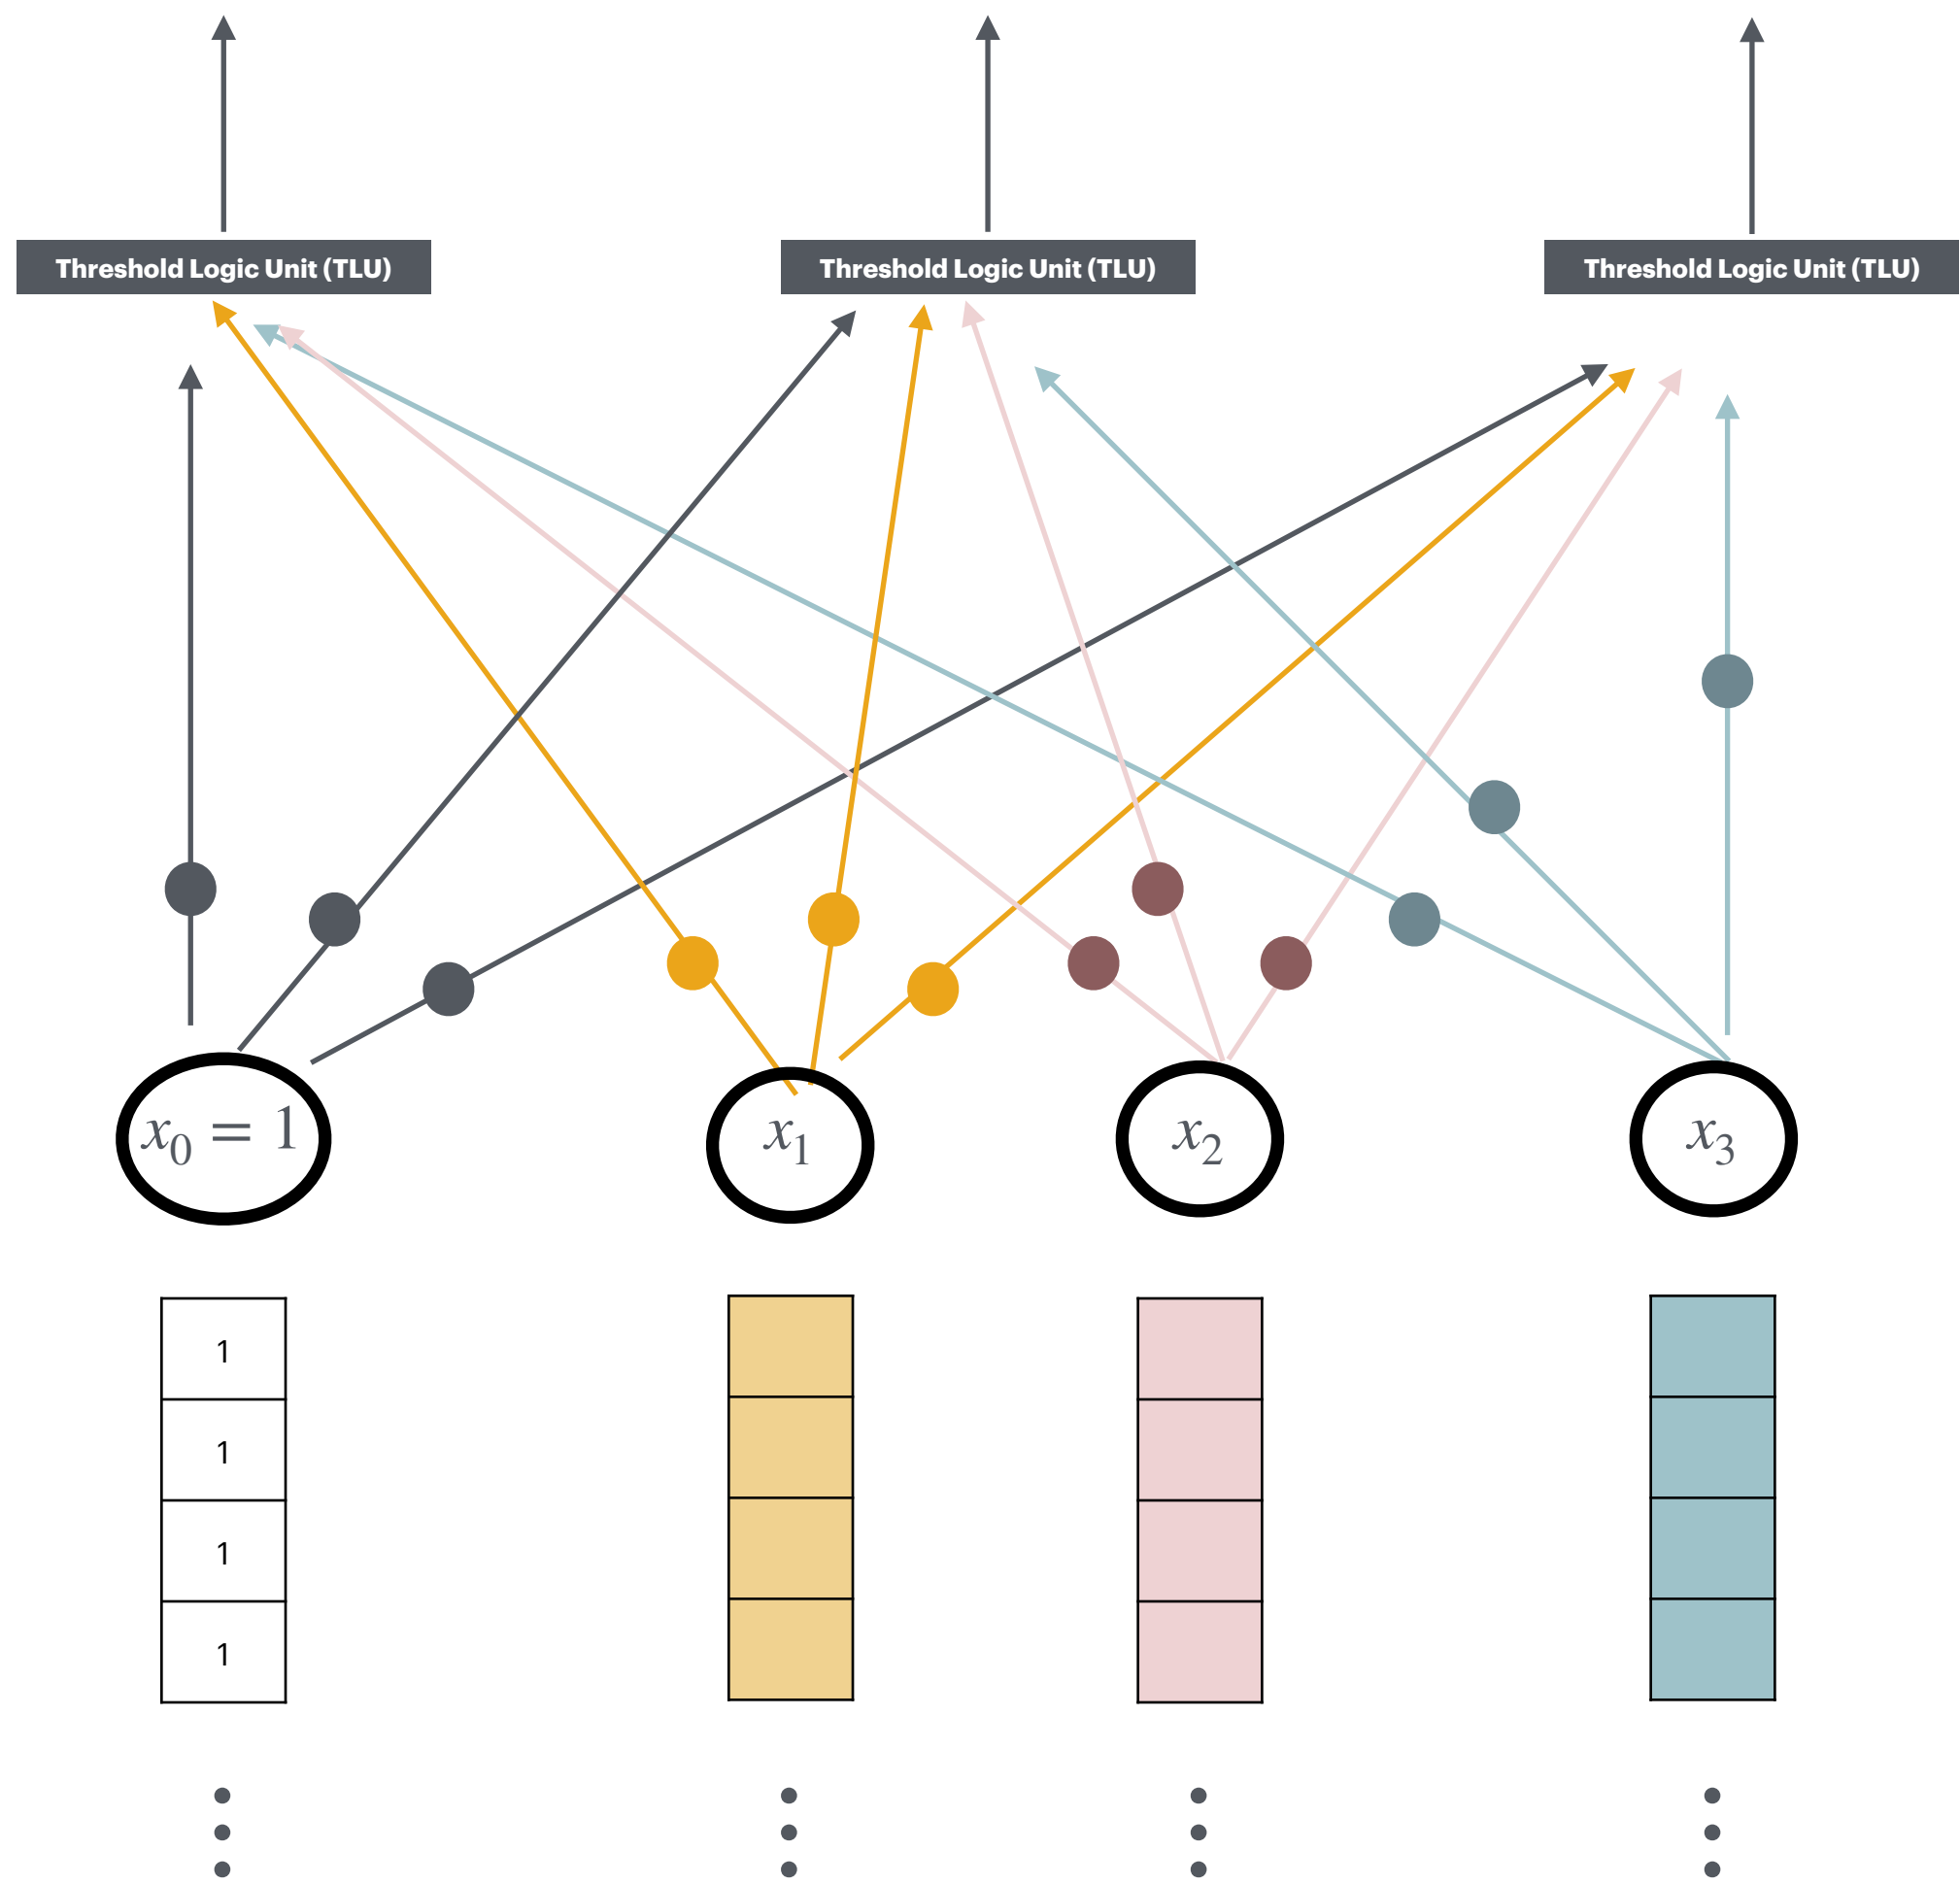


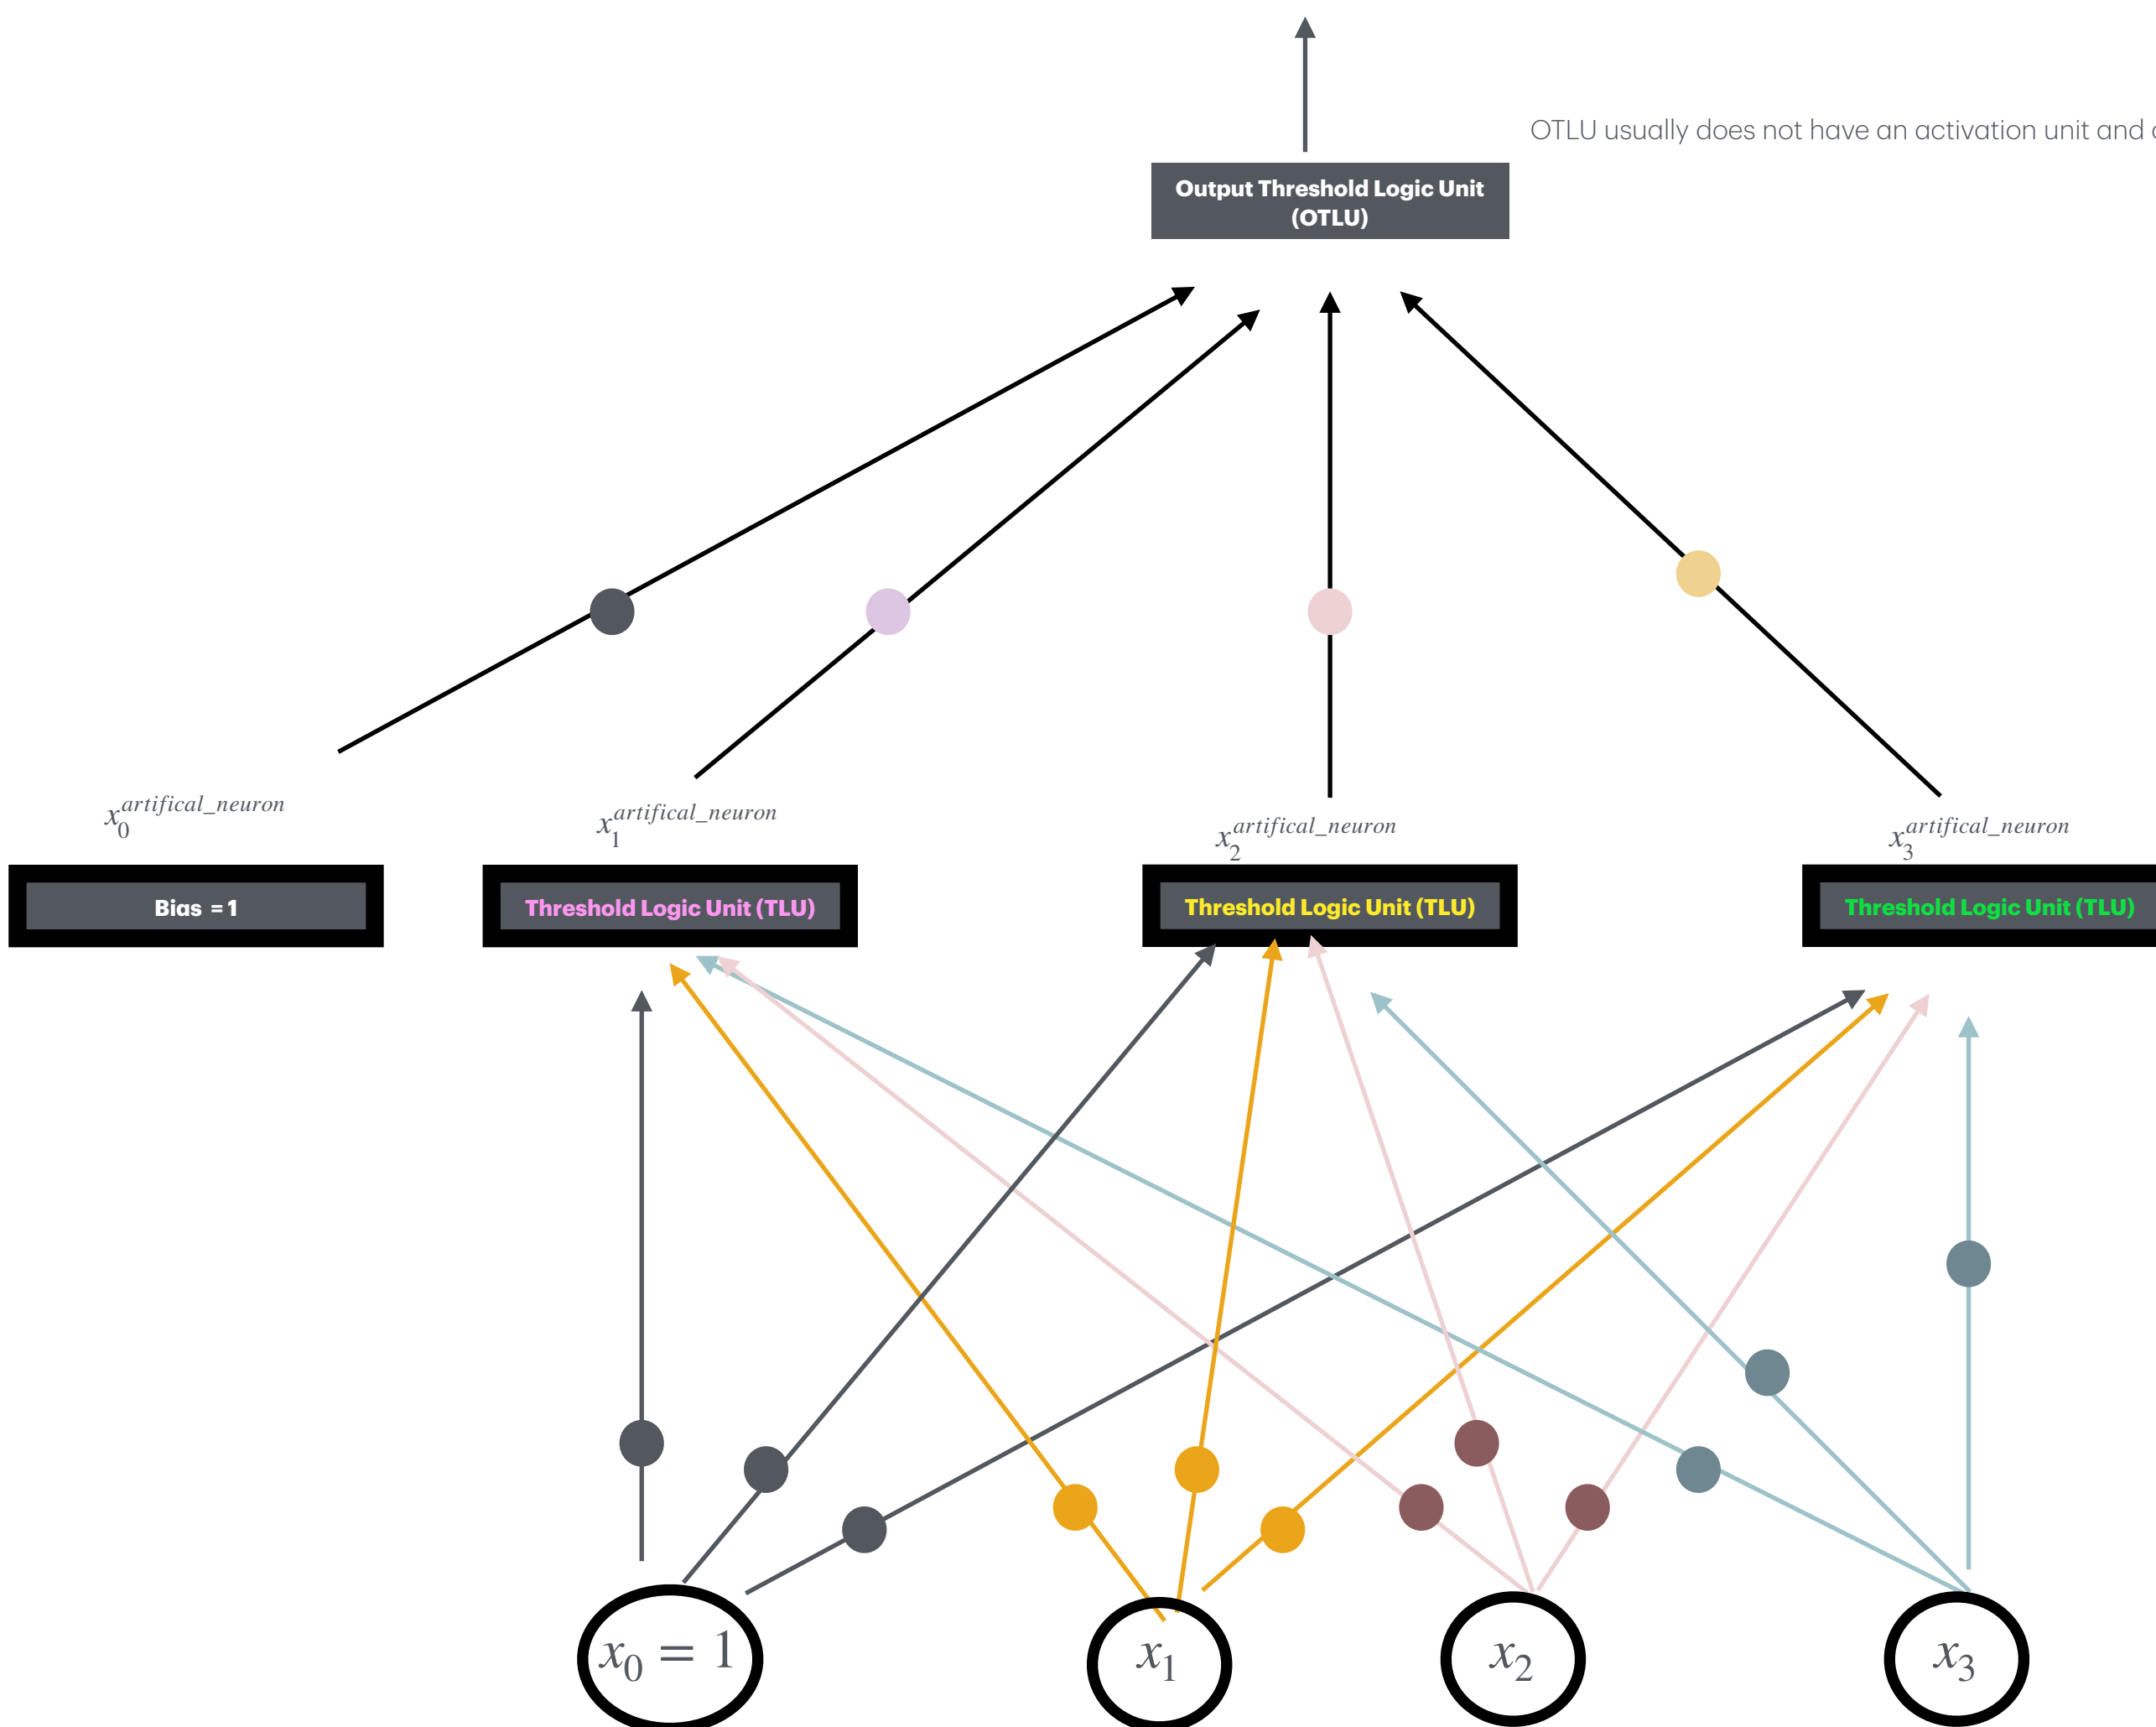
Note: Single Instance

$$\left(\begin{array}{|c|c|c|} \hline \text{yellow} & \text{pink} & \text{teal} \\ \hline \end{array} \begin{array}{|c|c|c|} \hline \text{yellow} & \text{yellow} & \text{yellow} \\ \hline \text{pink} & \text{pink} & \text{pink} \\ \hline \text{teal} & \text{teal} & \text{teal} \\ \hline \end{array} \right) + \begin{array}{|c|c|c|} \hline \text{grey} & \text{grey} & \text{grey} \\ \hline \end{array}$$
$$\left(\begin{array}{|c|c|c|} \hline \text{blue} & \text{blue} & \text{blue} \\ \hline \end{array} \right) + \begin{array}{|c|c|c|} \hline \text{grey} & \text{grey} & \text{grey} \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline \text{grey} & \text{grey} & \text{grey} \\ \hline \end{array}$$

Next layer receives single instances vector

Logical Computation: Perception

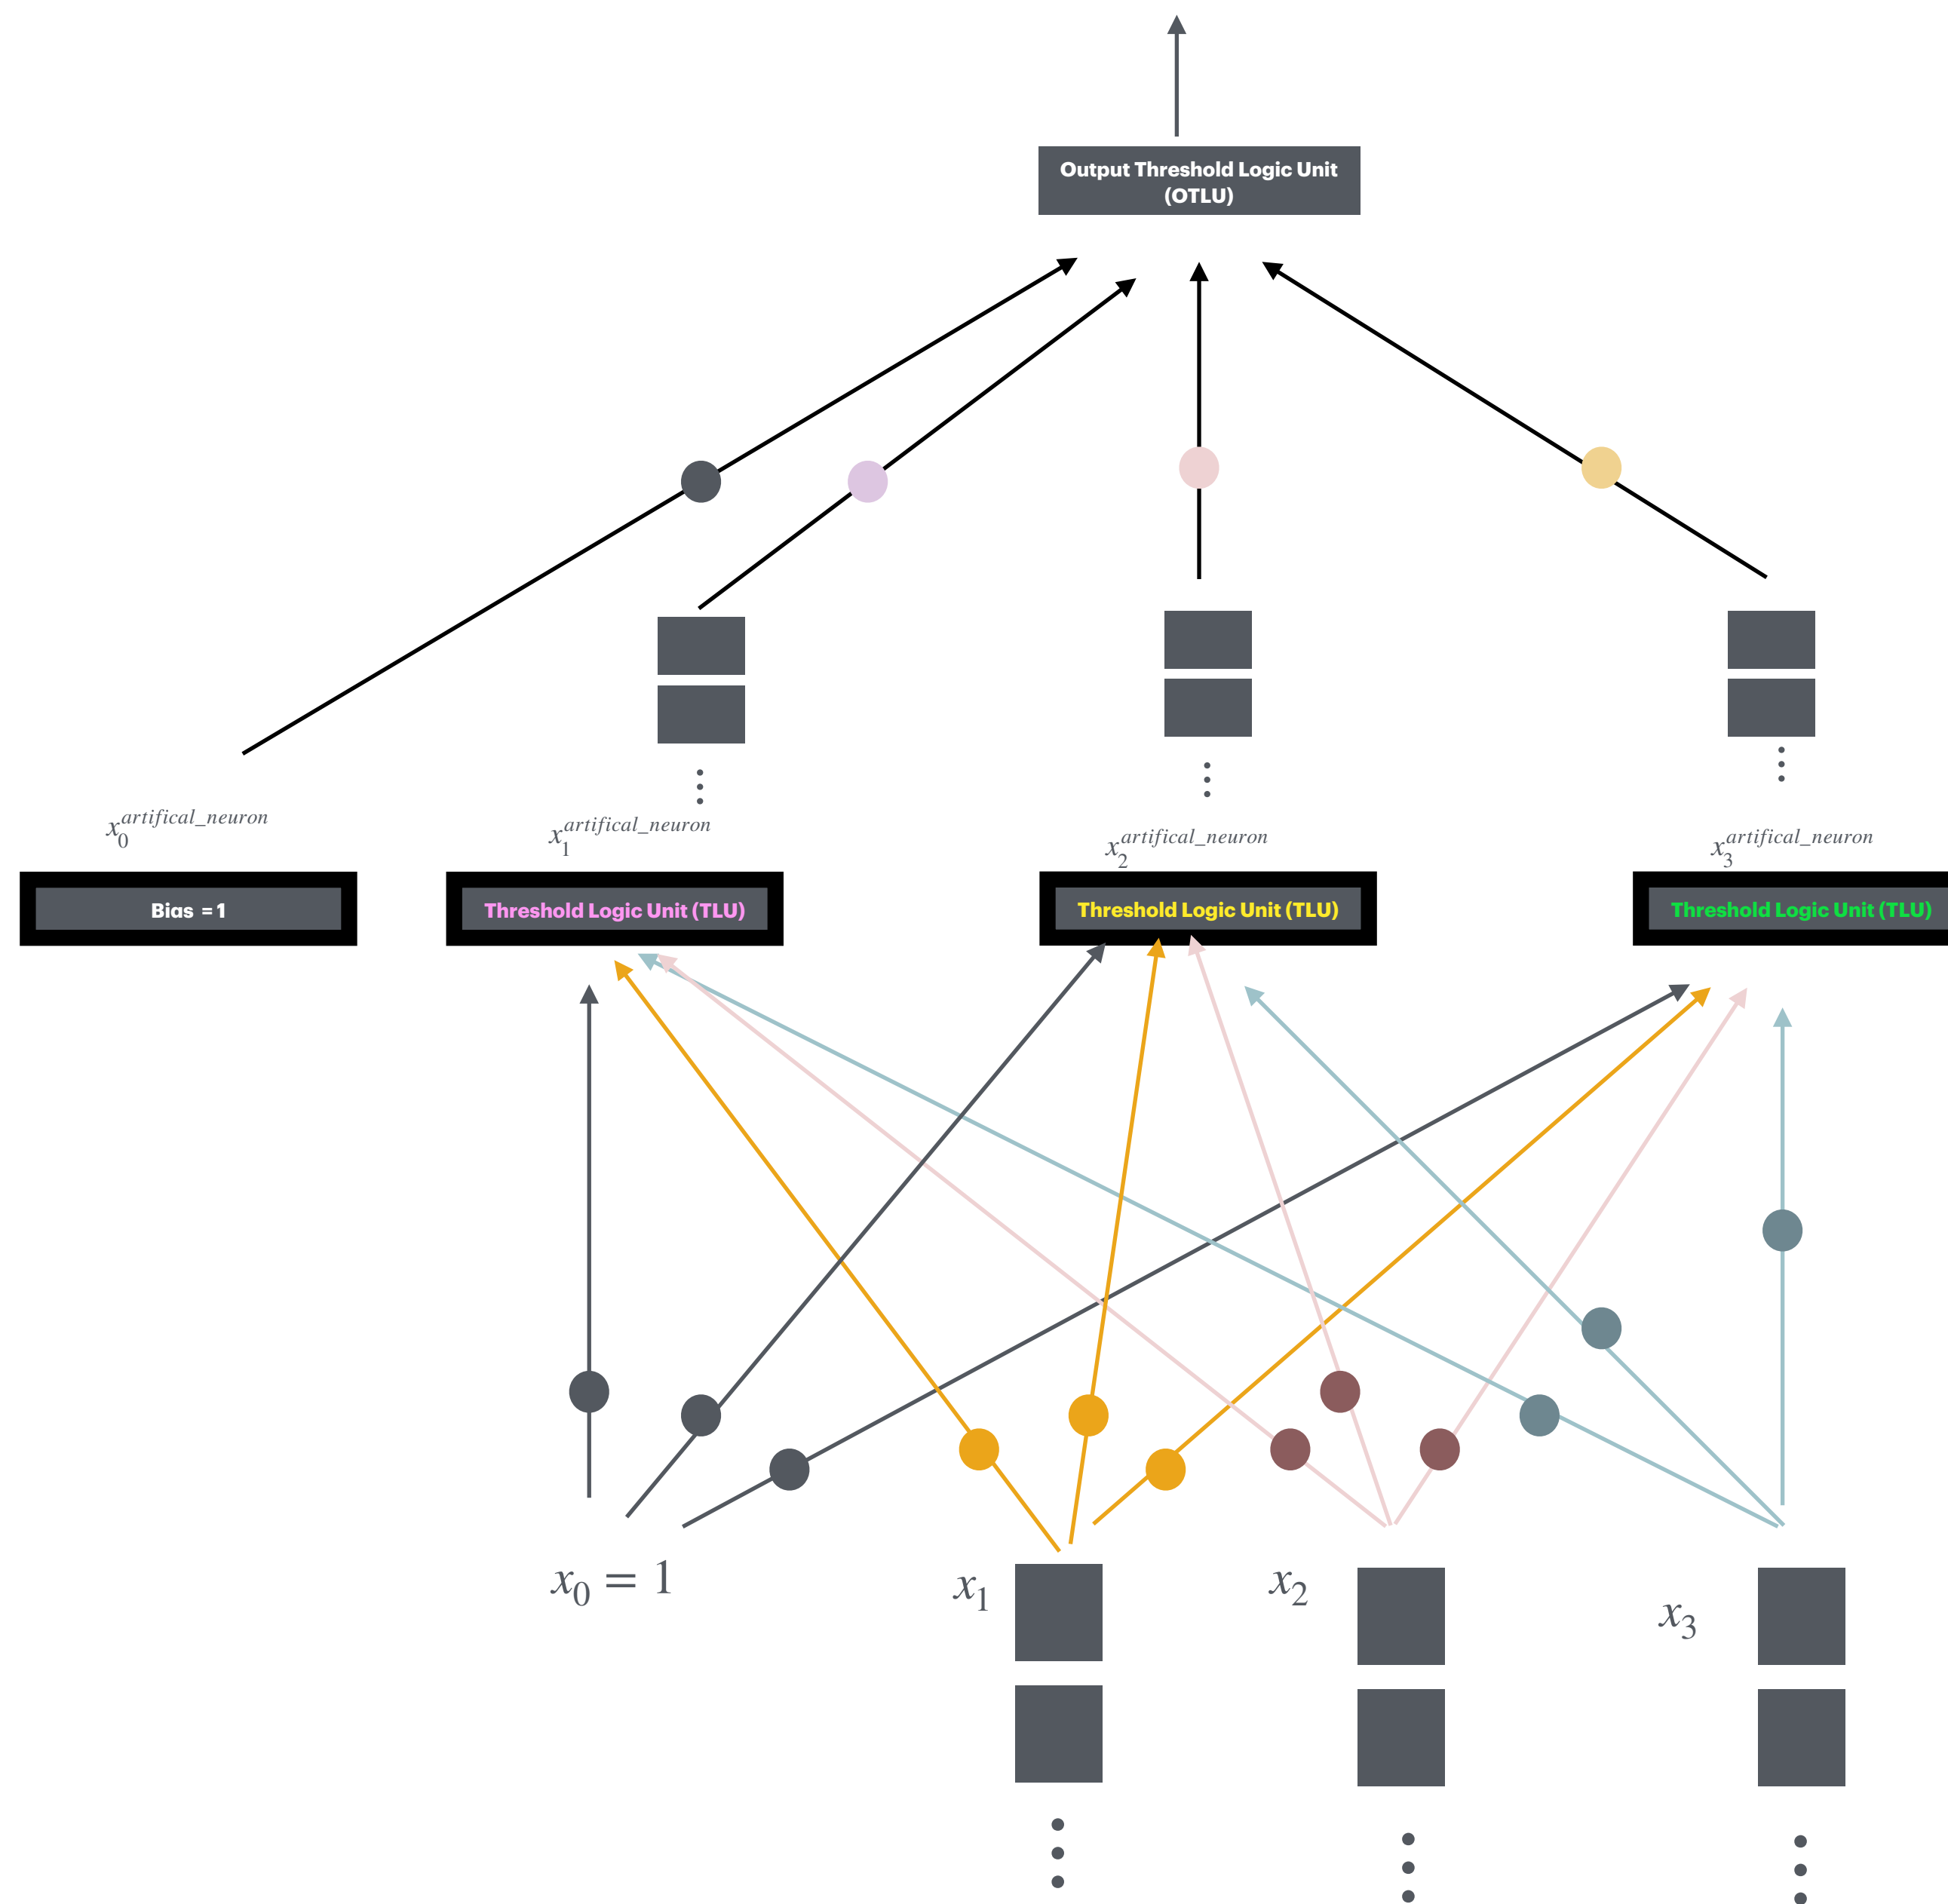




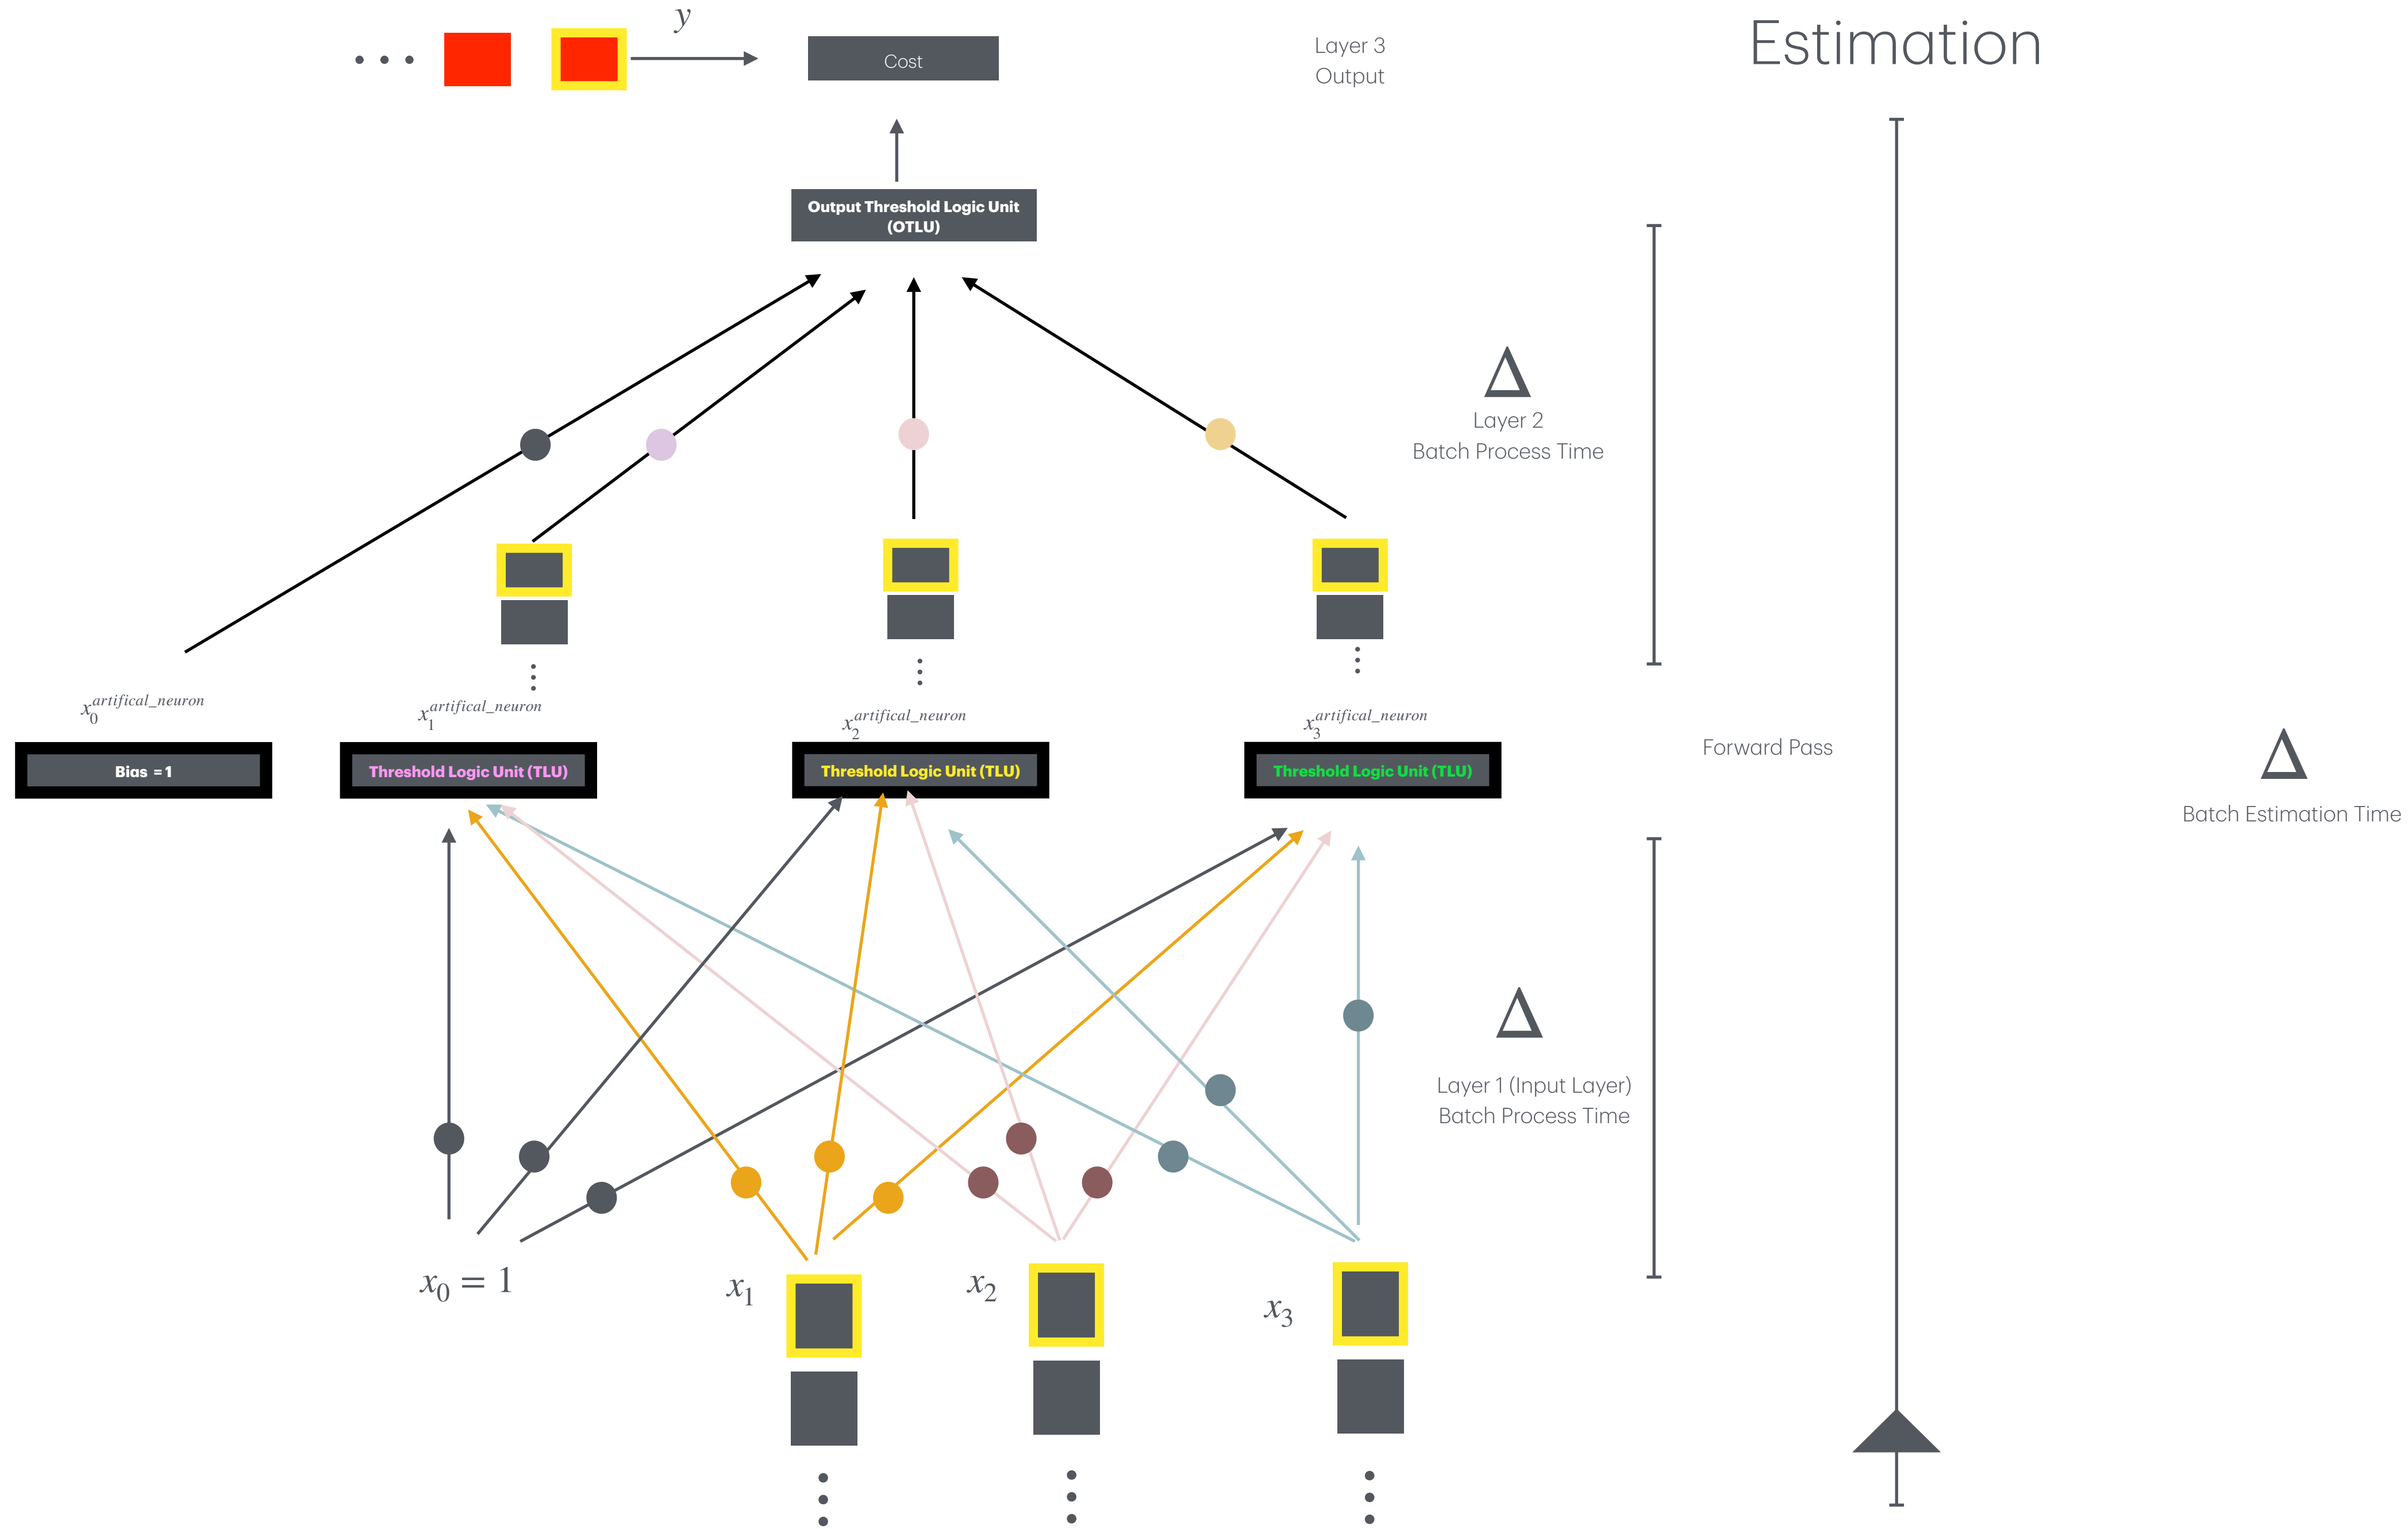
OTLU usually does not have an activation unit and outputs the weighted sum

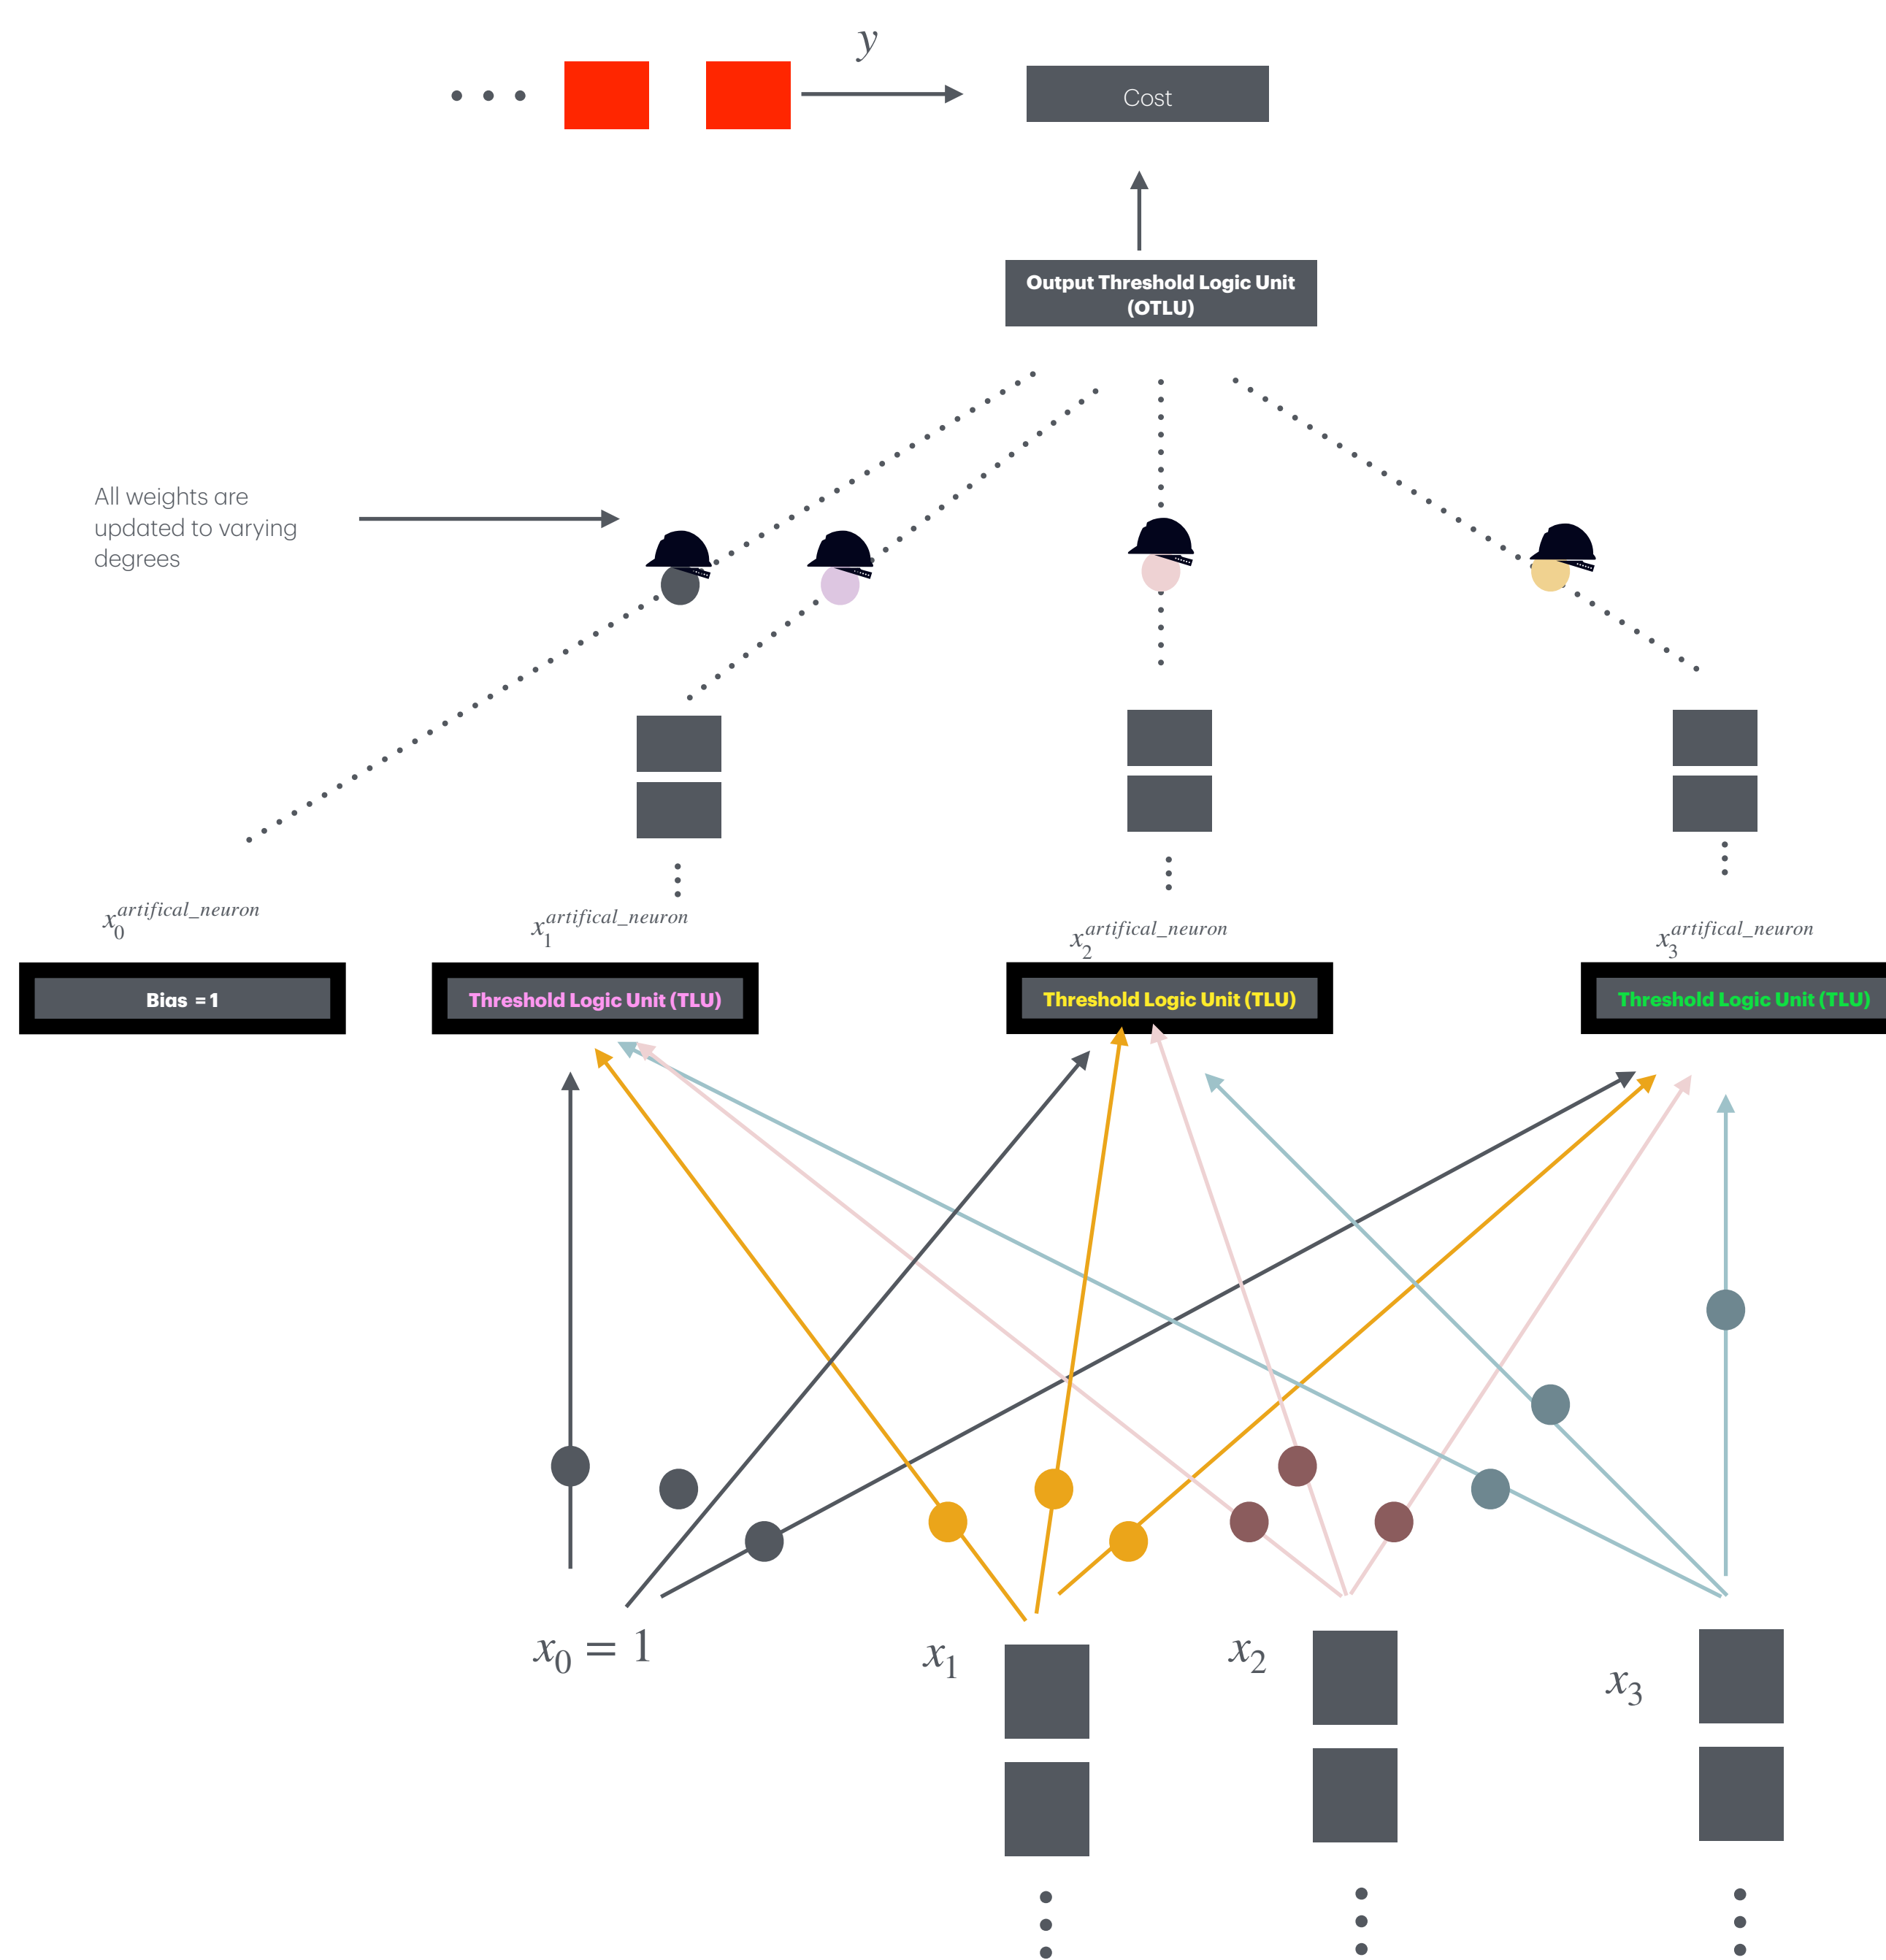
$$\begin{pmatrix} \text{pink} & \text{yellow} & \text{green} \\ \text{purple} & \text{pink} & \text{yellow} \end{pmatrix} + \text{grey} = \text{OTLU}$$

$$\begin{pmatrix} \text{pink} \\ \text{yellow} \\ \text{green} \end{pmatrix} + \text{grey} = \text{OTLU}$$



Networks handles mini-batch at a time





Reverse Pass

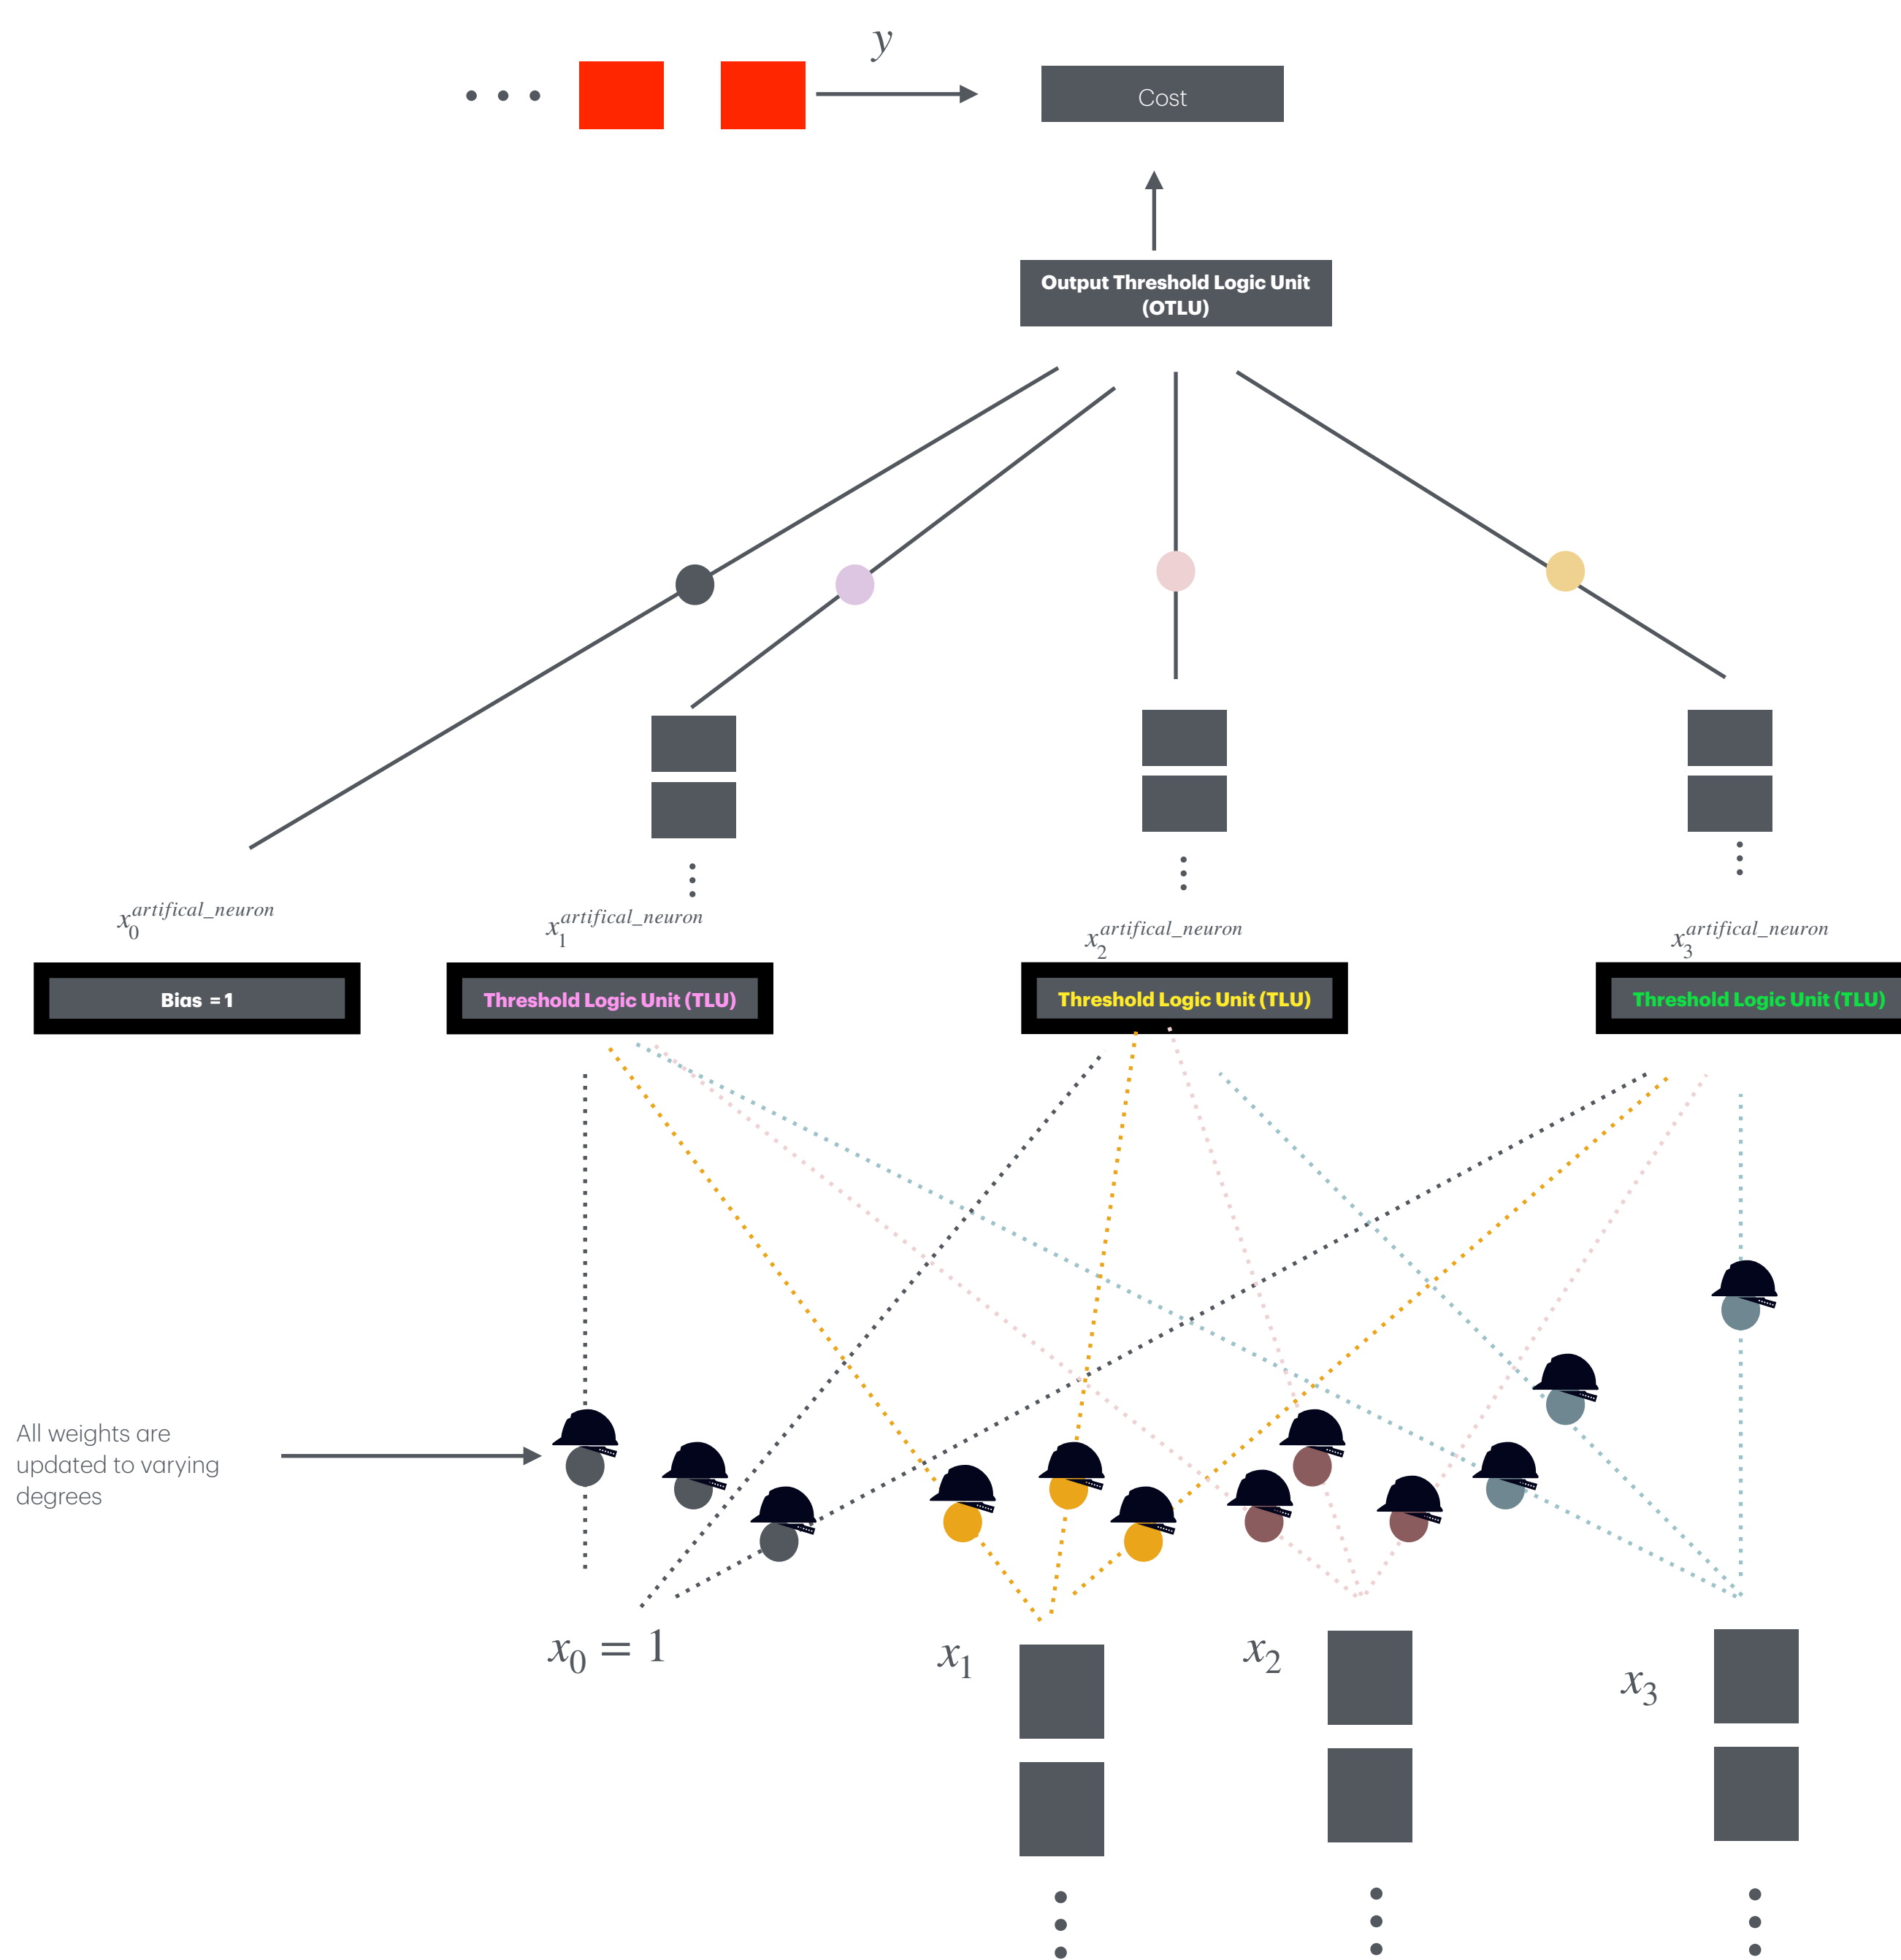
How much **layer 2 connections** contributed to high cost (i.e. high error)

Cost/Error gradients are measured across connections (weights)

$$\frac{Cost}{w_{some_connection}}$$

Gradient Descent performed on all connections (weights) using error gradients

Note: Input batch persistence is required for reverse algorithm



Reverse Pass

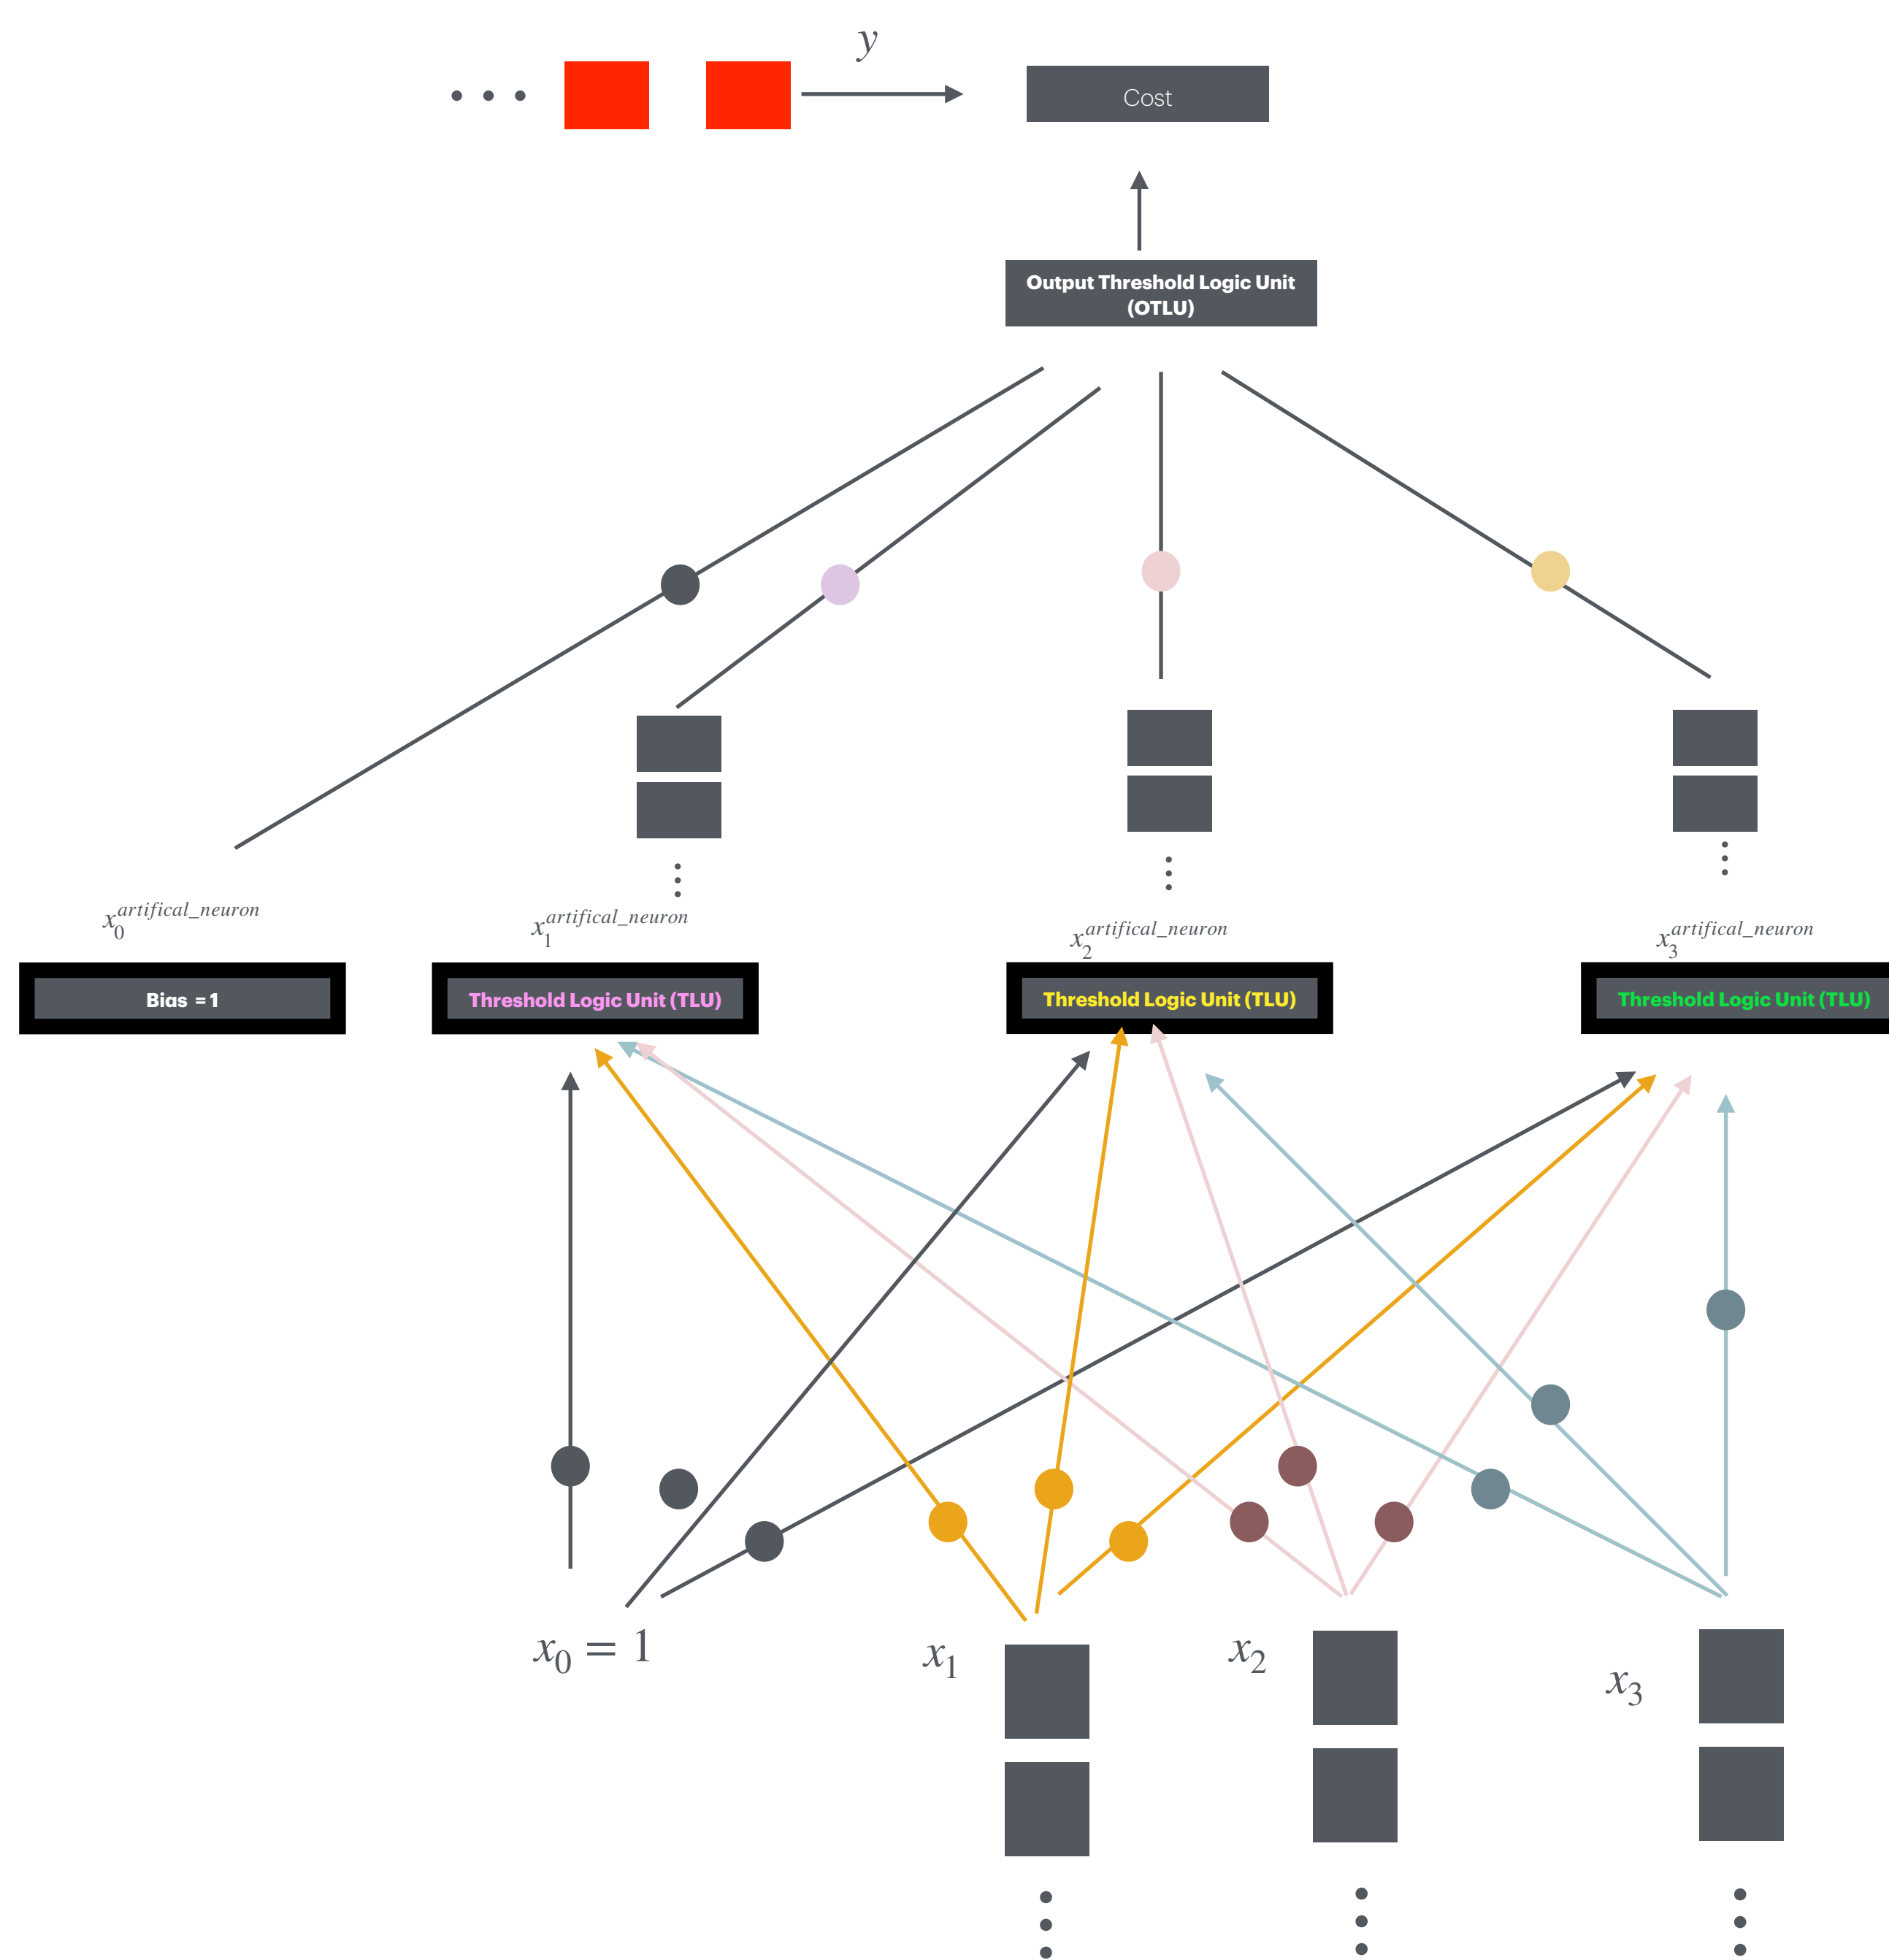
How much **layer 1 connections** contributed to high cost (i.e. high error)

Cost/Error gradients are measured across connections (weights)

Gradient Descent performed on all connections (weights) using error gradients

Note: Input batch persistence is required for reverse algorithm

$$\frac{Cost}{W_{some_connection}}$$



Reverse Pass

Stop after performing weight update using Gradient Descent on all the connections in each layer

: Activation Function



Activation Function:

Linear Regression/Classifiers :

- **Heaviside**
- **Sign Function**

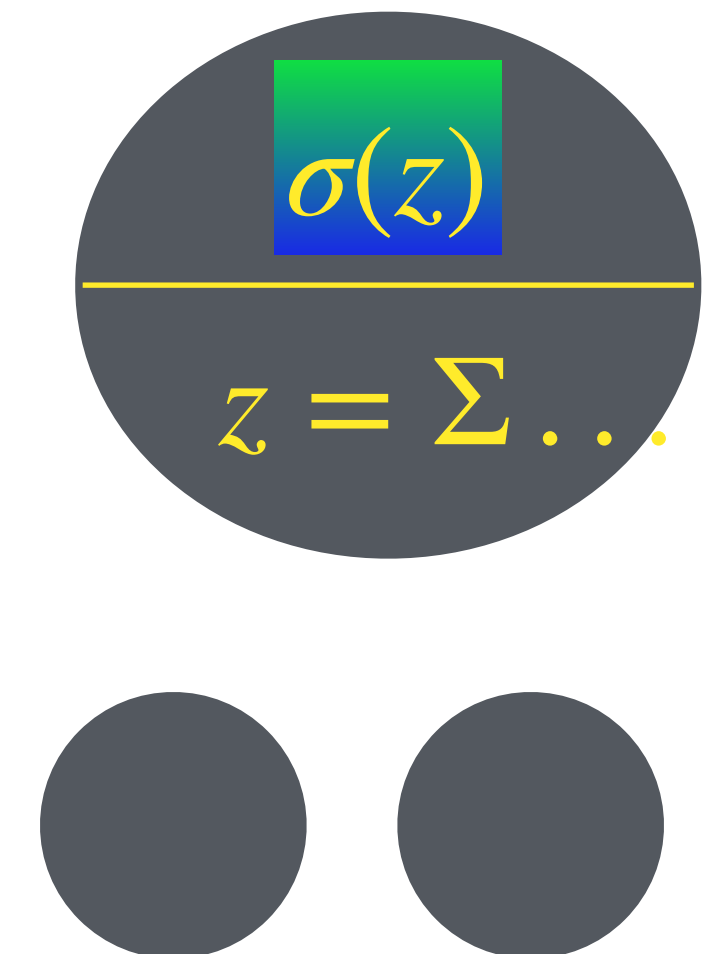
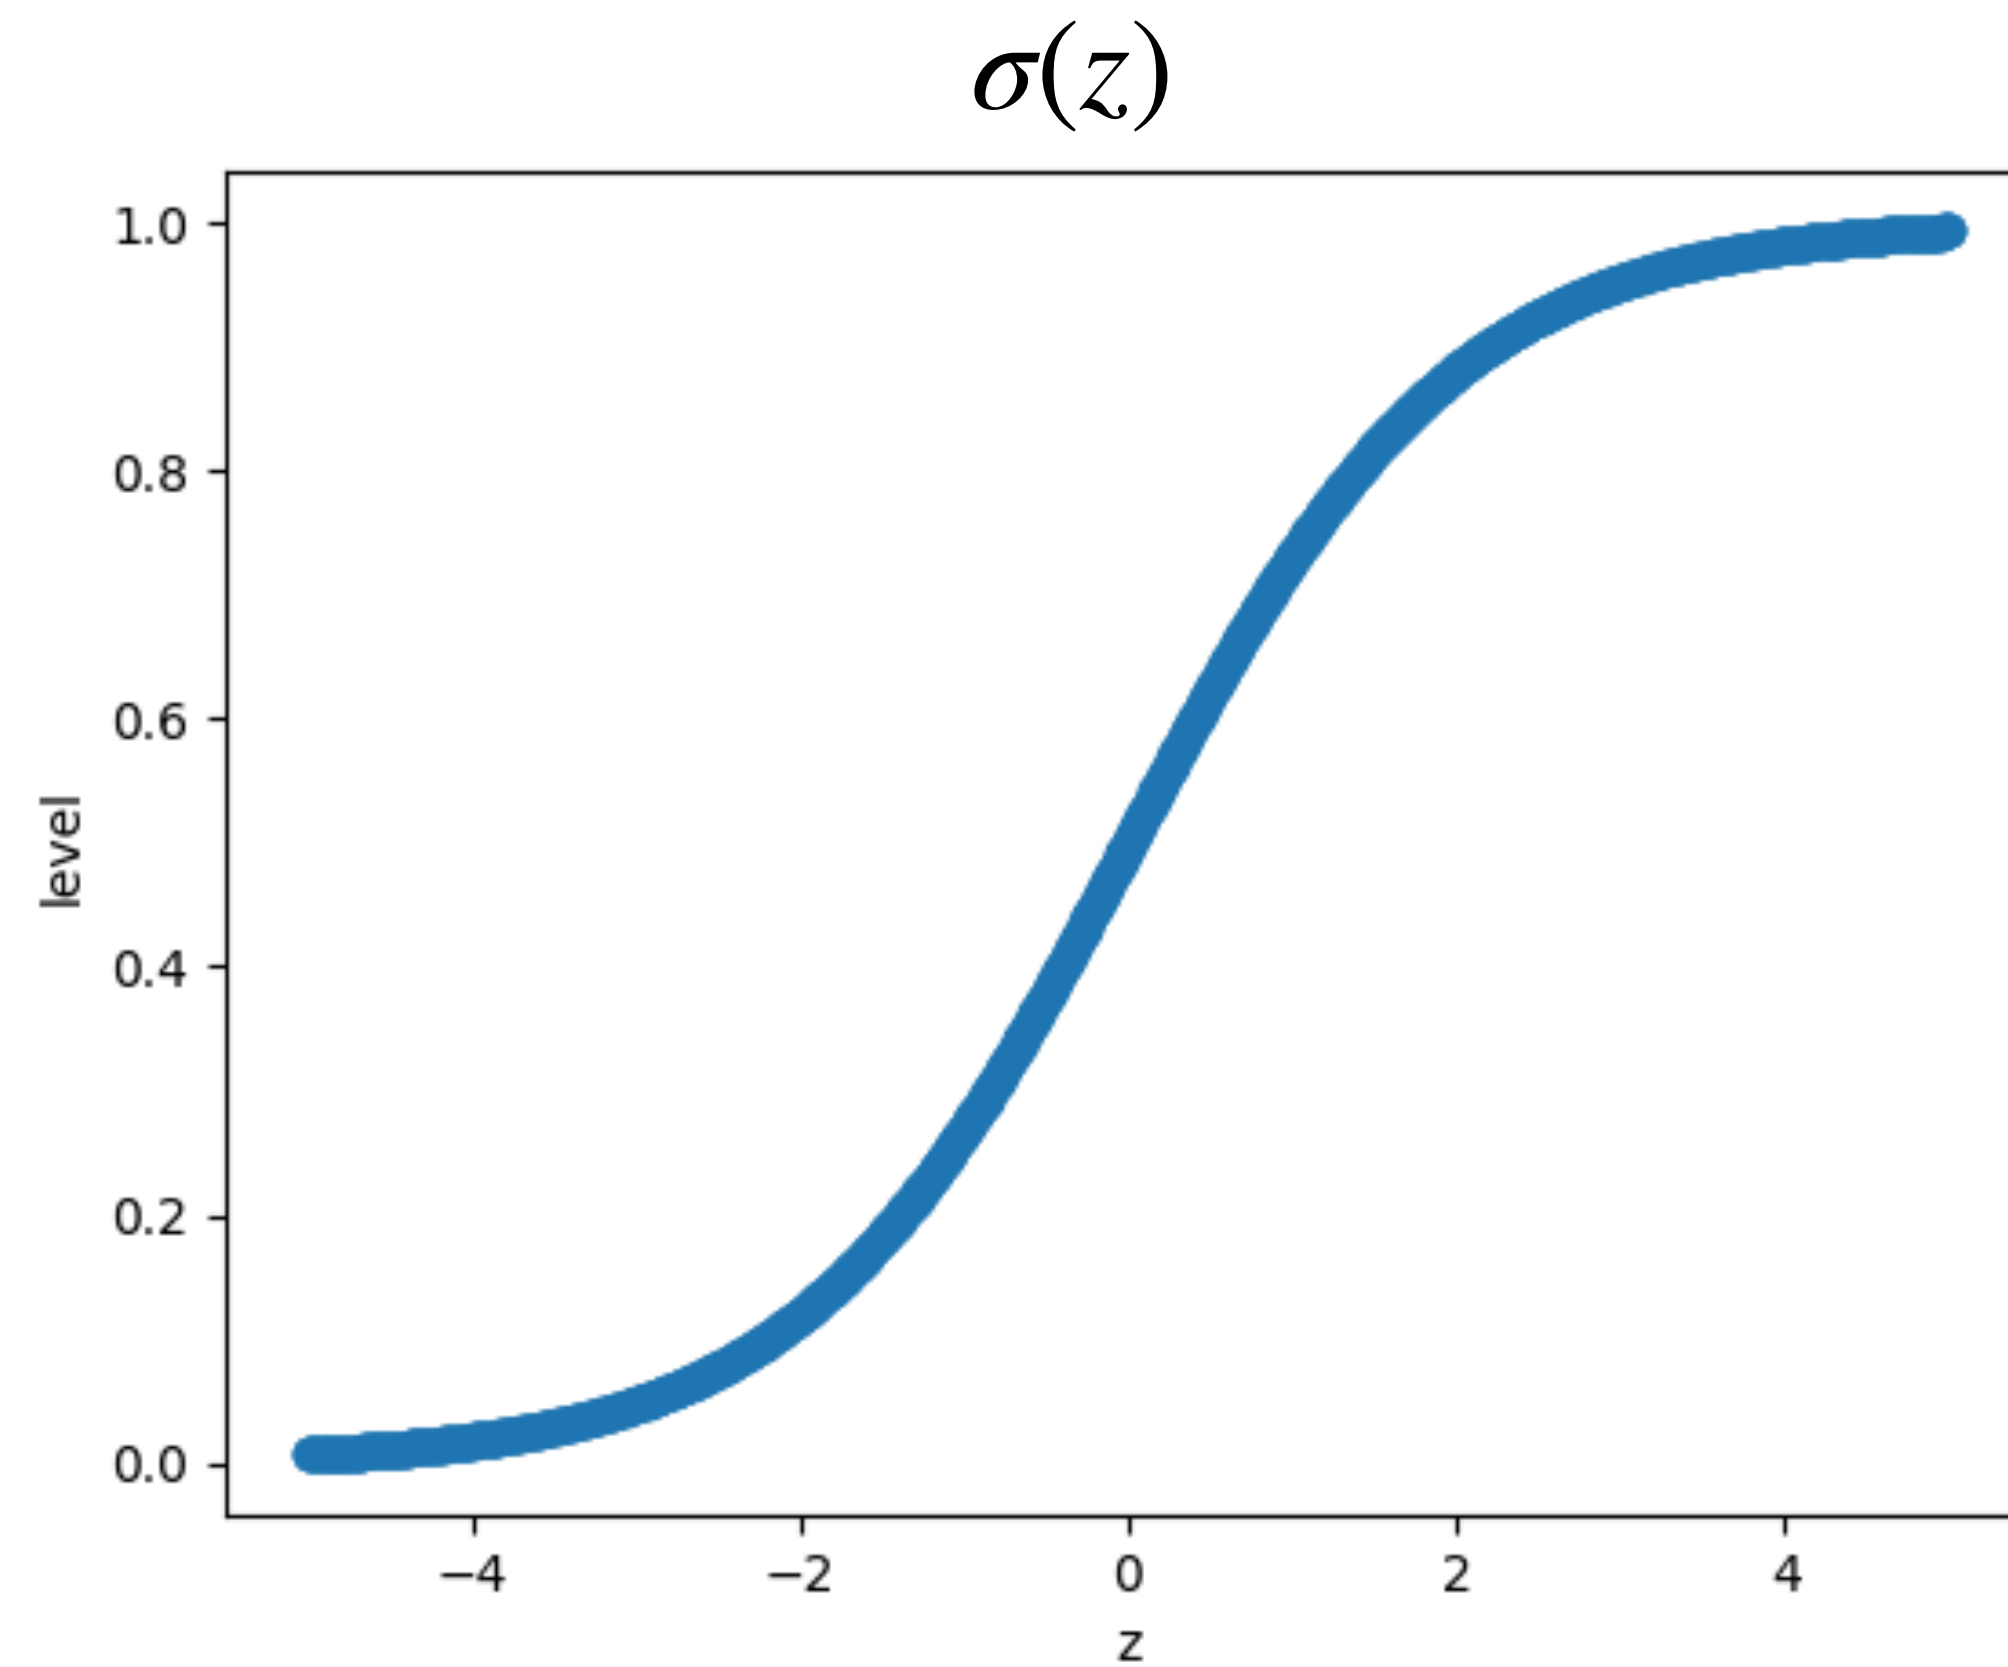
Nonlinear Regression/Classifiers :

- **Sigmoid Function**
- **Hyperbolic Tangent Function**
- **Rectified Linear Unit Function**

Non-Linear activation functions can be used on linearly models as well


Activation Function

Biological neurons have been observed to implement a roughly **sigmoid activation** function



Forward-Mode AutoDifferentiation

$g(x, y) = 5 + xy$

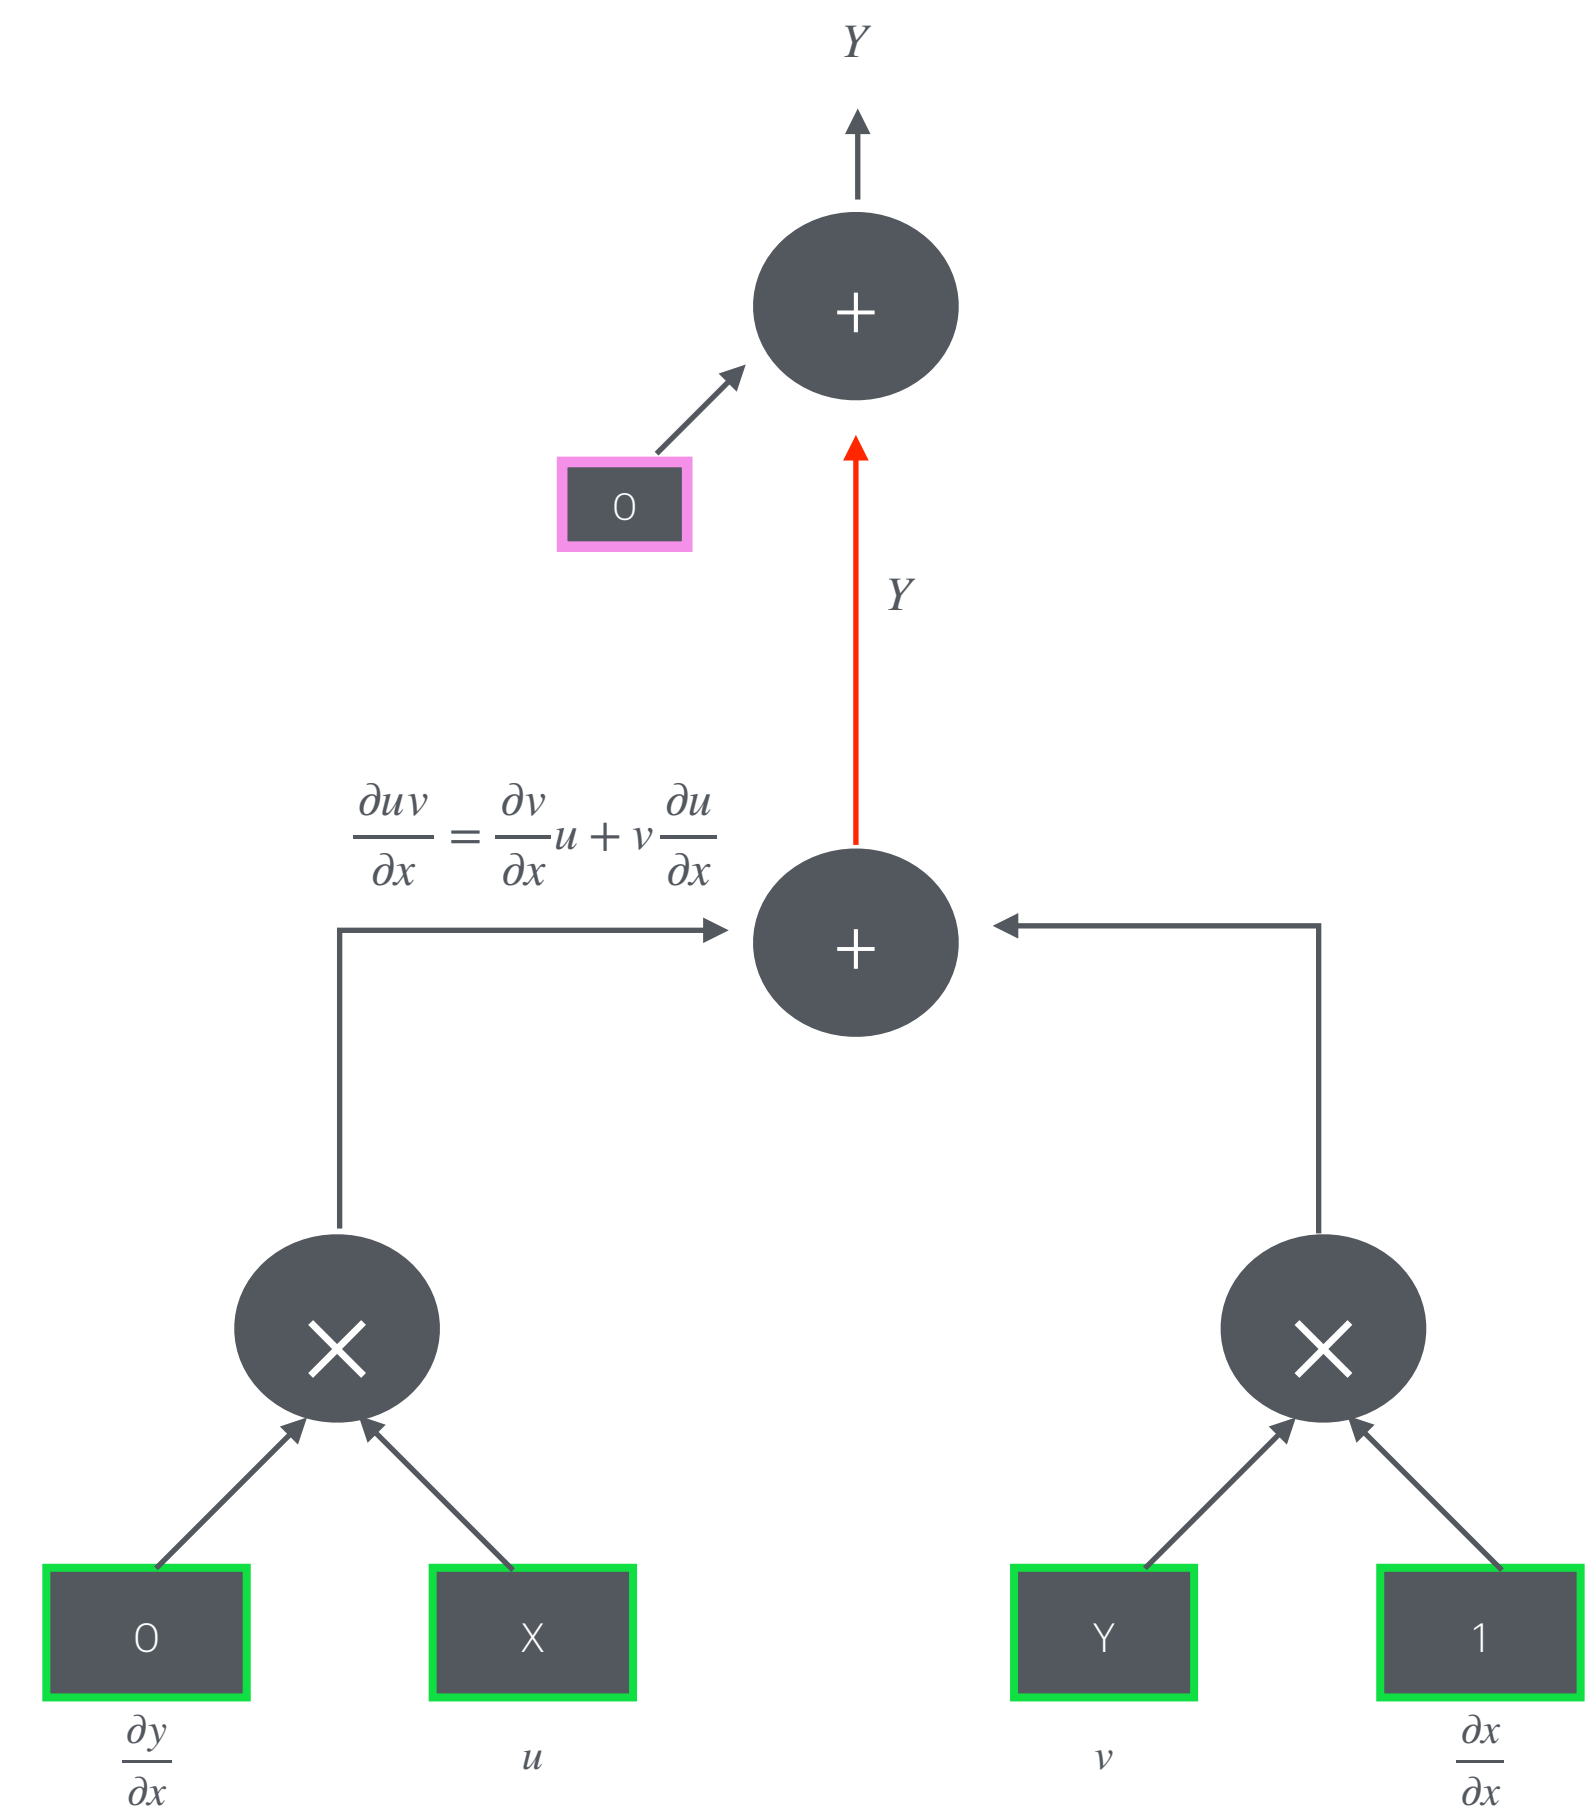


Forward-Mode AutoDifferentiation


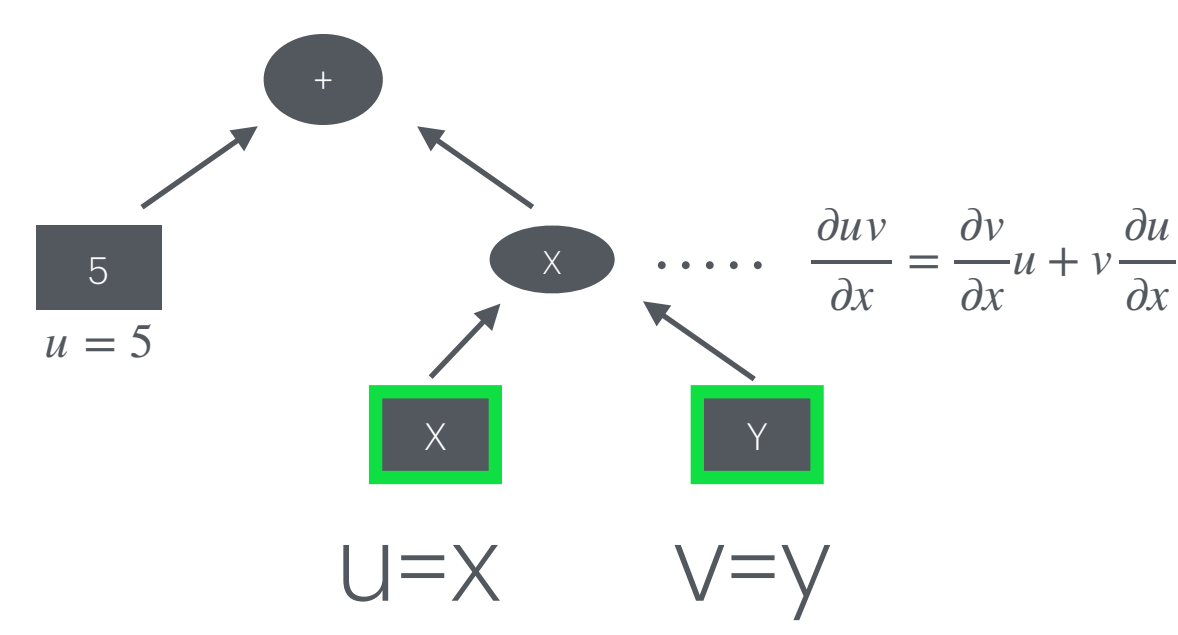
$g(x, y) = 5 + xy$

Partial Derivative $\frac{g(x, y)}{\partial x}$

Symbolic Differentiation (created from AutoDiff)

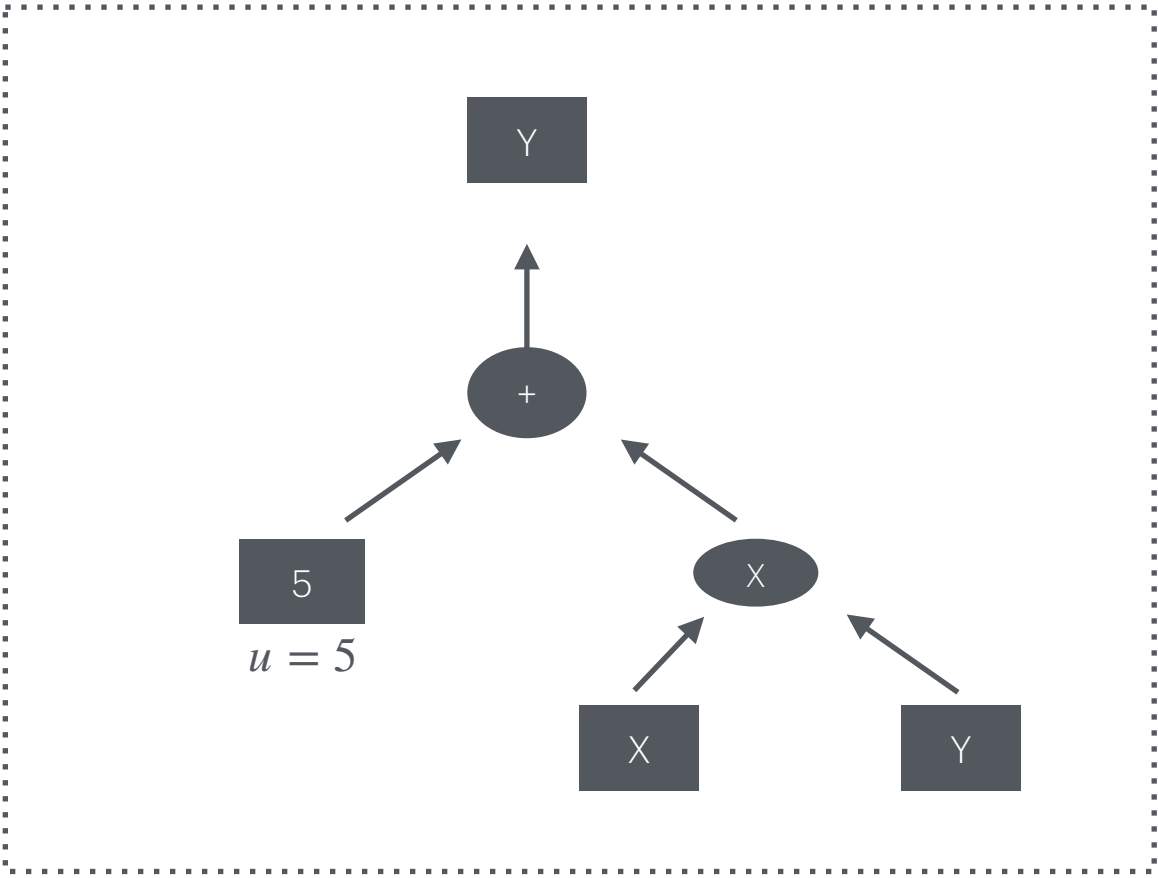
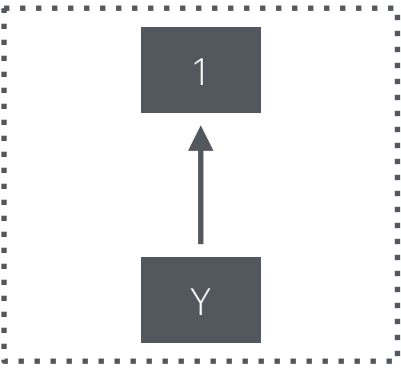
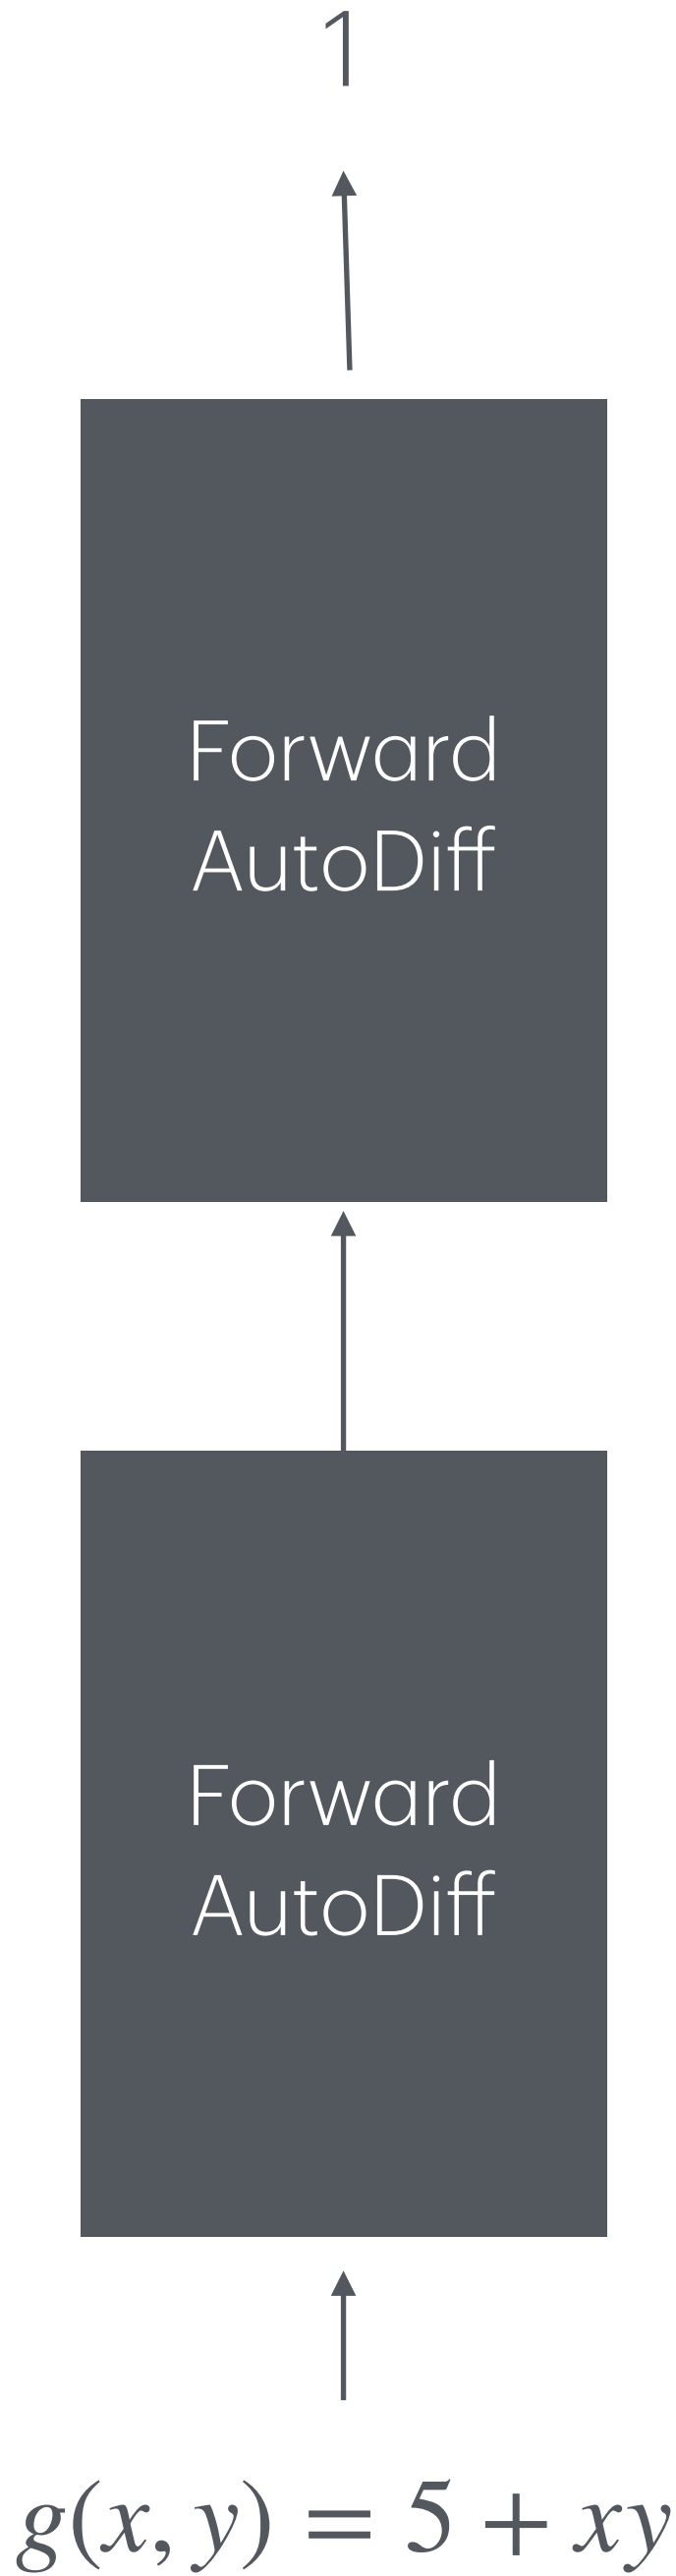


AutoDiff Computation Graphs



Forward-Mode AutoDifferentiation

Chain Forward Mode AutoDiffs



$$\frac{\partial f(3,4)}{\partial x} = ?$$

$$f(3 + \epsilon, 4) = f(3, 4) + f'(3, 4)\epsilon$$

$$f(3 + \epsilon, 4) = f(3, 4) + \frac{\partial f(3, 4)}{\partial x}\epsilon$$

Run forward diff to find real and ϵ components

If ϵ is a infinitesimal number with $\epsilon^2 = 0$, dual numbers can be used to solve forward-mode autodiff

Rule: $h(a + \epsilon) = h(a) + h'(a)\epsilon$

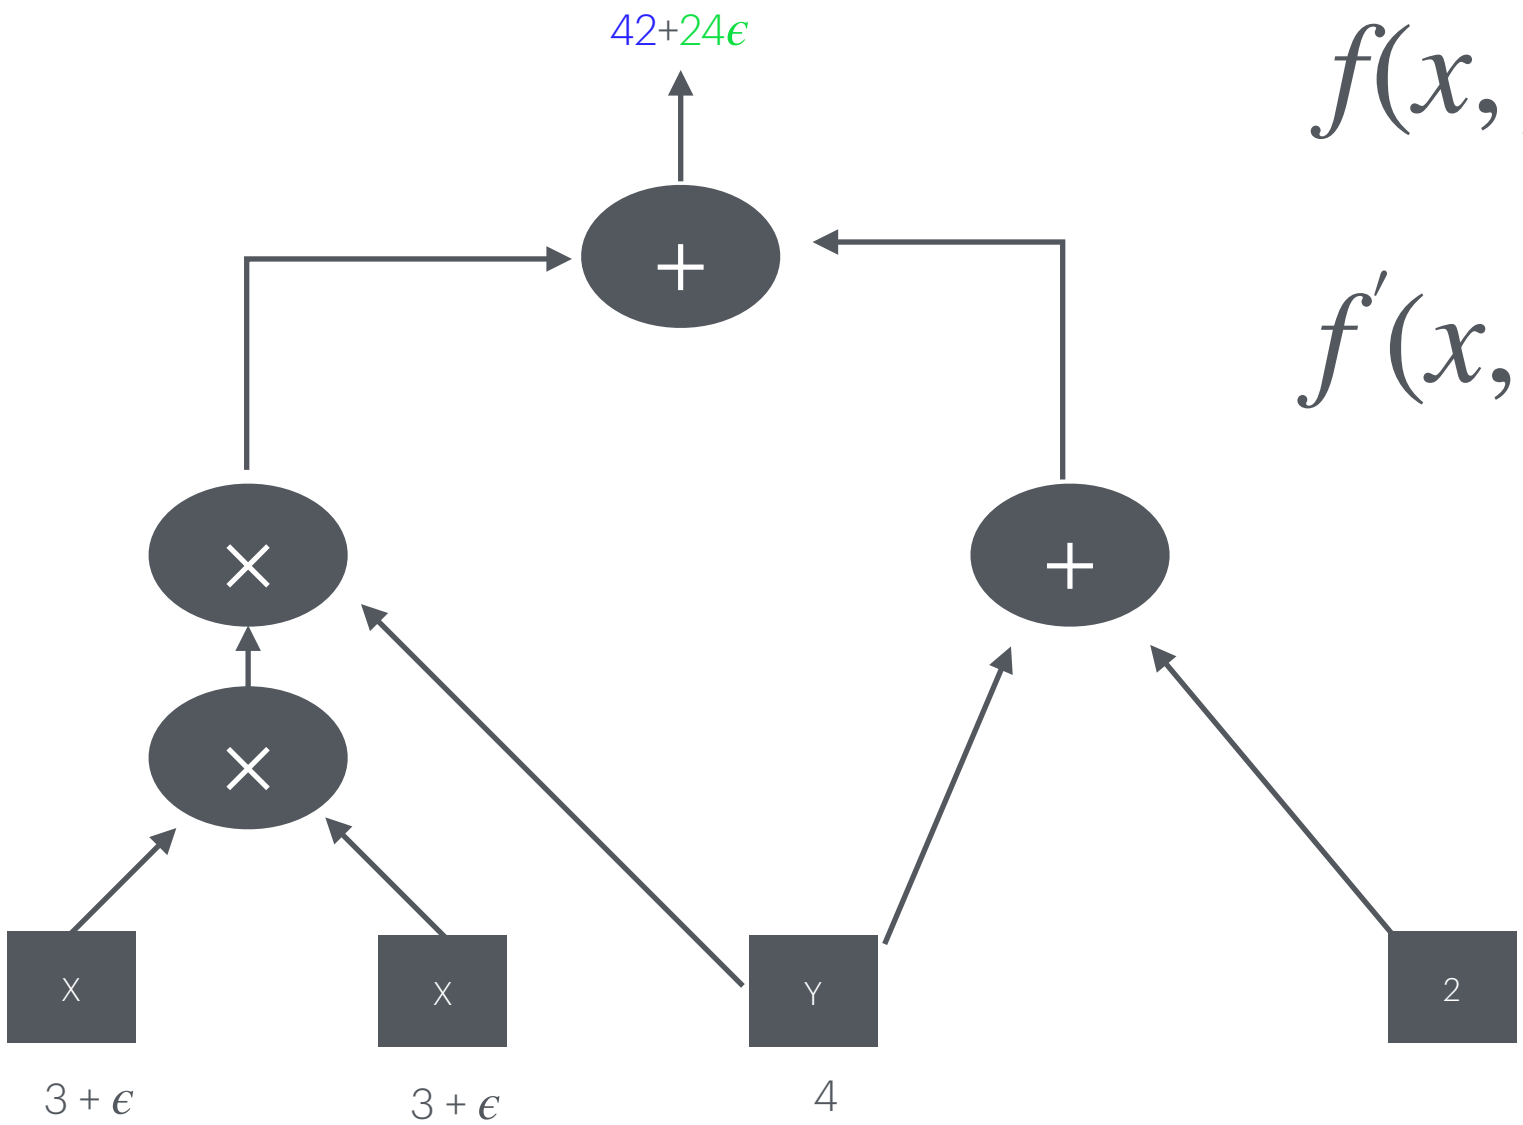


$$f(x, y) = x^2y + y + 2$$

$f(x, y) = 42$ real component

$f'(x, y) = 24$ ϵ component

Note: $\epsilon^2 = 0$

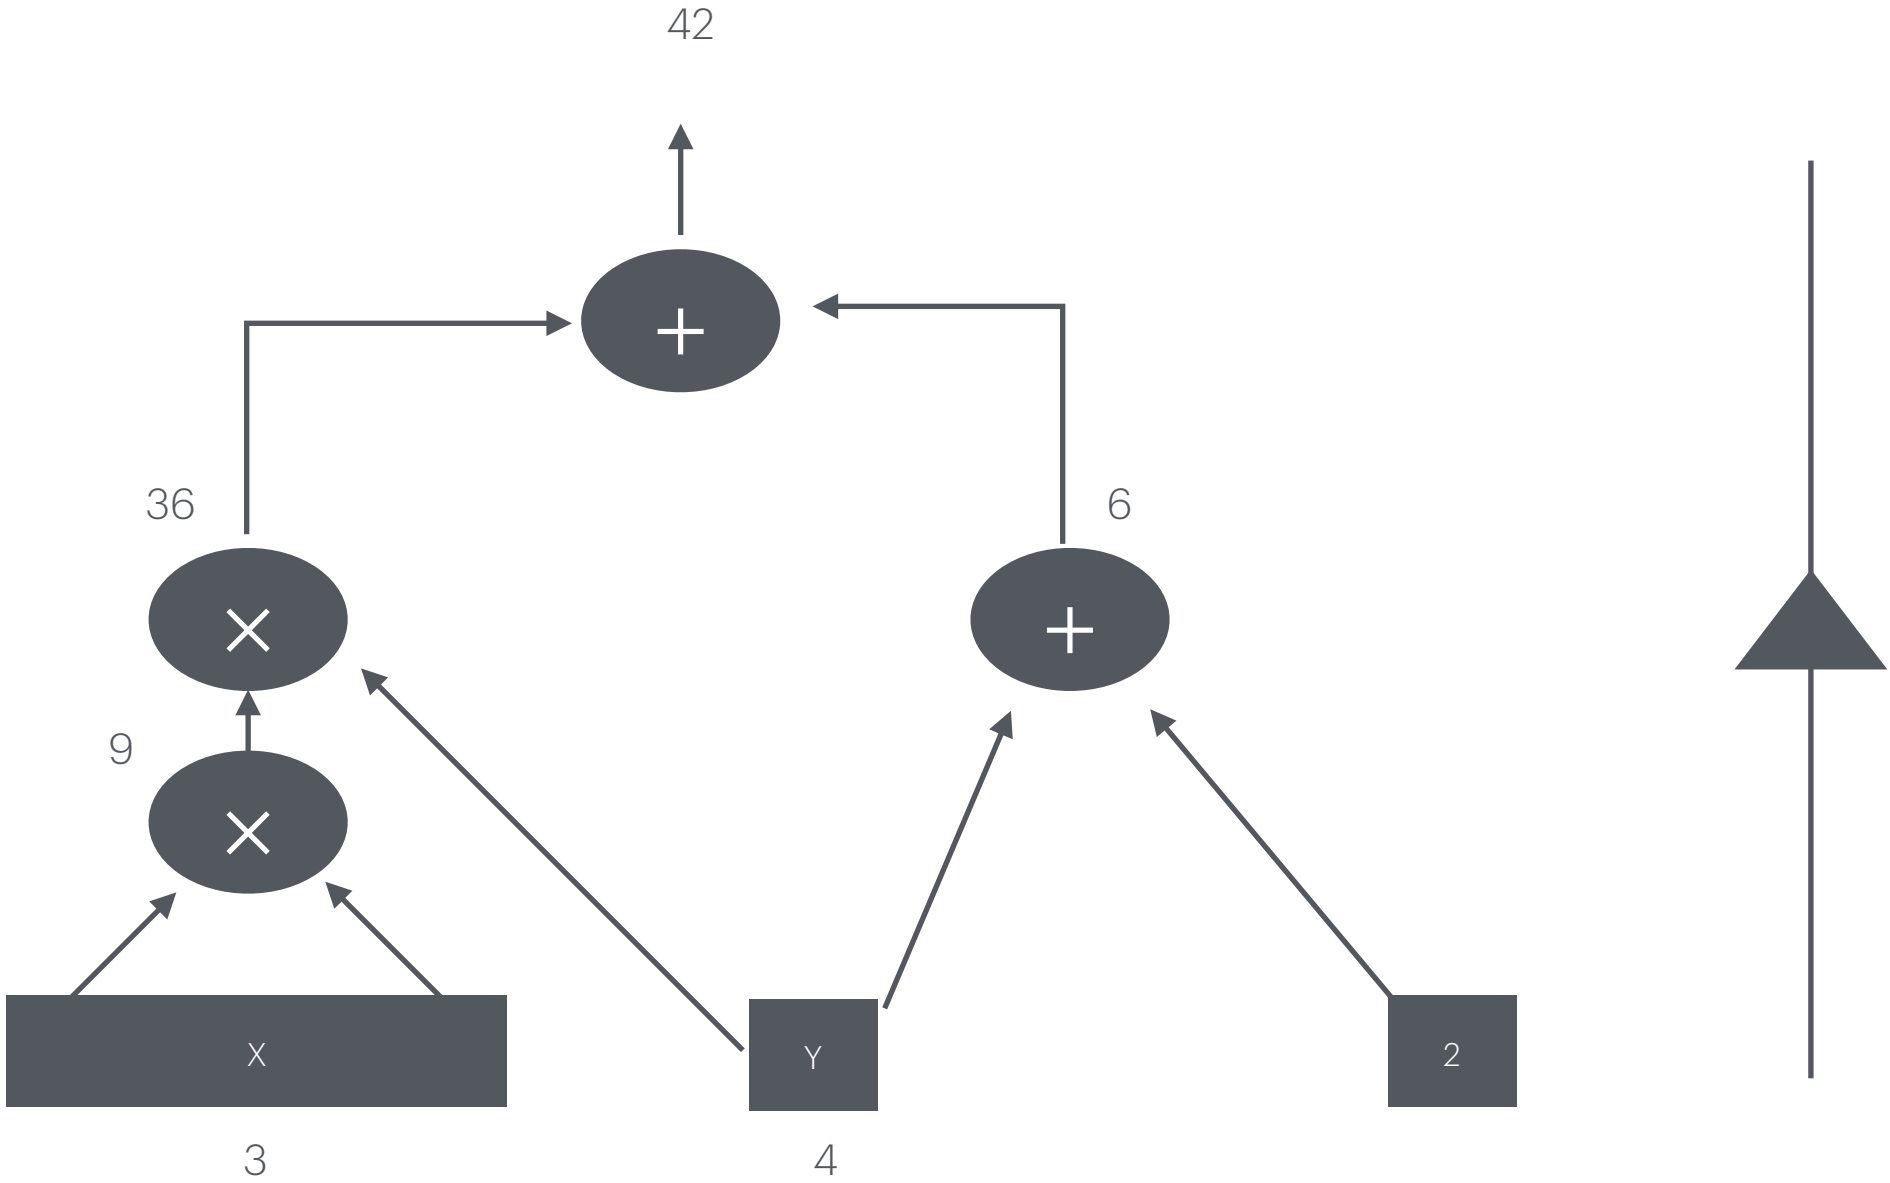


Partial derivative with respect to y requires the same process

Reverse-Mode AutoDifferentiation

$f(x, y) = x^2y + y + 2$

$\frac{\partial f(3,4)}{\partial x} = ?$



n_7 change caused by n_5

$f(x, y)$ change caused by n_7

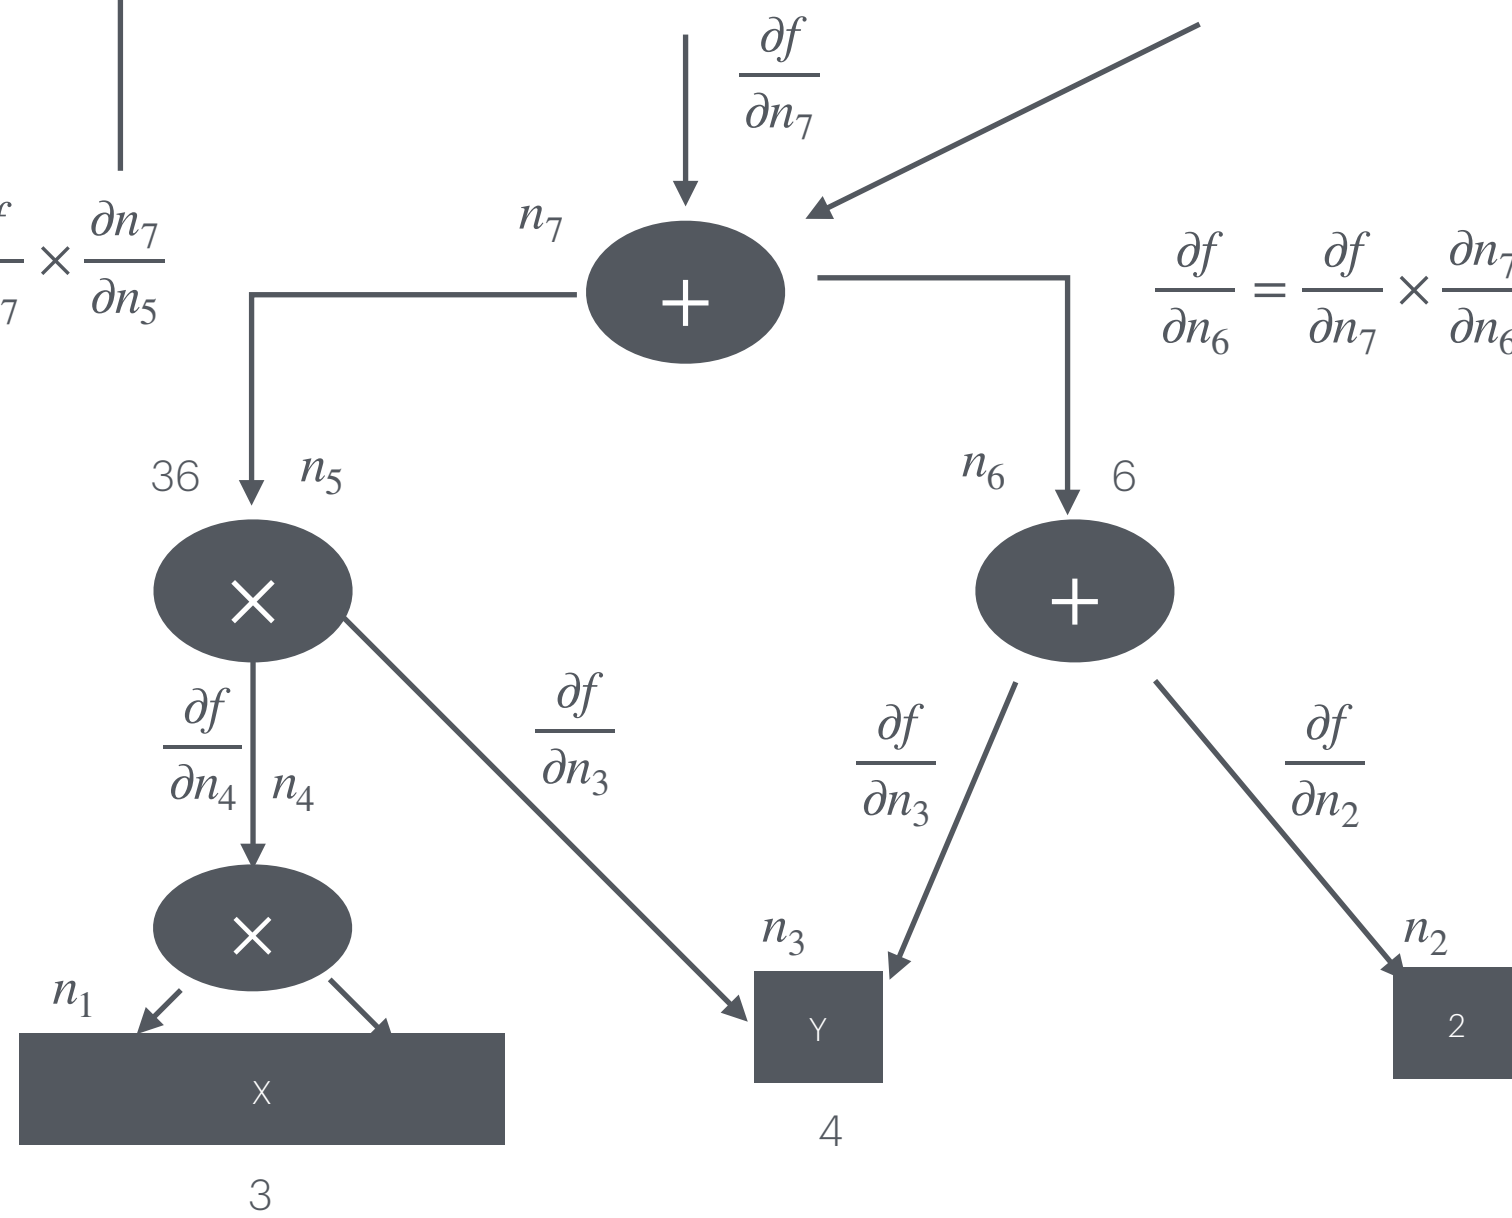
$f(x, y)$ change caused by n_5

$$\frac{\partial f}{\partial n_5} = \frac{\partial f}{\partial n_7} \times \frac{\partial n_7}{\partial n_5}$$


Output n_x at each node. $\frac{\partial f}{\partial n_7} = 1$

$$\frac{\partial f}{\partial n_6} = \frac{\partial f}{\partial n_7} \times \frac{\partial n_7}{\partial n_6}$$

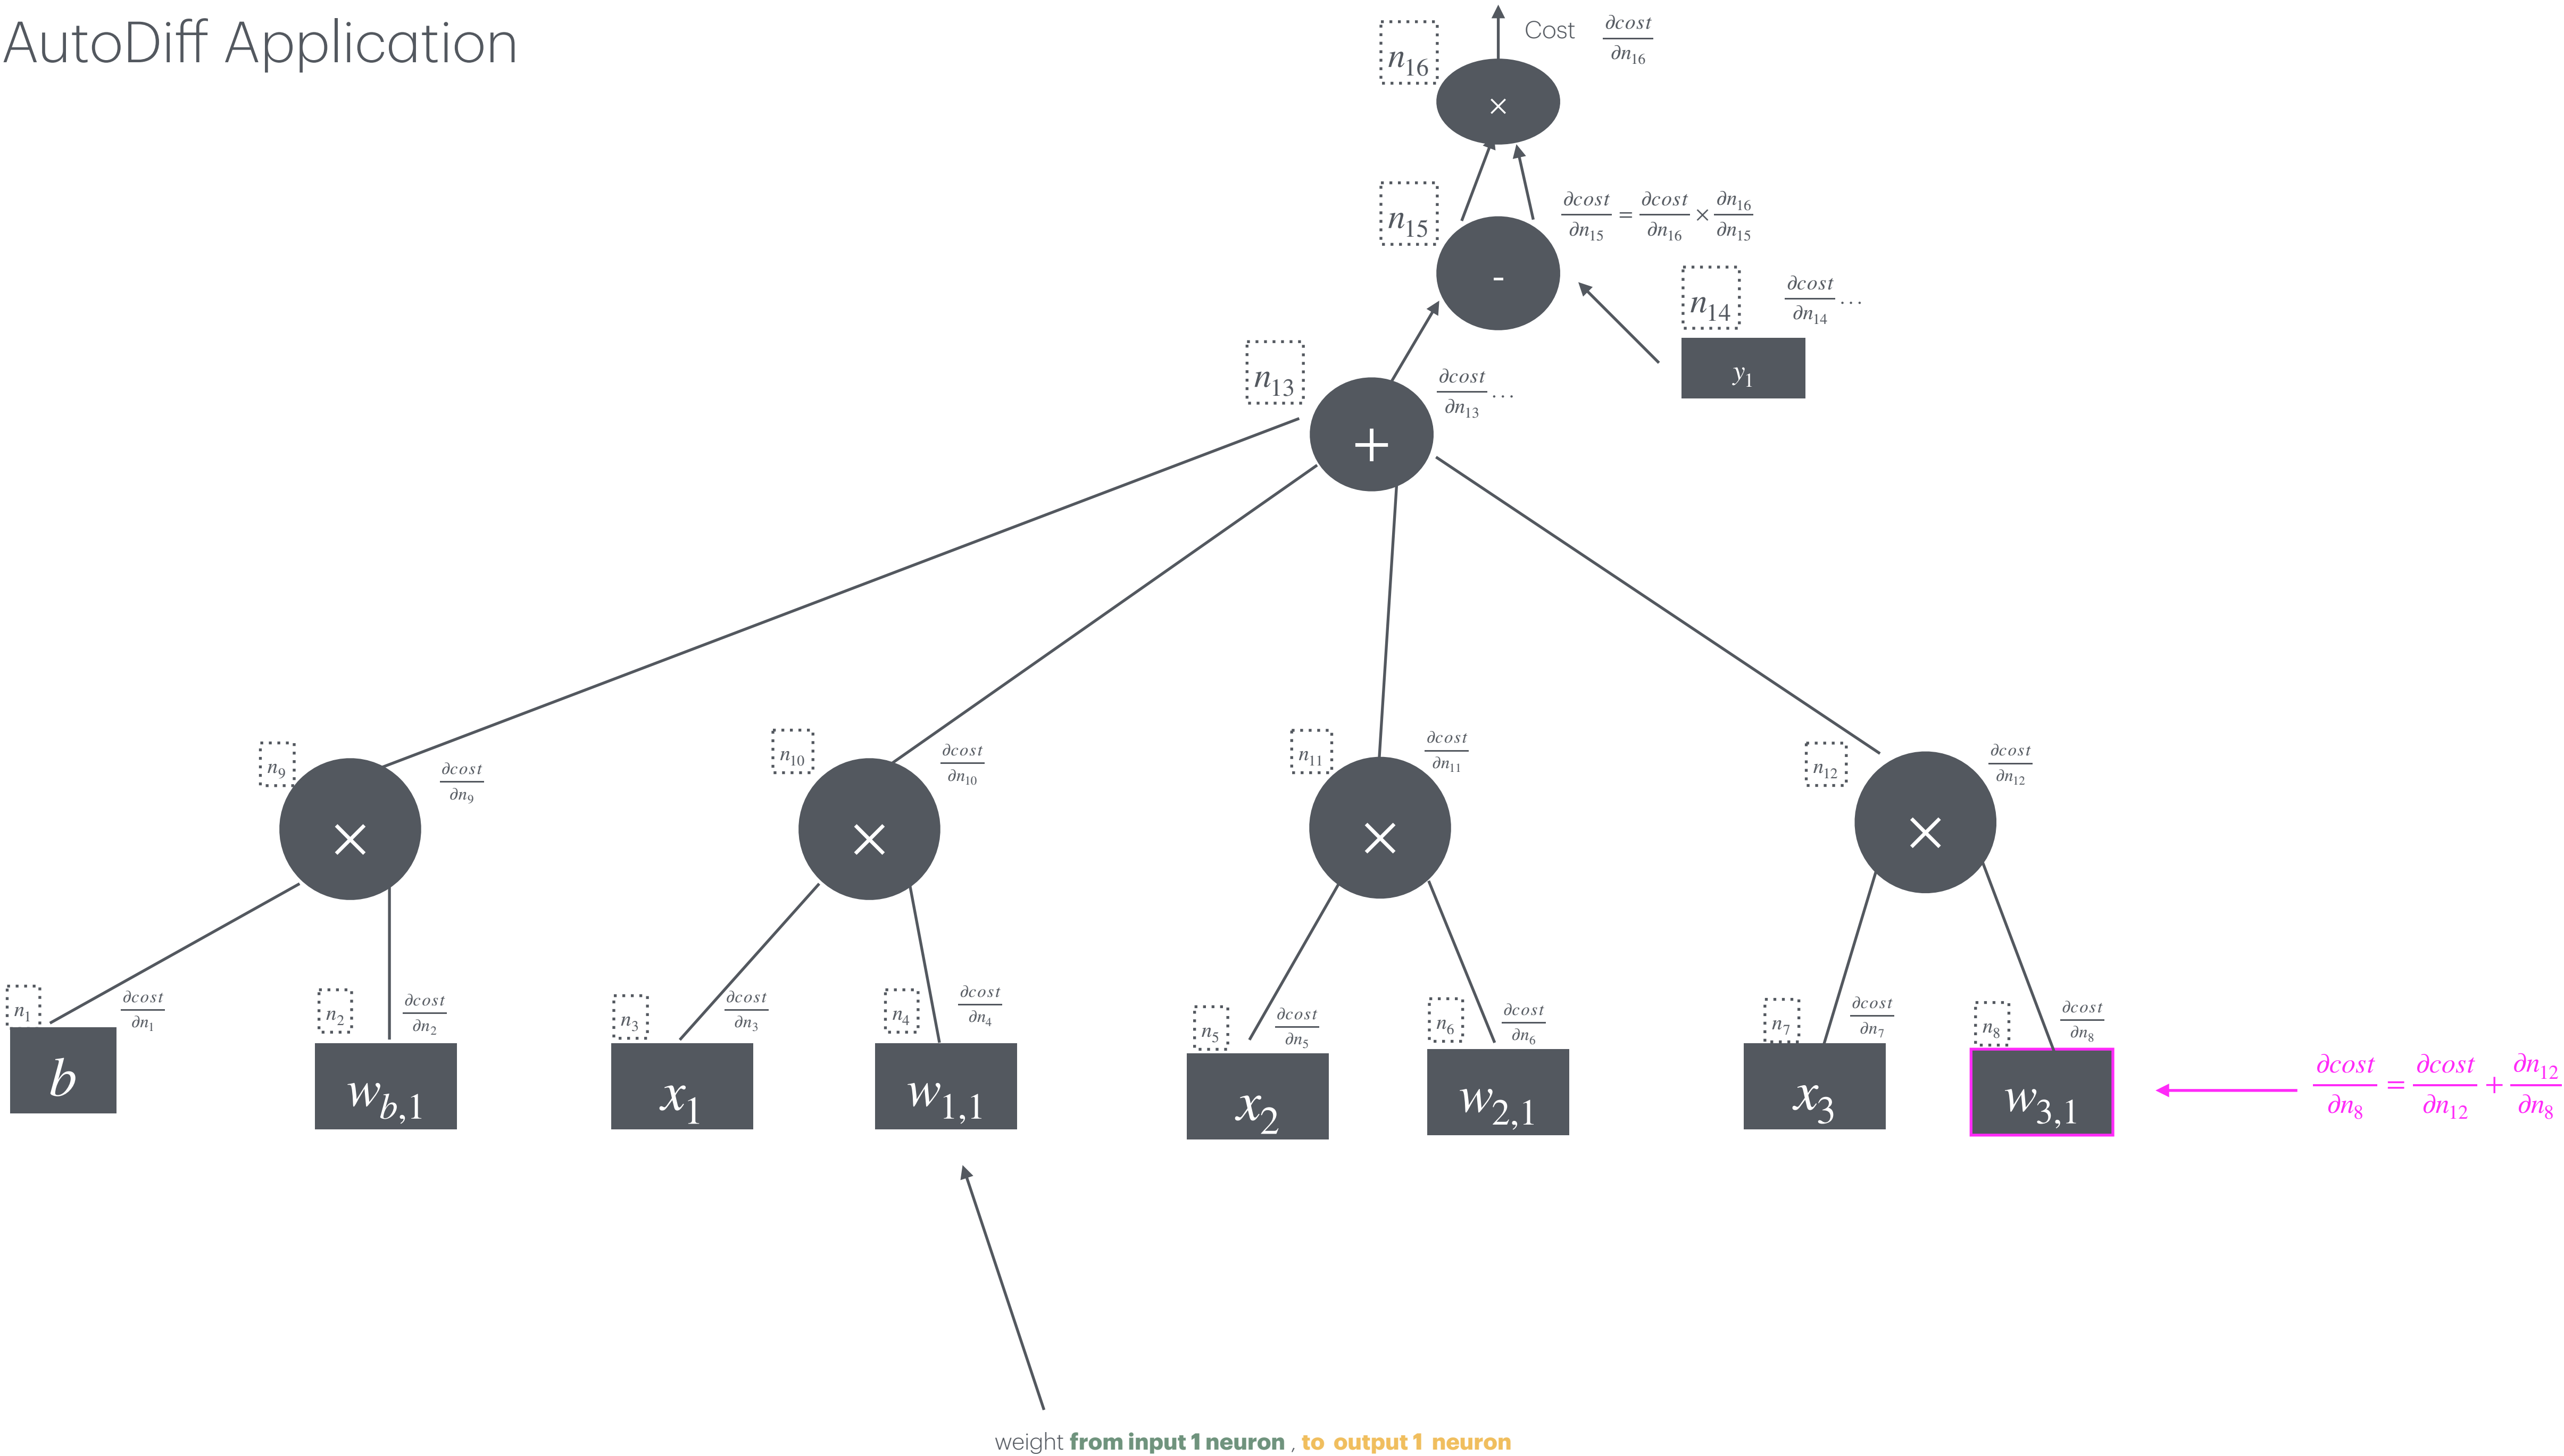
$$\frac{\partial f}{\partial n_1}$$



AutoDiff Application



AutoDiff Application



Tensorflow

Keras(API)

Keras(implementation)

Libraries containing Keras

- ★ Tensorflow
- Microsoft Cognitive Toolkit
- Theano
- ★ Keras(API)
- ★ PyTorch

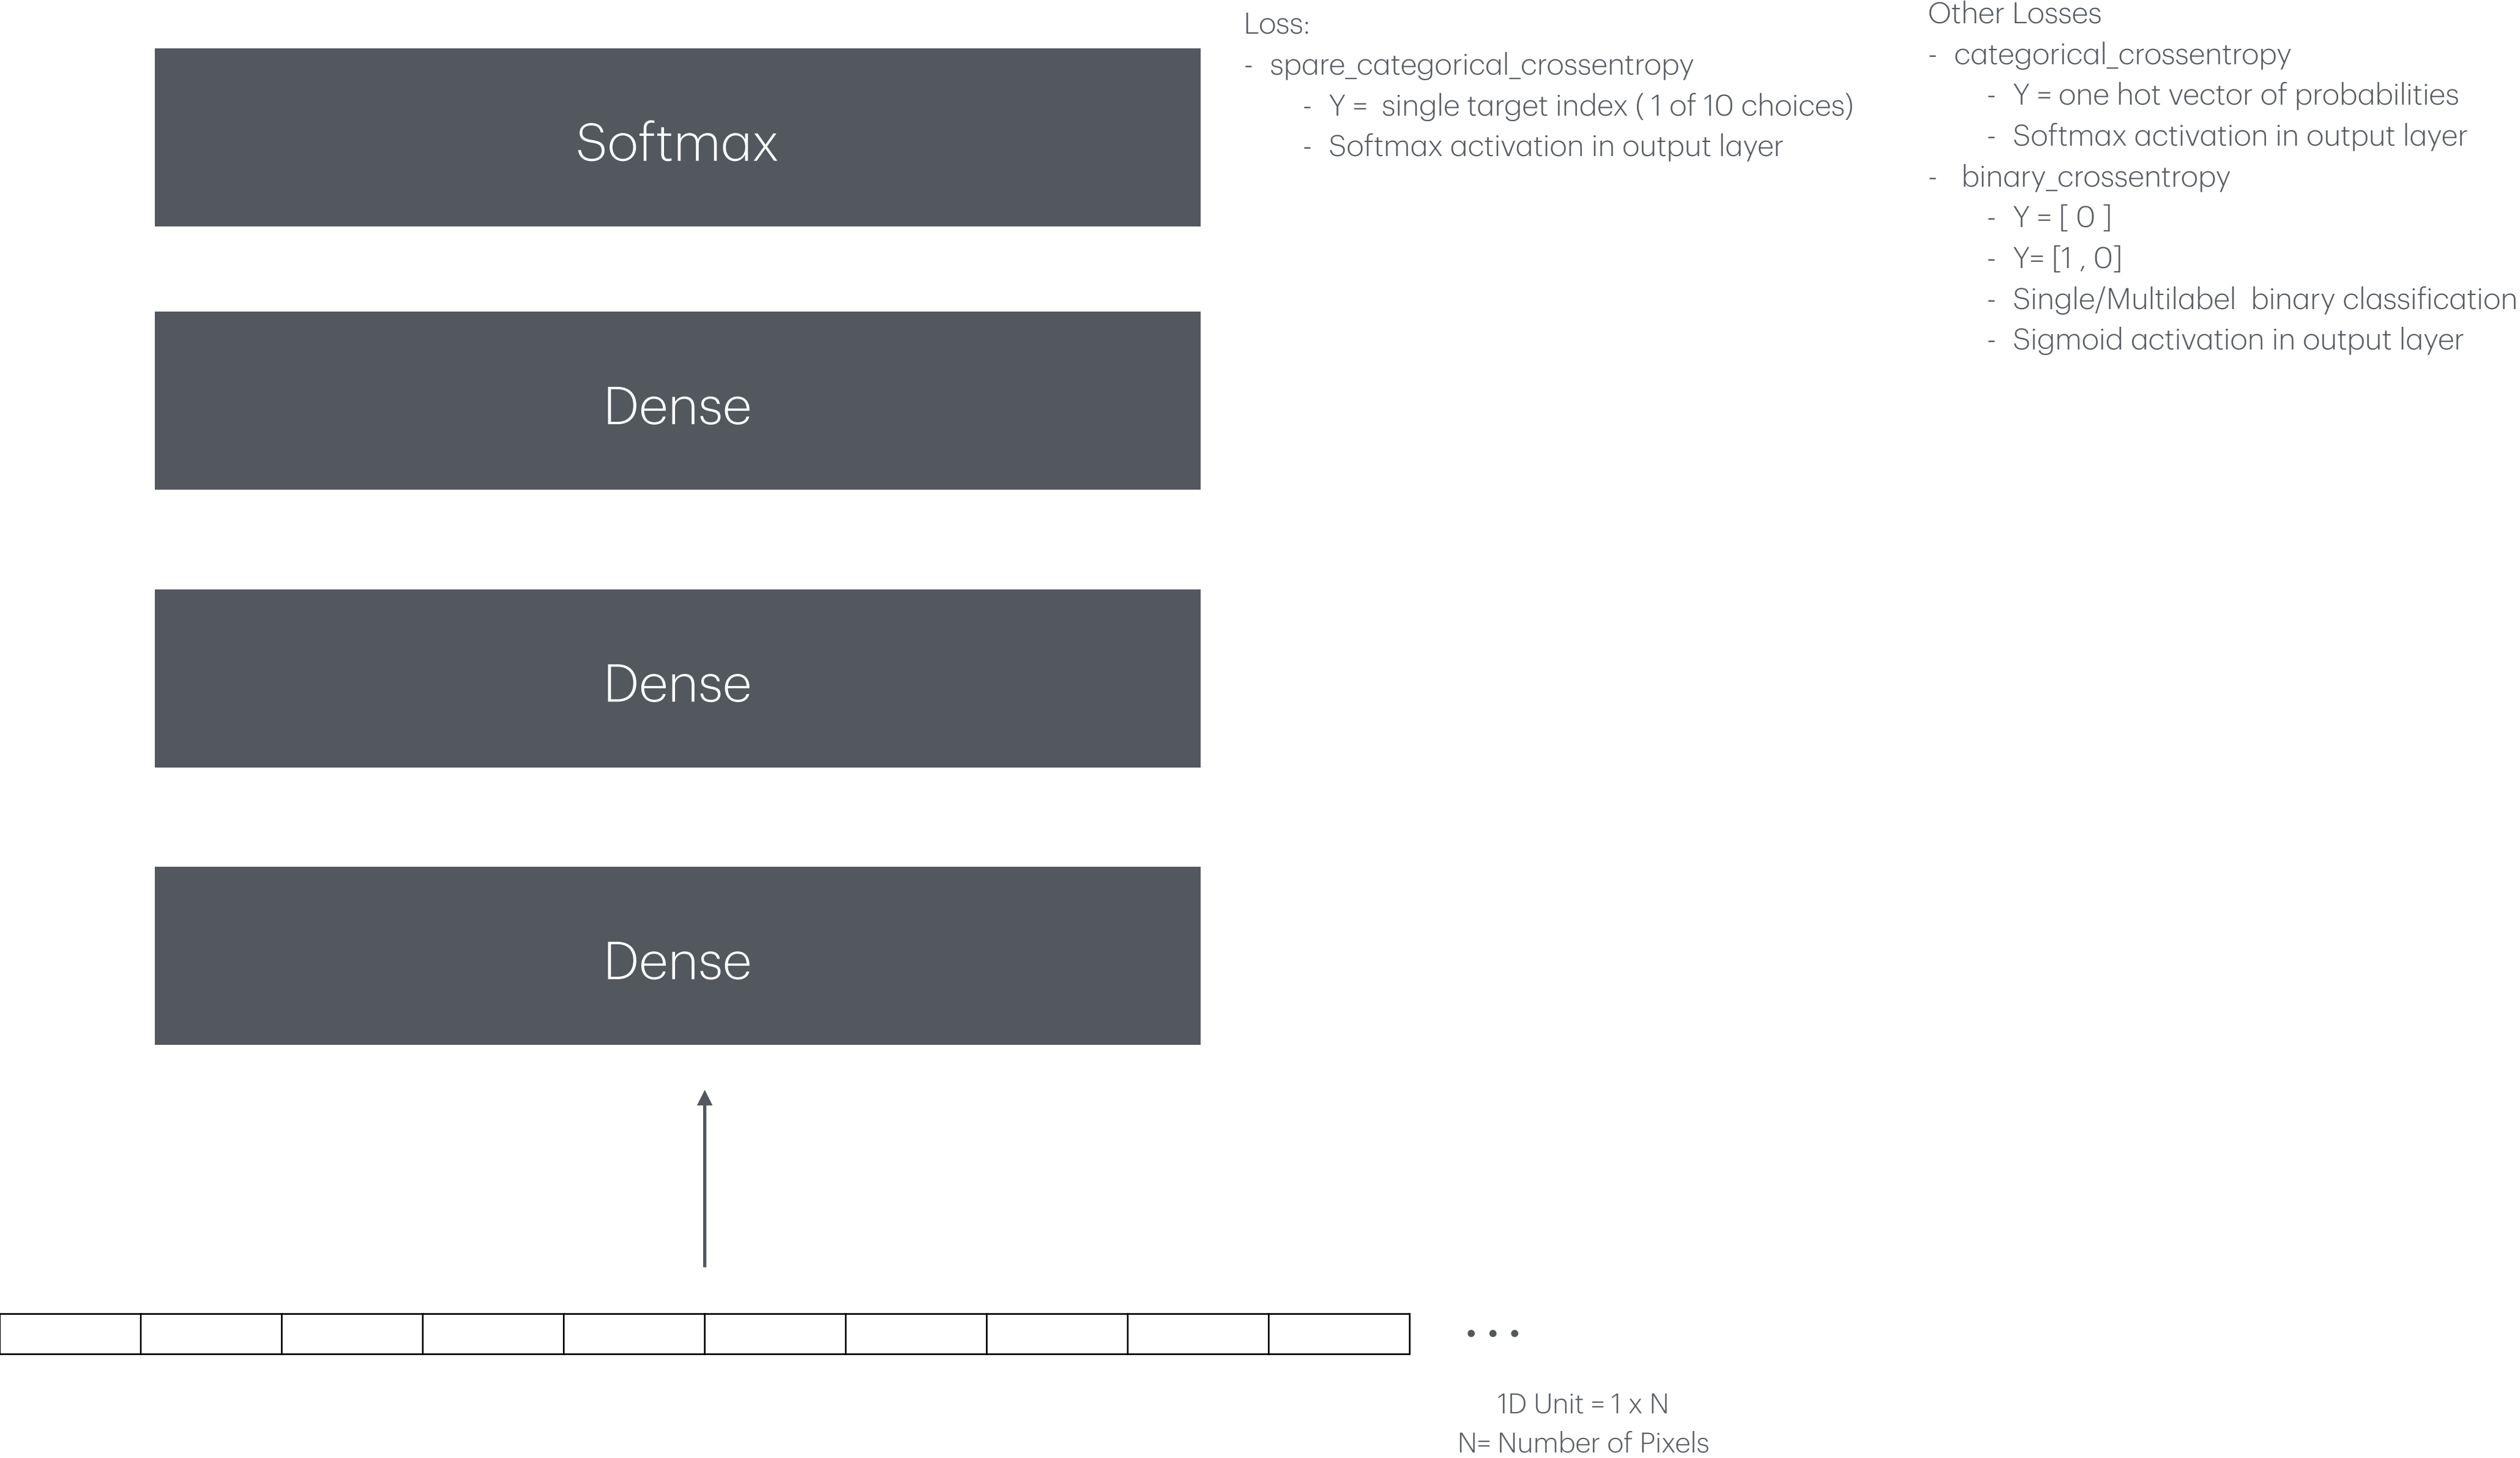
★ – *popular*

Others containing Keras

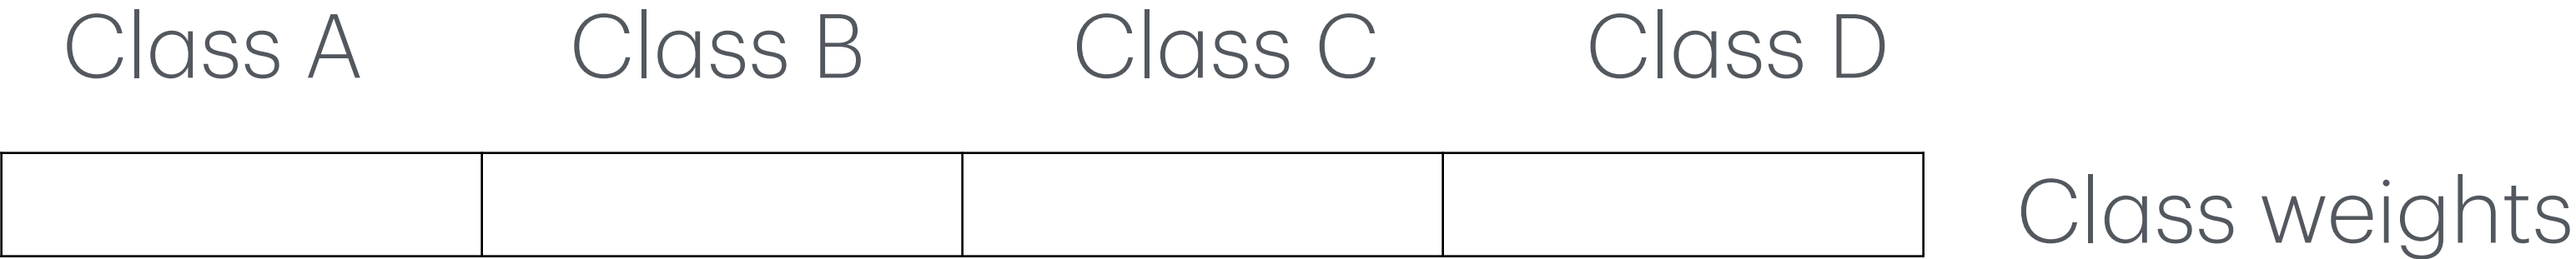
- Javascript/Typescript
- PlaidML
- Apple's Core ML
- Apache MXNet



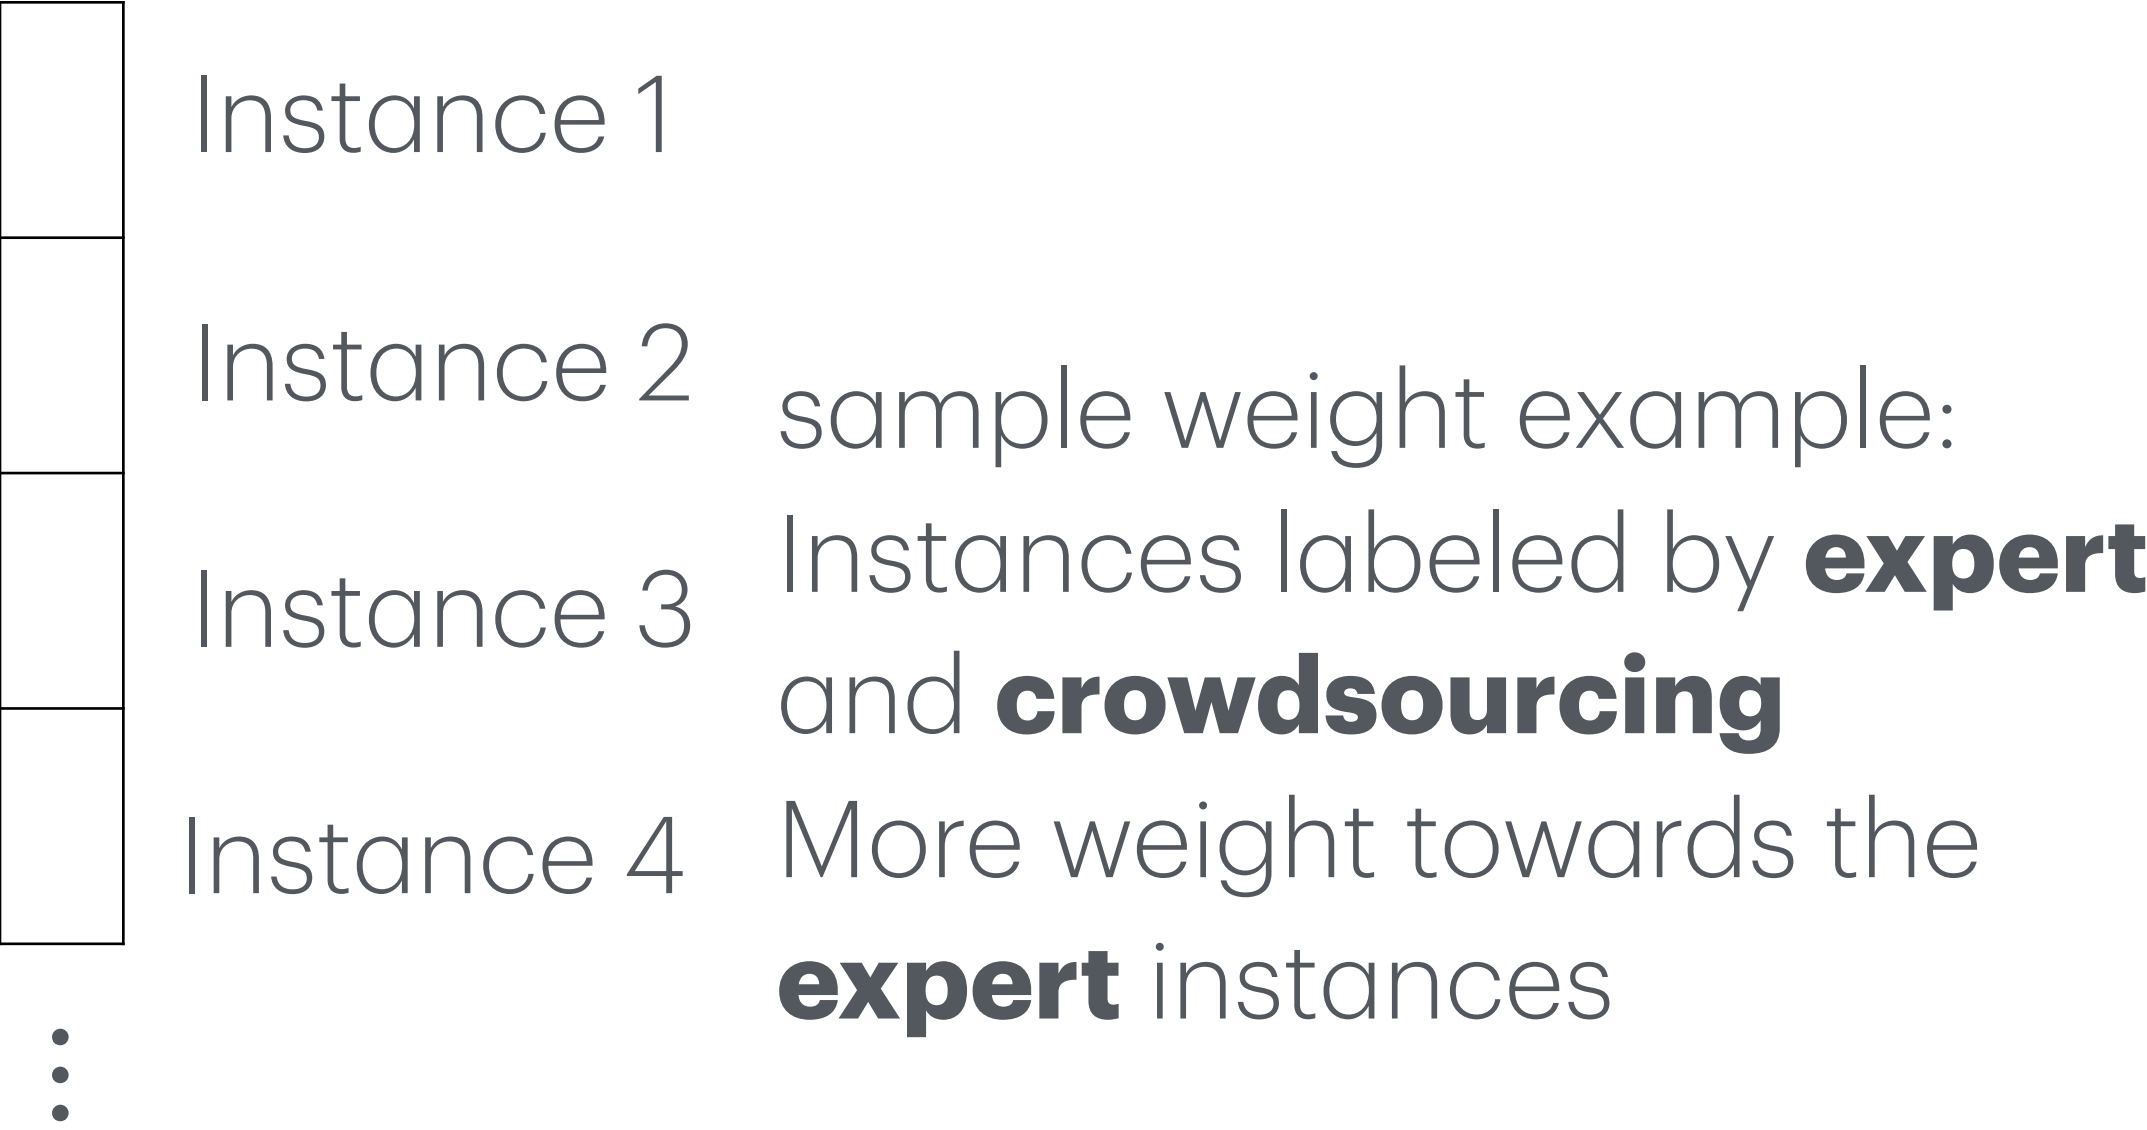
Sequential Model: Classify Fashion MNIST



Dealing with skewed data



class weight example:
Class A **overrepresented** in
dataset. Give **more weight** to
Class B,C, and D



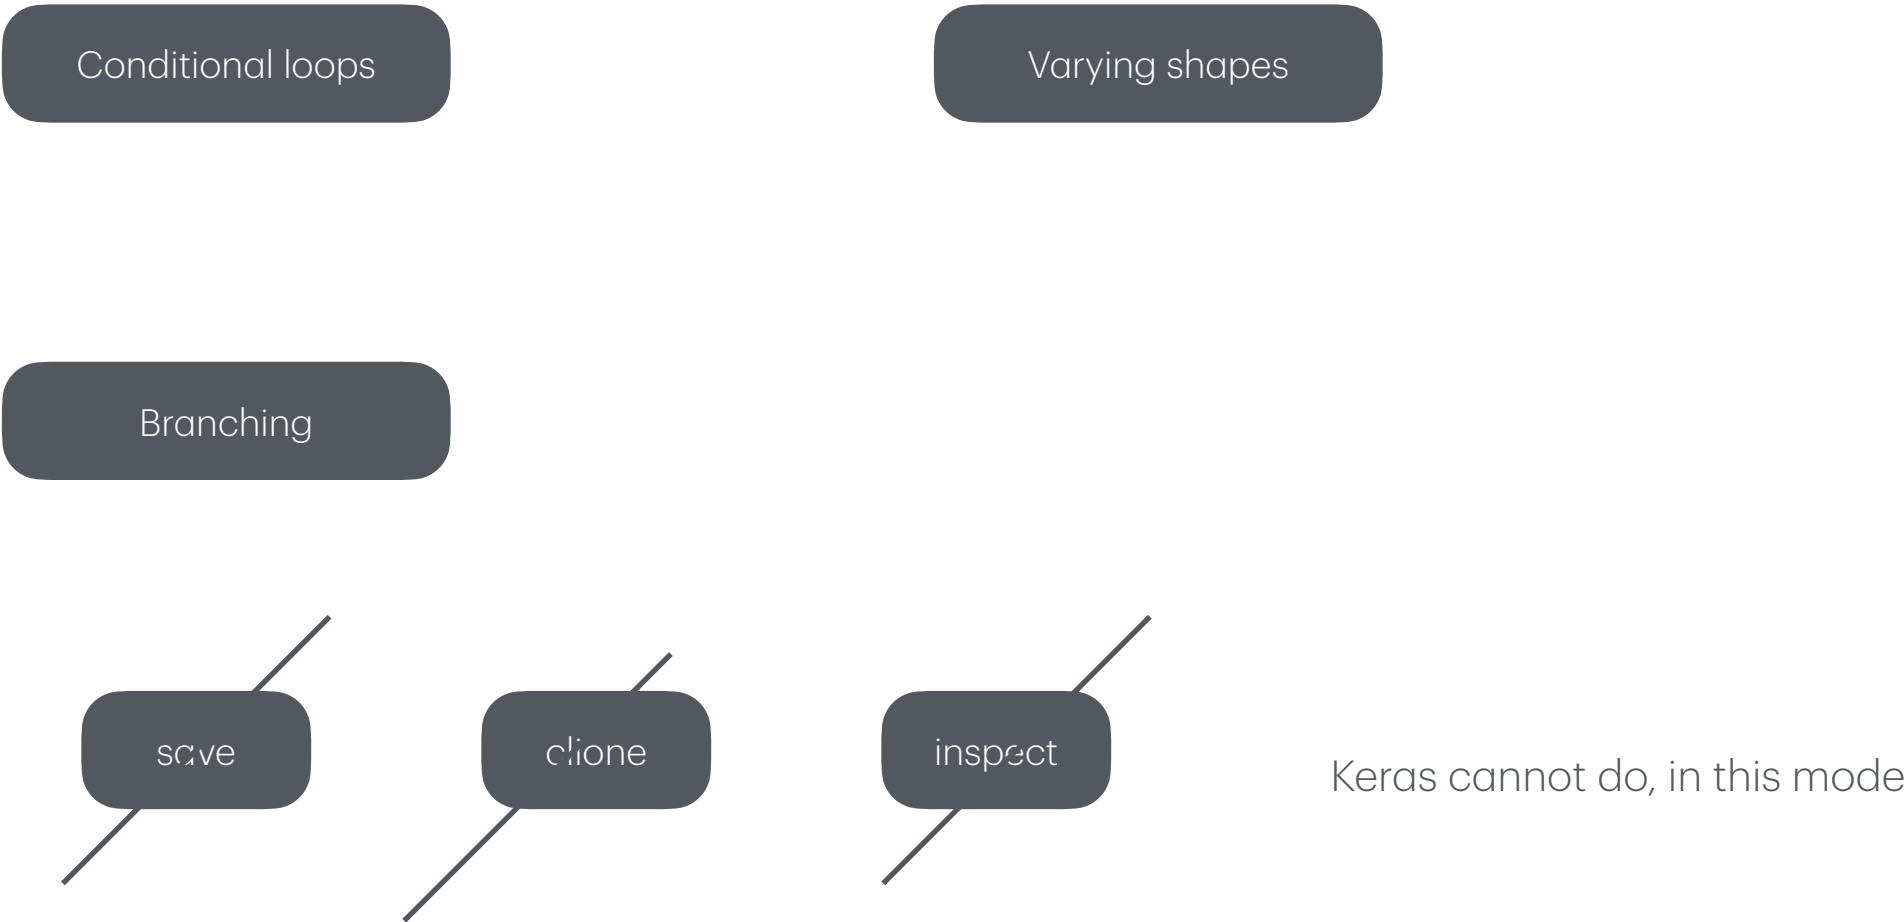
Samples weights

Saving

- Functional
- Sequential



- Subclassing

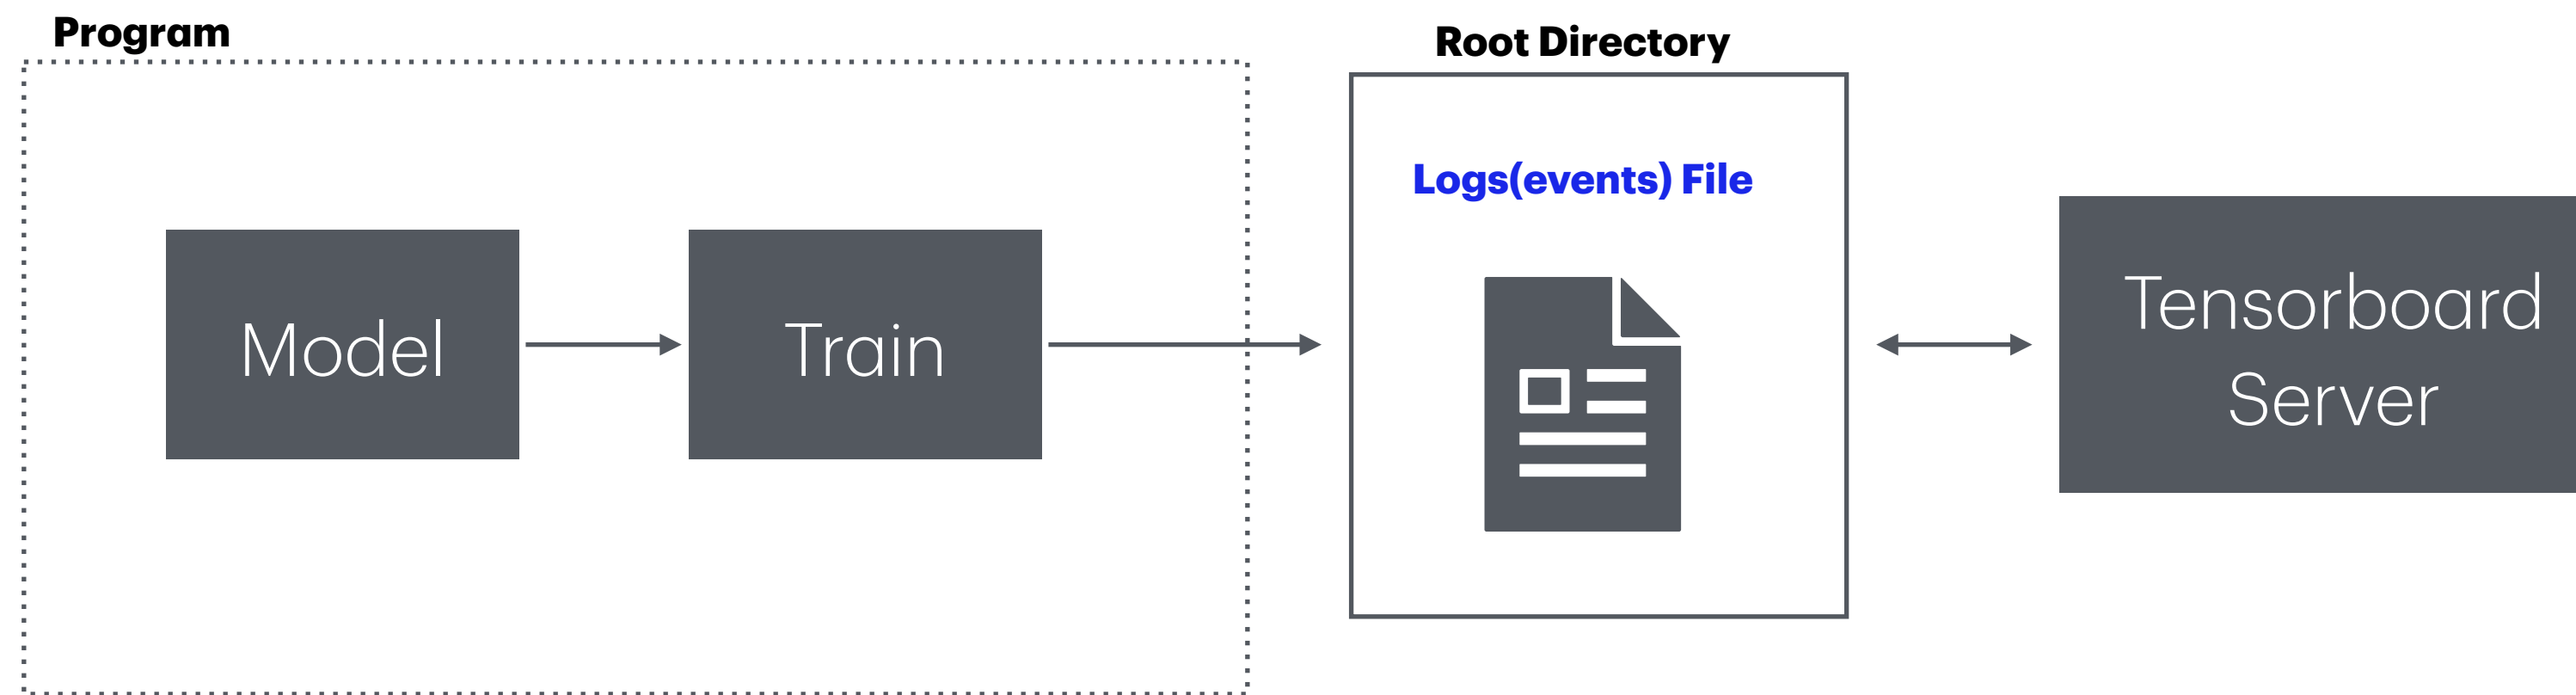


Save and load model weights yourself

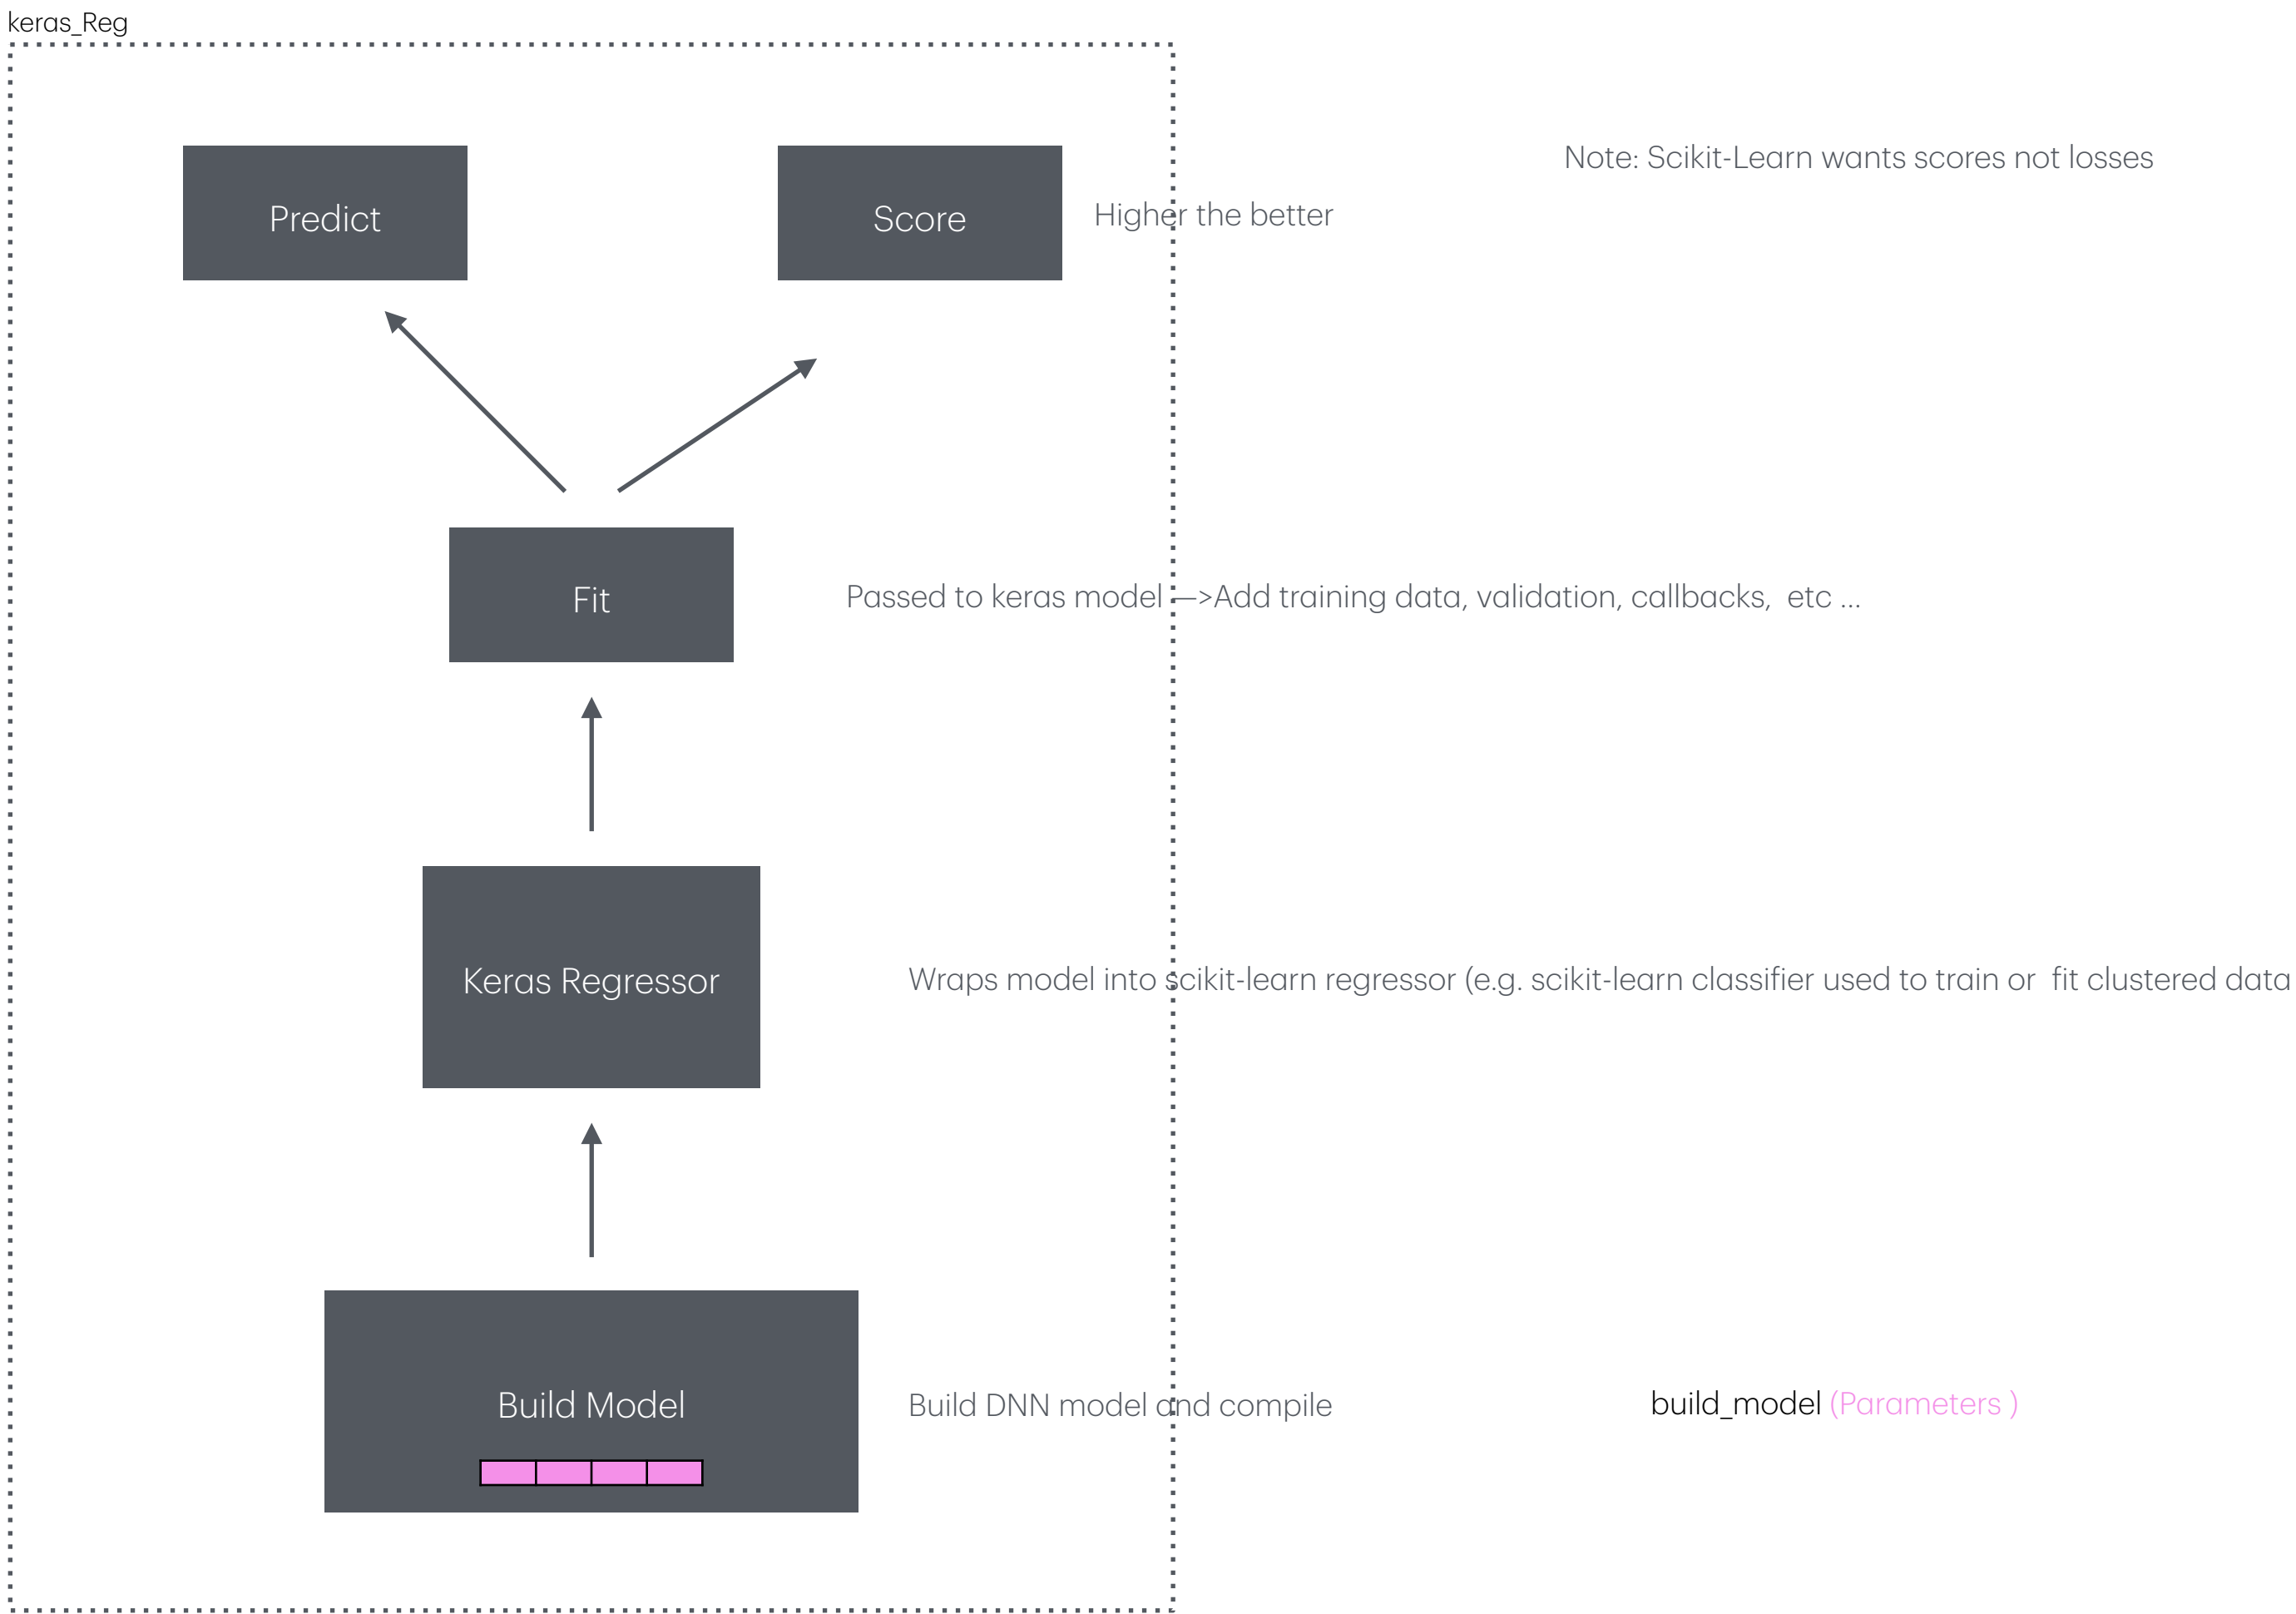
Saving



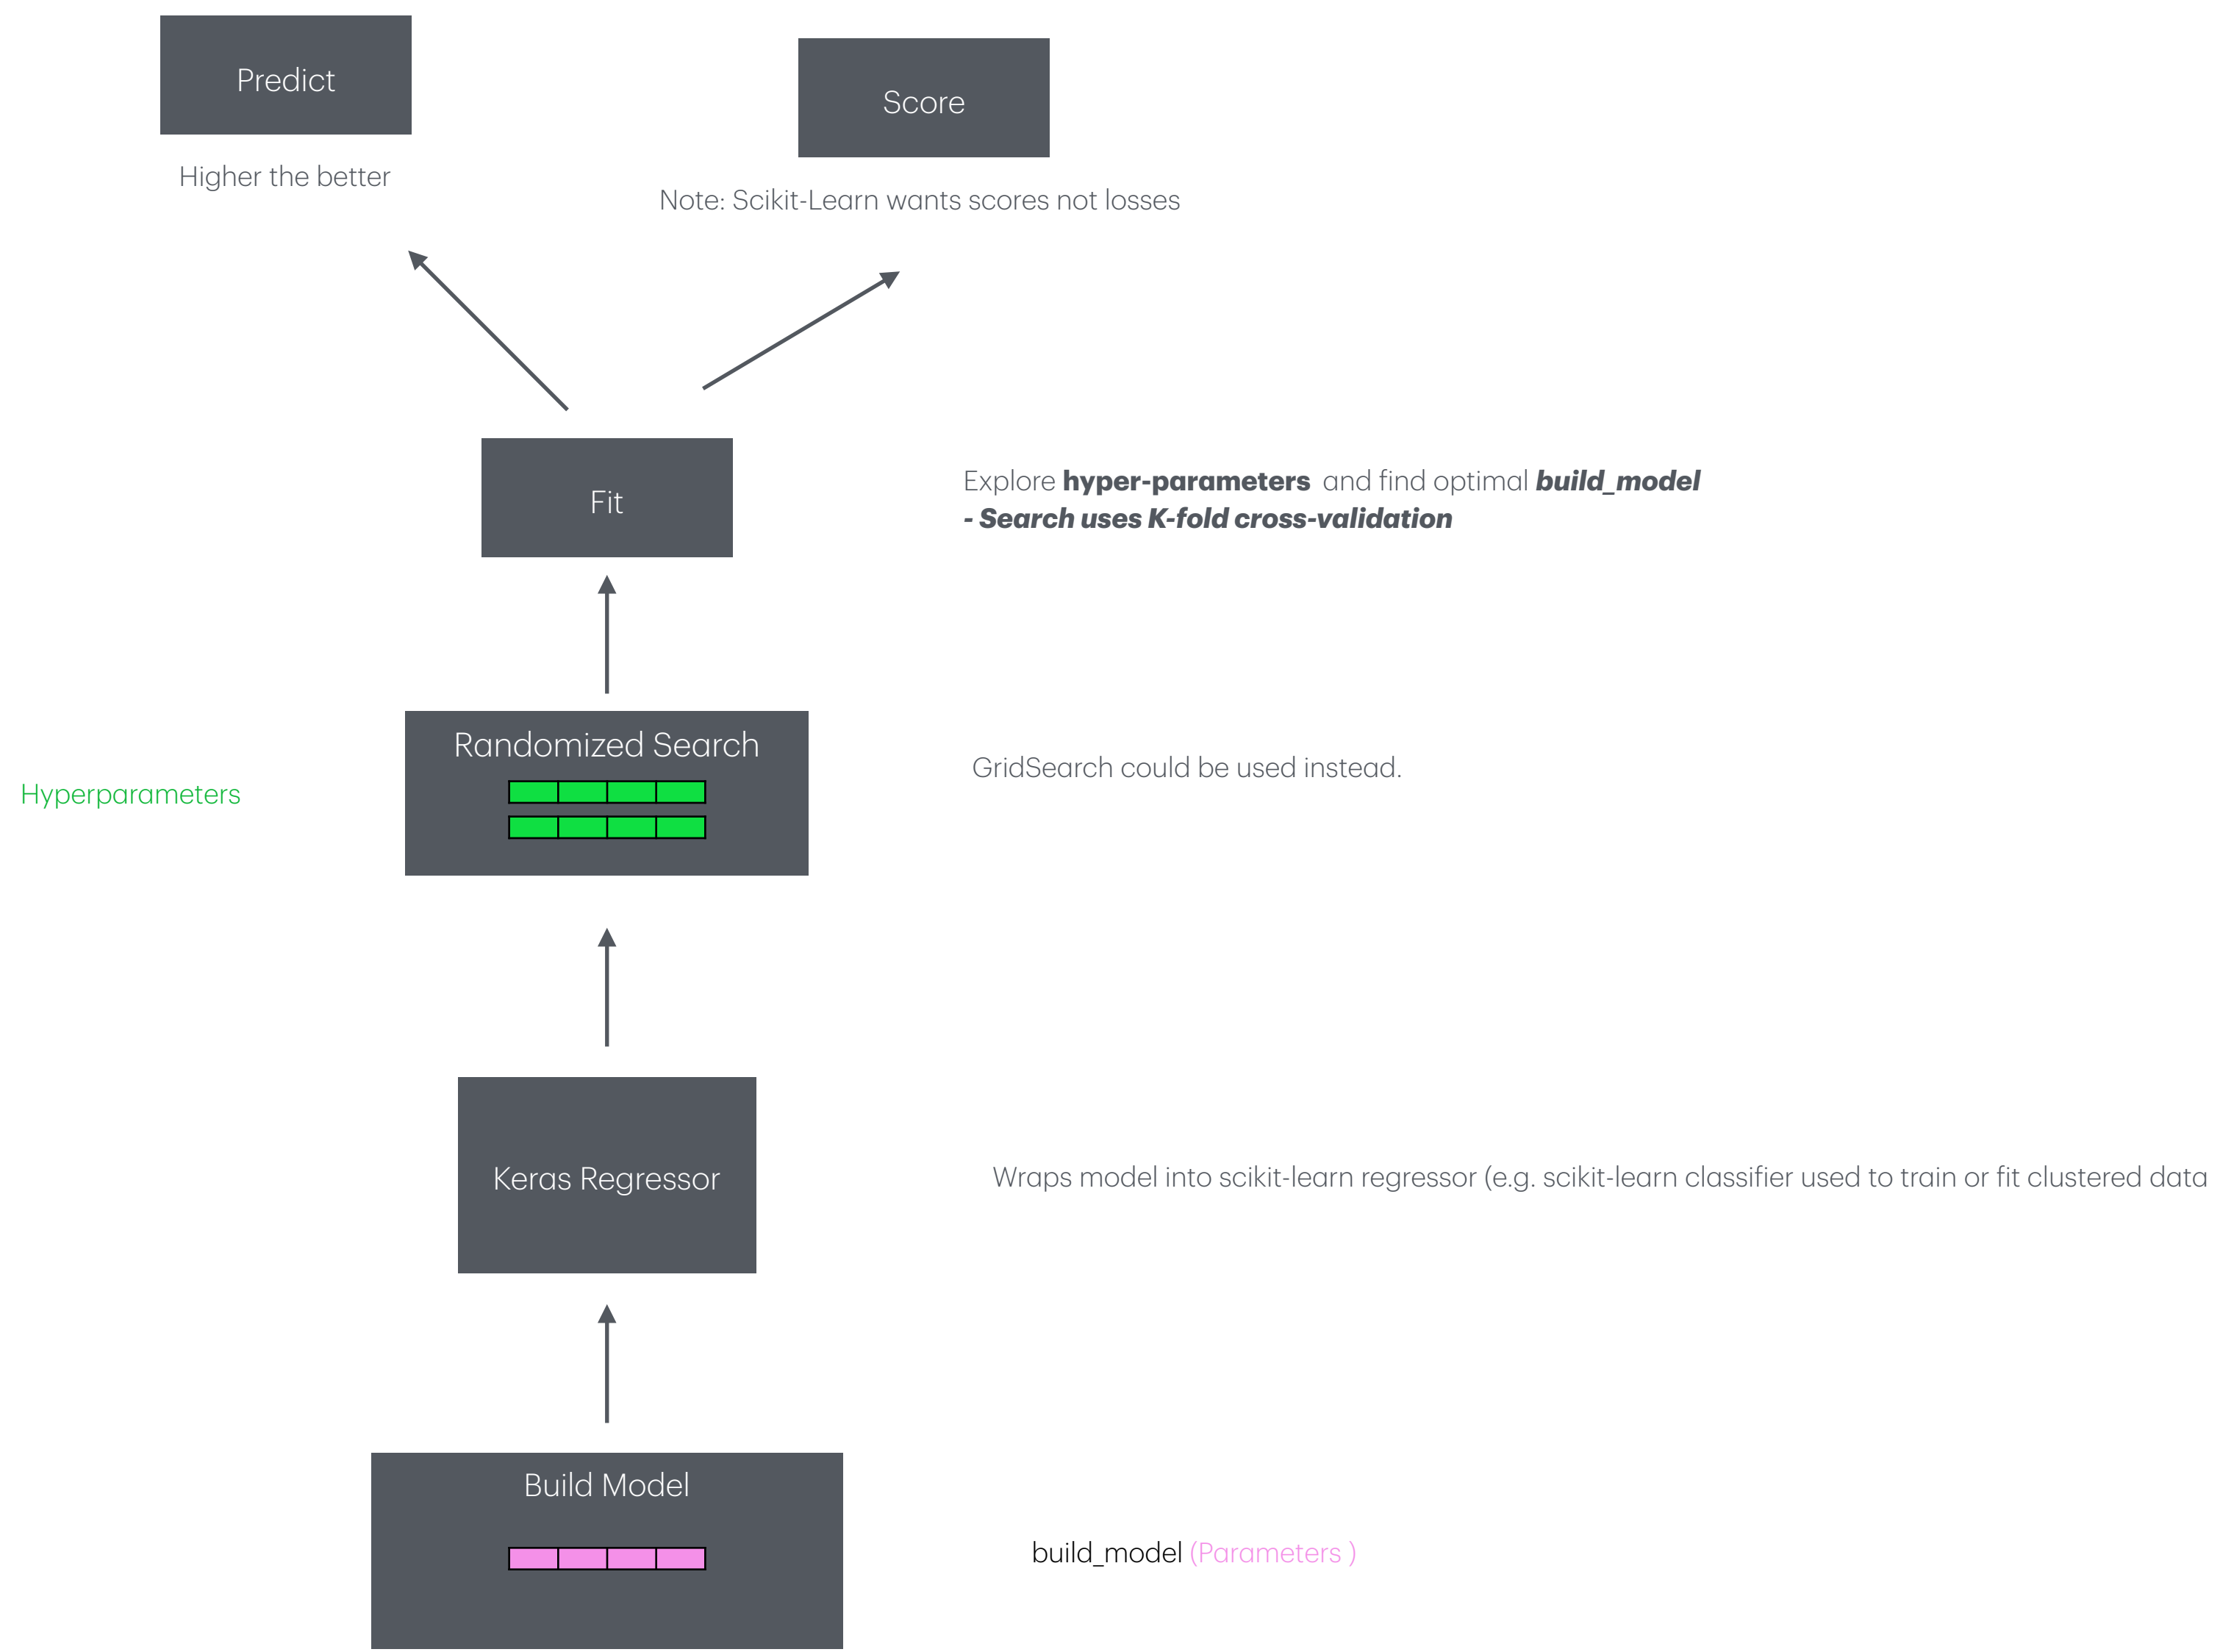
Tensorboard




Fine-Tuning Neural Networks



Fine-Tuning Neural Networks



Networks learn in hierarchical way



Complex problems deep networks
Have higher parameter efficiency.
Fewer neurons needed per layer



Shallow network can solve many
problems with enough neurons



Epoch
000,101

Learning rate
0.03

Activation
Tanh

Regularization
None

Regularization rate
0

Problem type
Classification

DATA

Which dataset do you want to use?



Ratio of training to test data: 60%

Noise: 5

Batch size: 19

REGENERATE

FEATURES

Which properties do you want to feed in?

X_1



X_2



X_1^2



X_2^2



X_1X_2



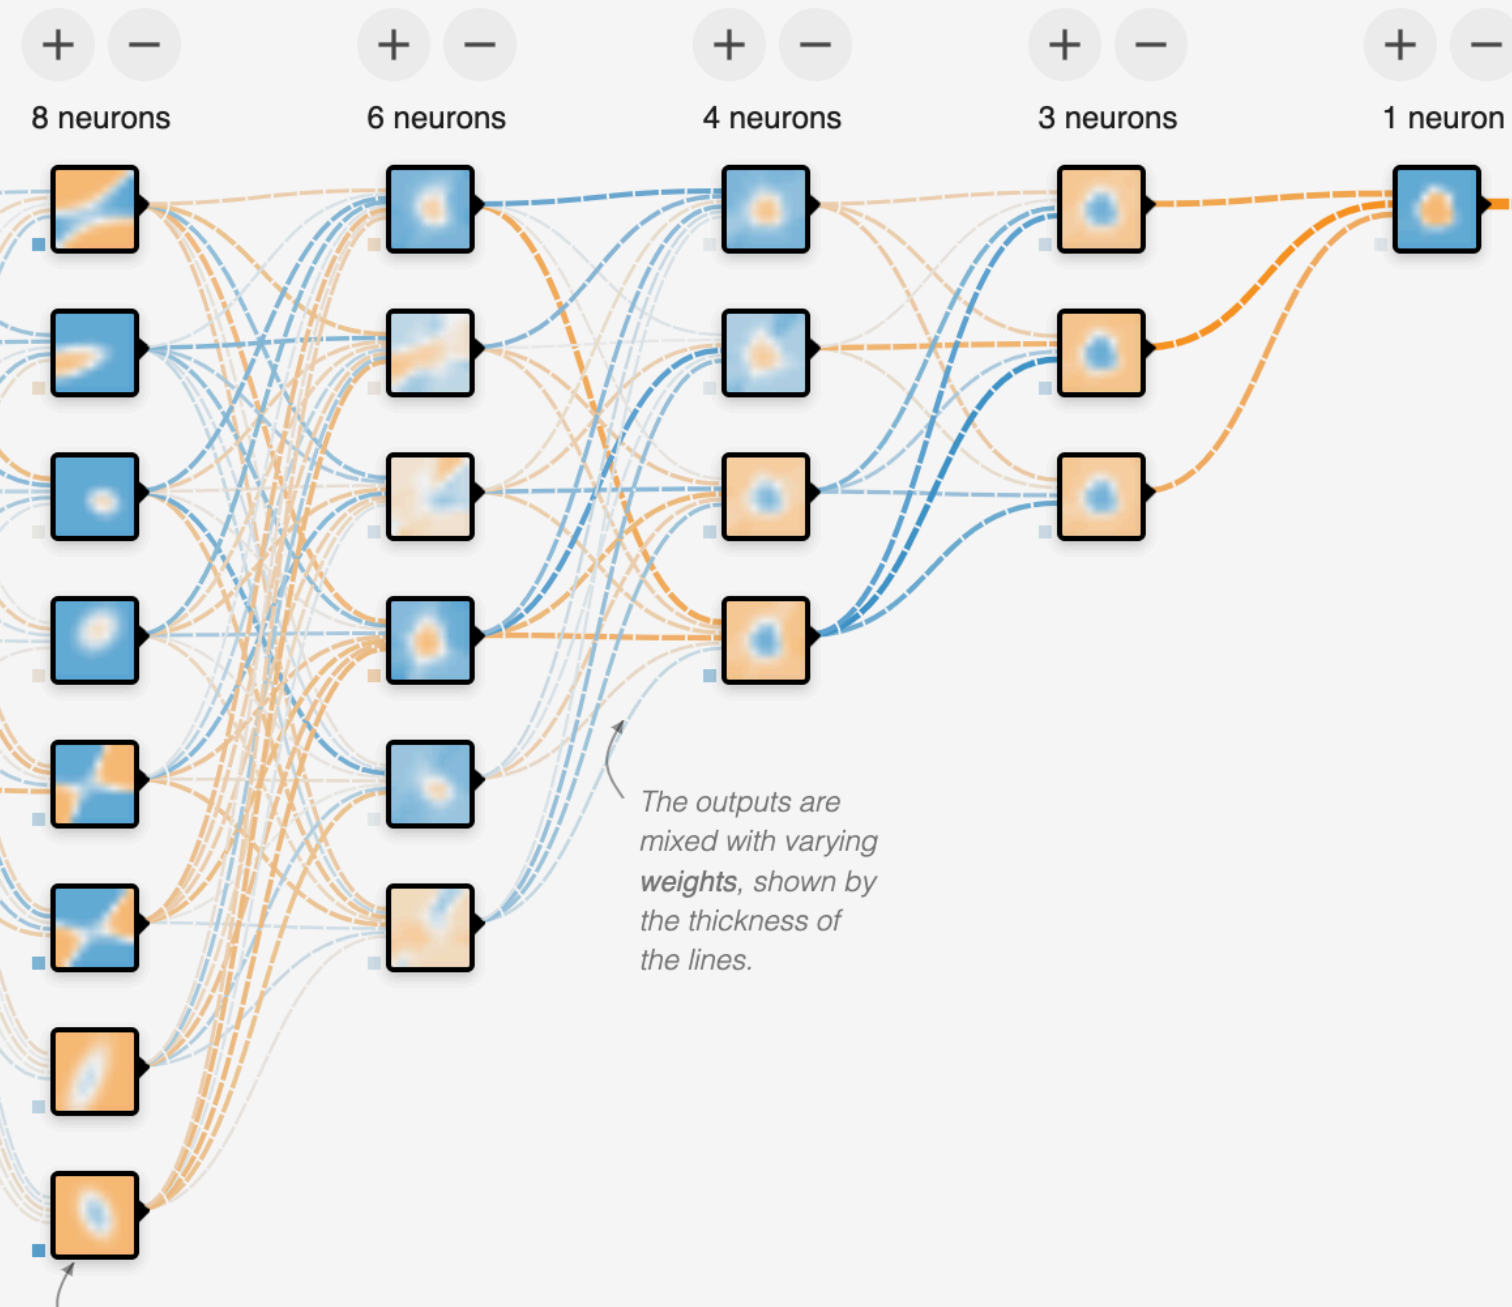
$\sin(X_1)$



$\sin(X_2)$




5 HIDDEN LAYERS




OUTPUT

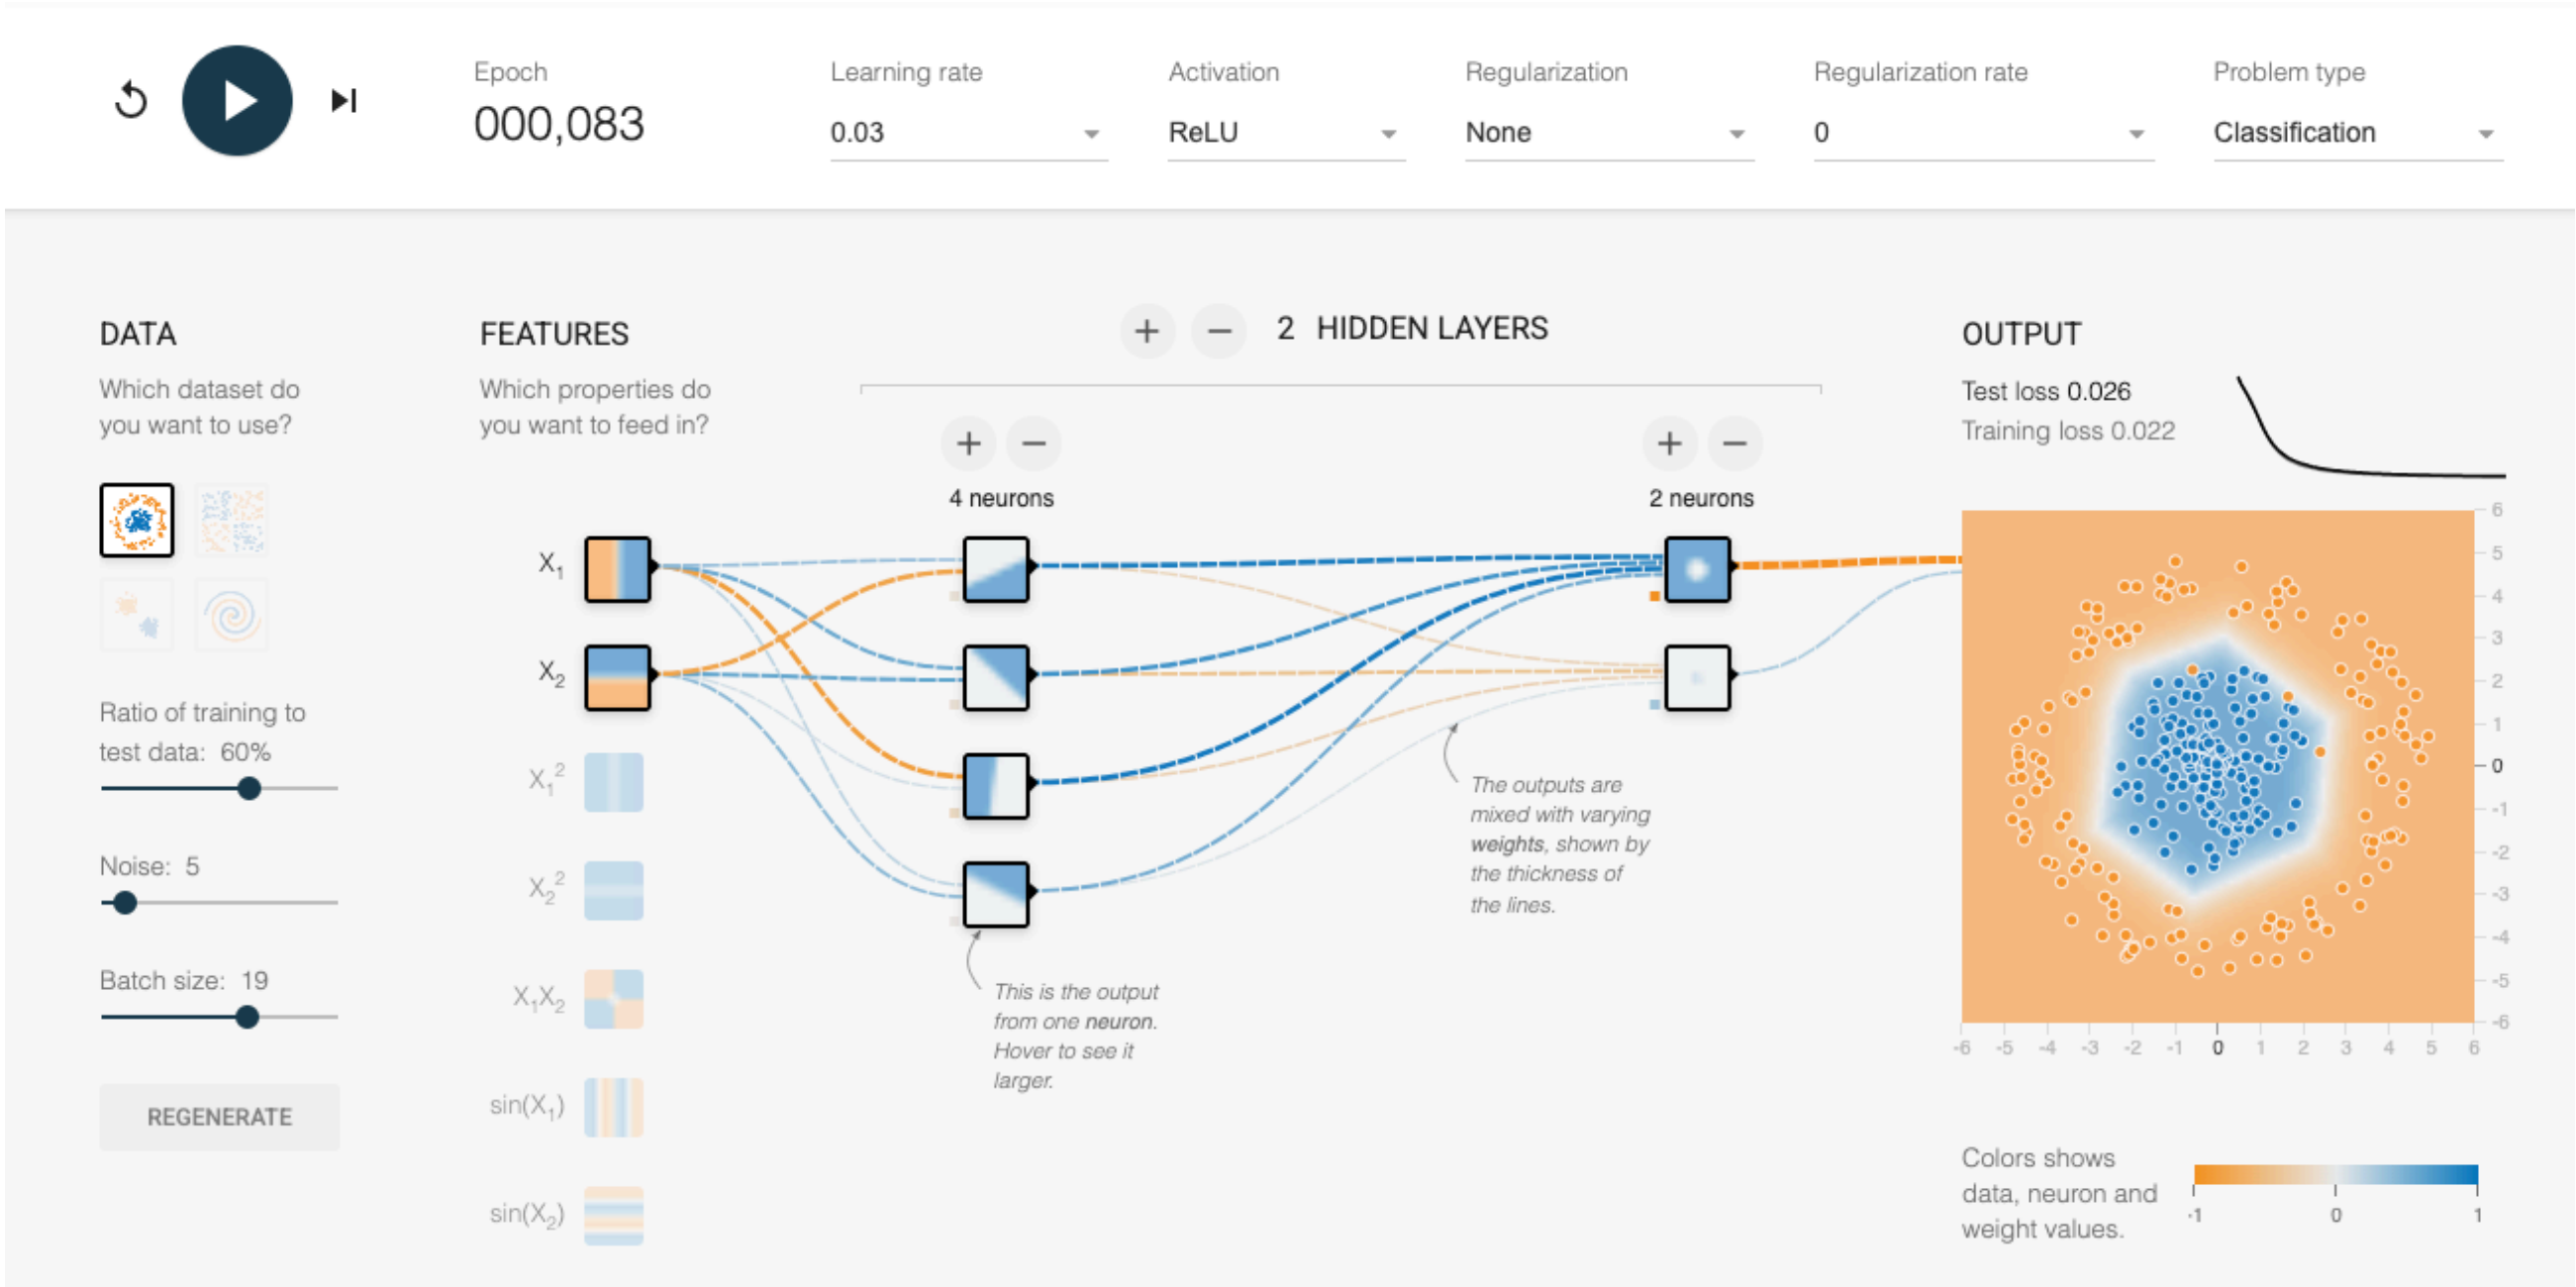
Test loss 0.014
Training loss 0.009



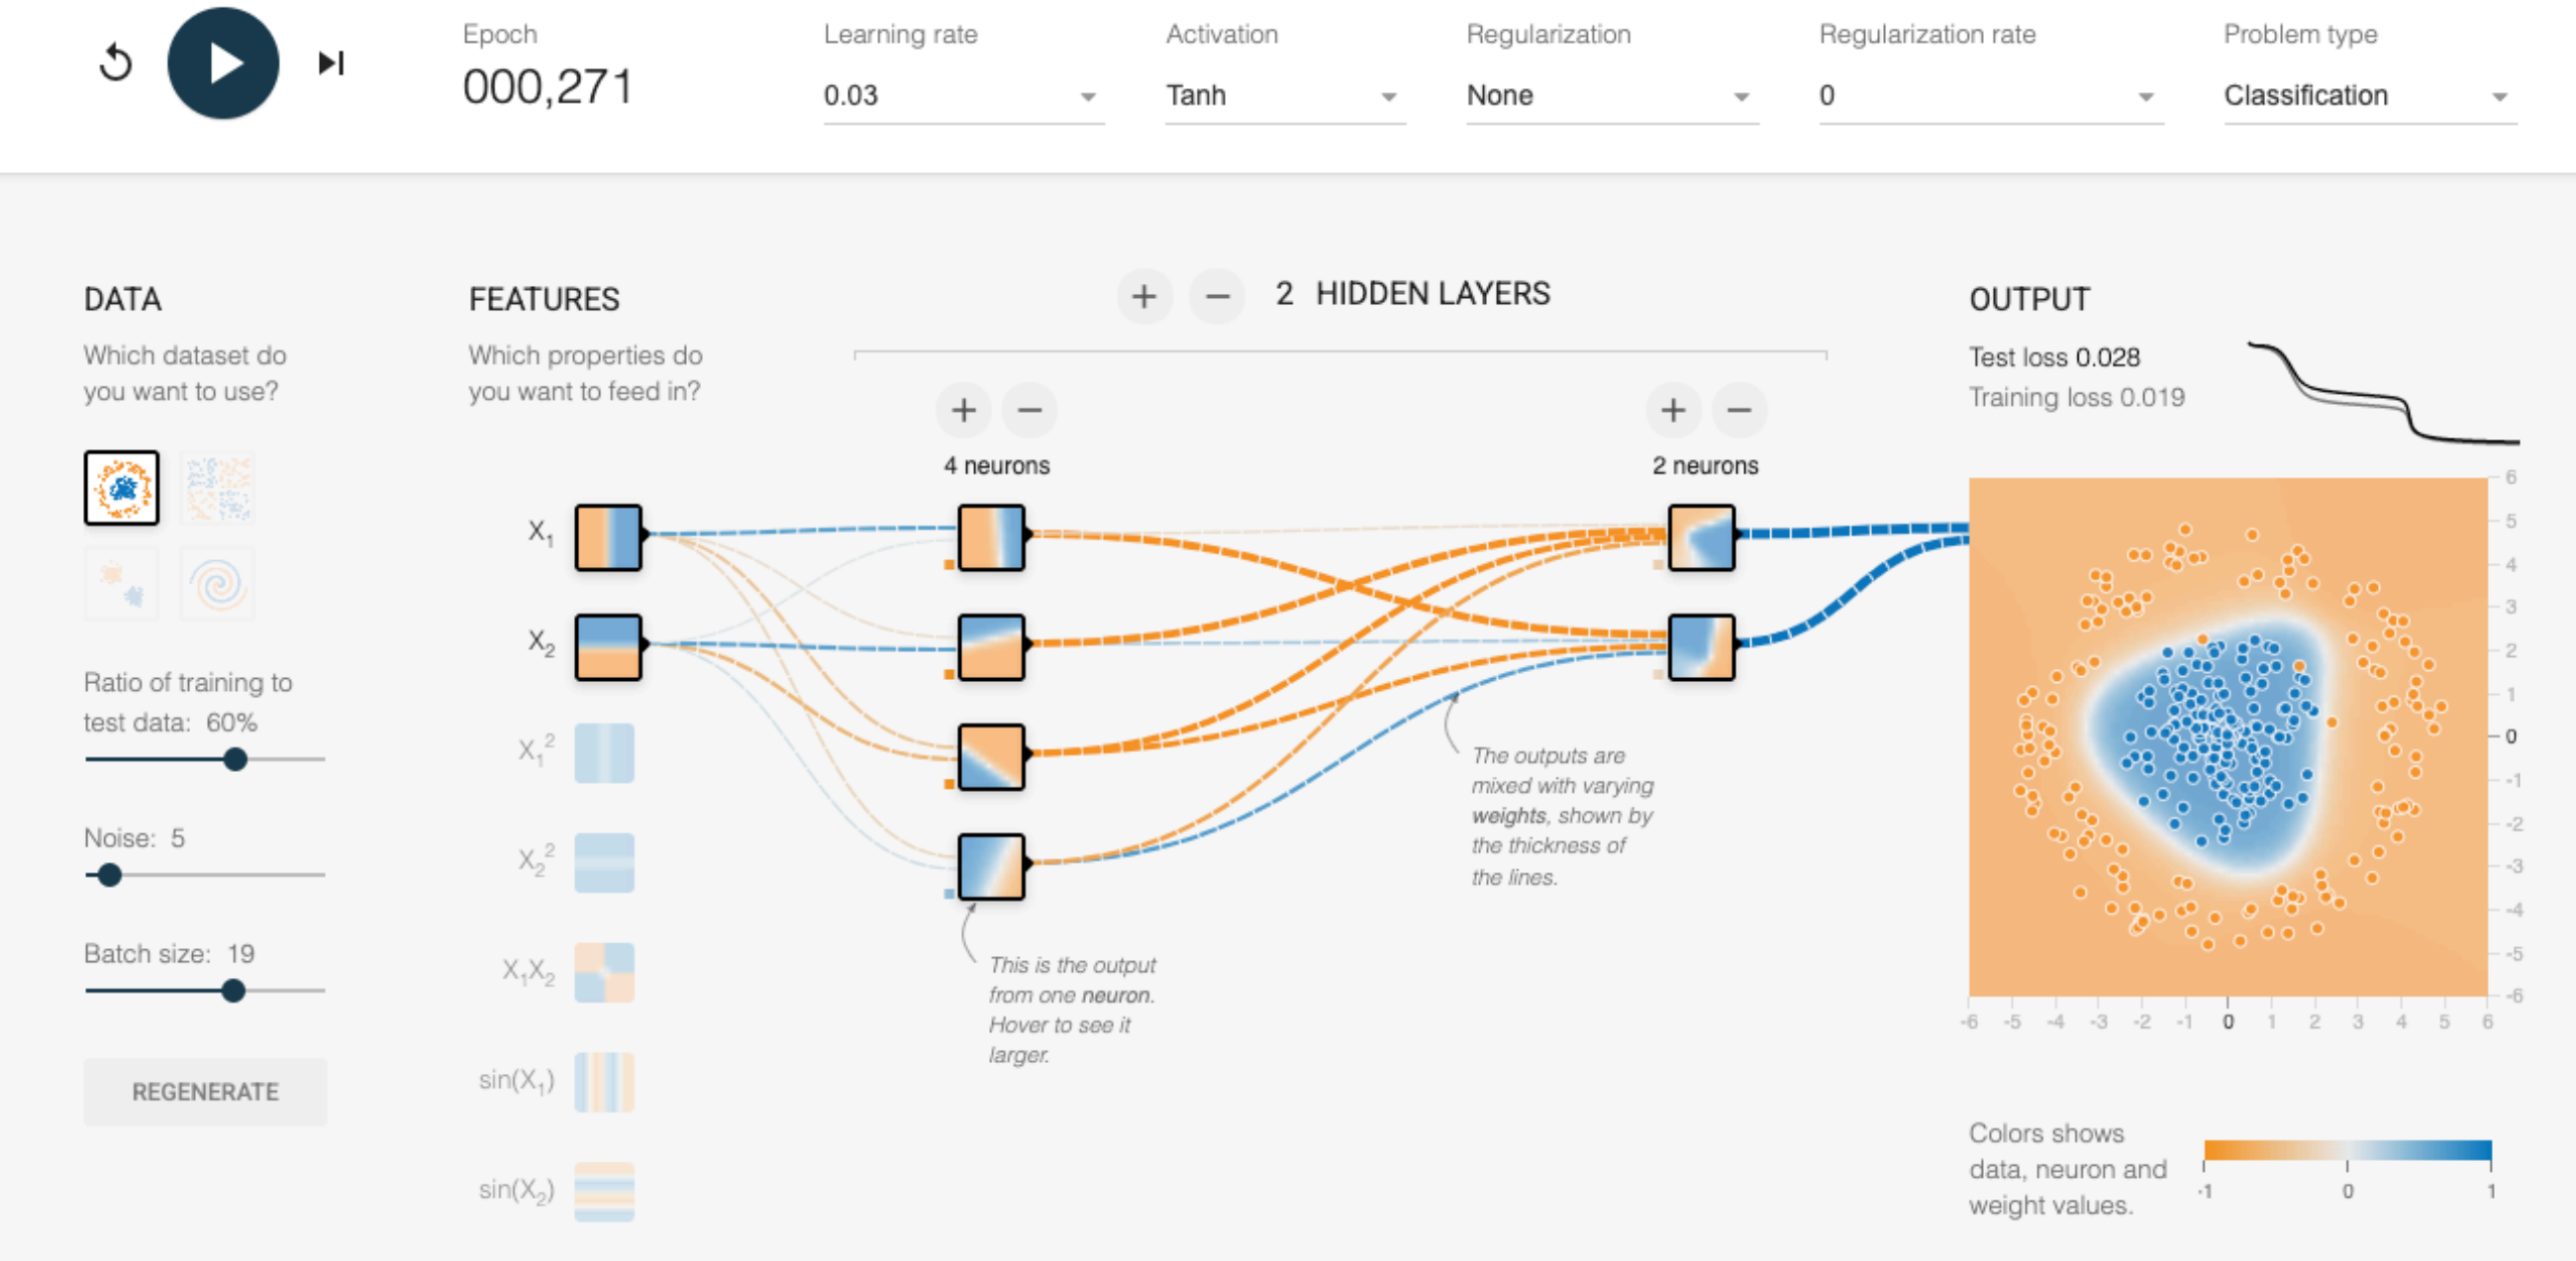
Colors shows data, neuron and weight values.



☐ Show test data ☐ Discretize output



***ReLU faster,
notice linear
boundaries***

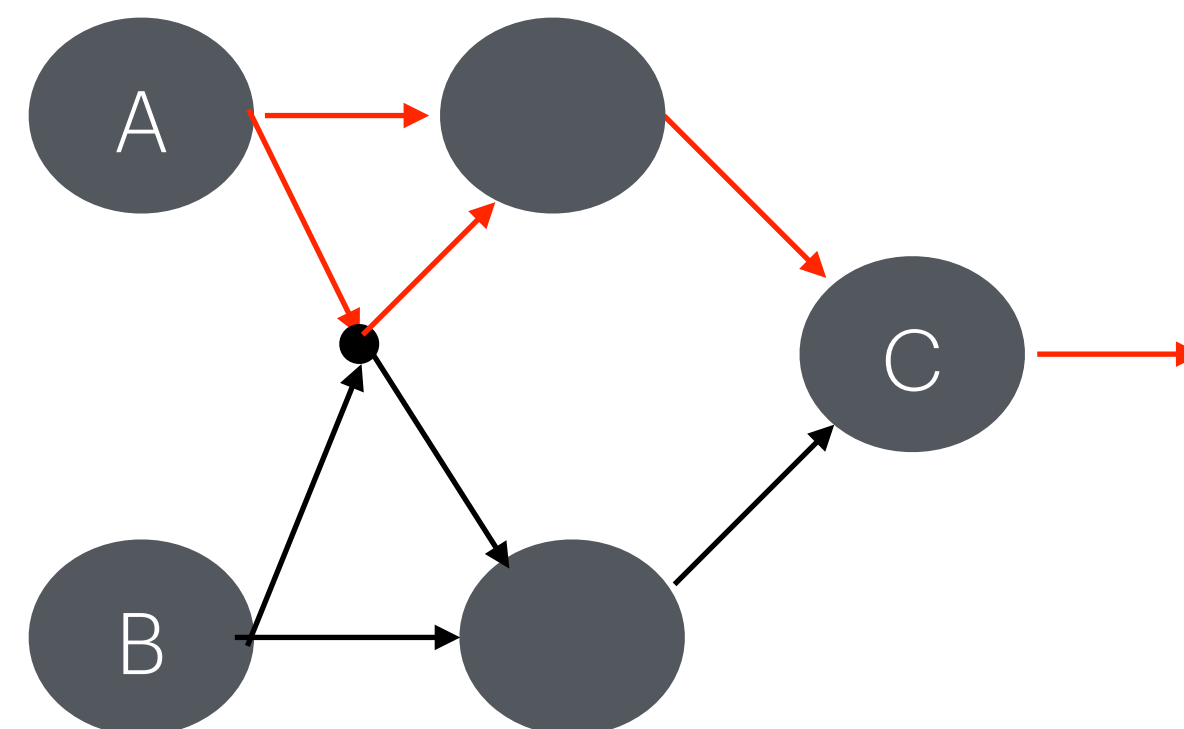
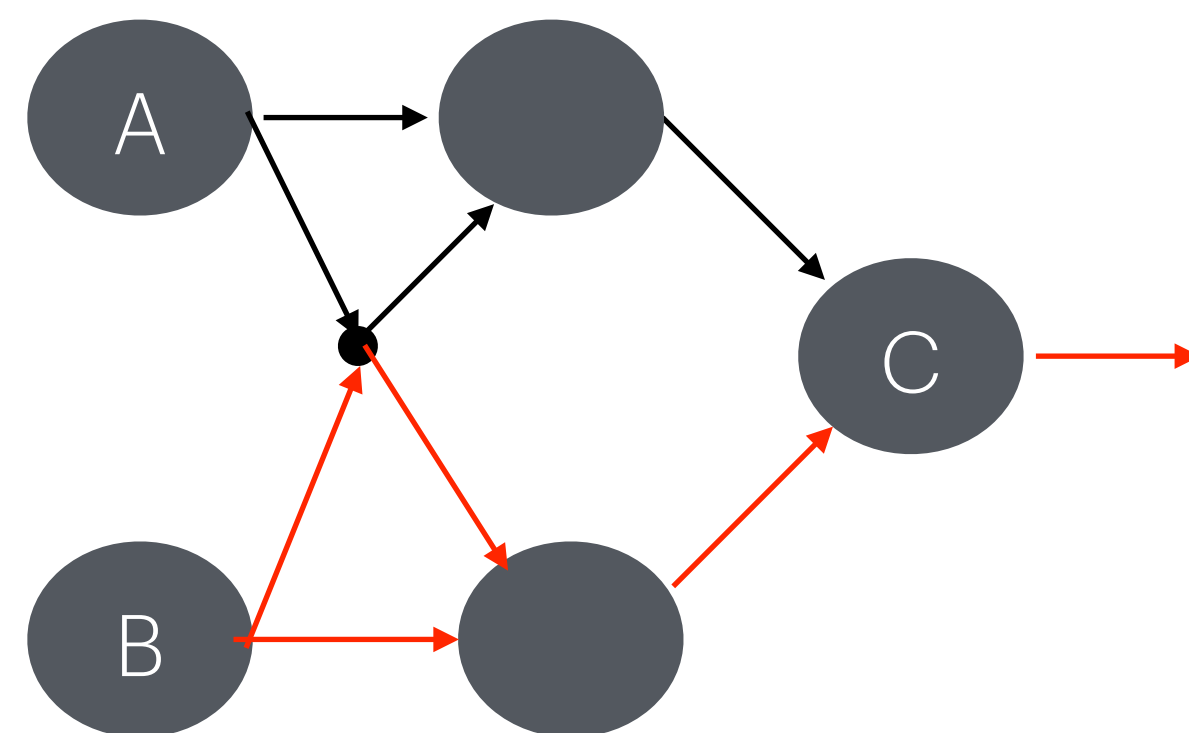
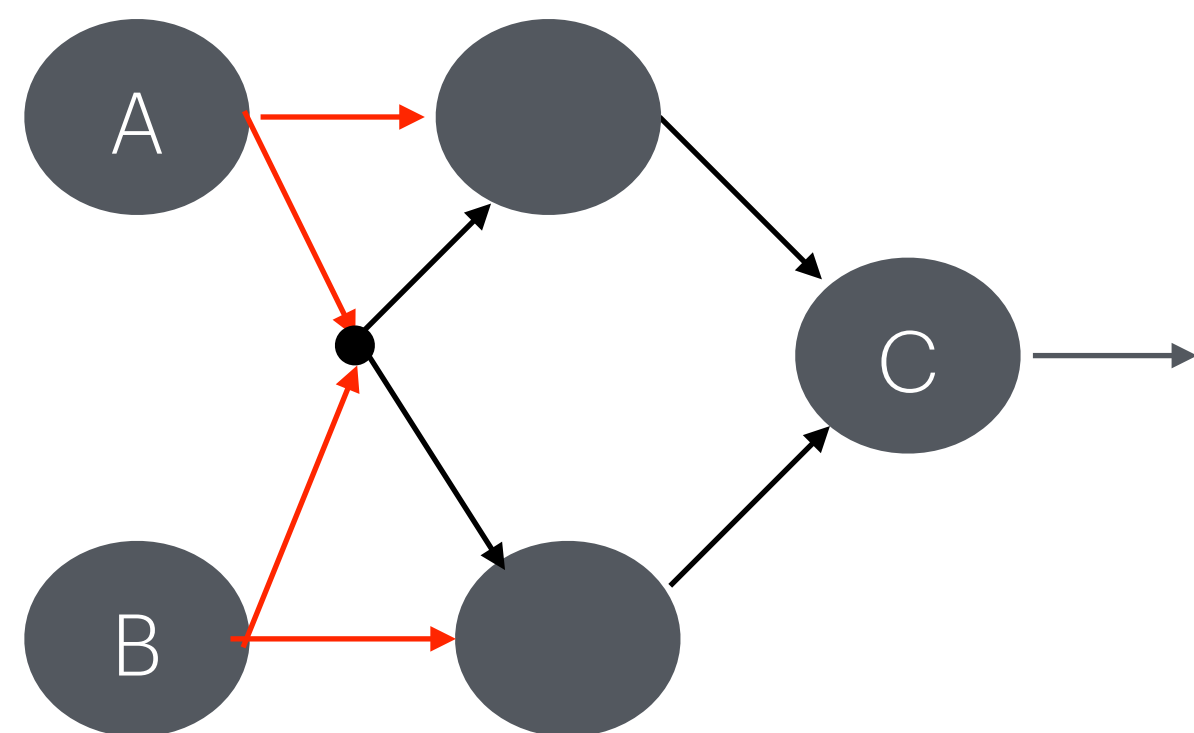


***TANU takes time
to converge on a
solution***

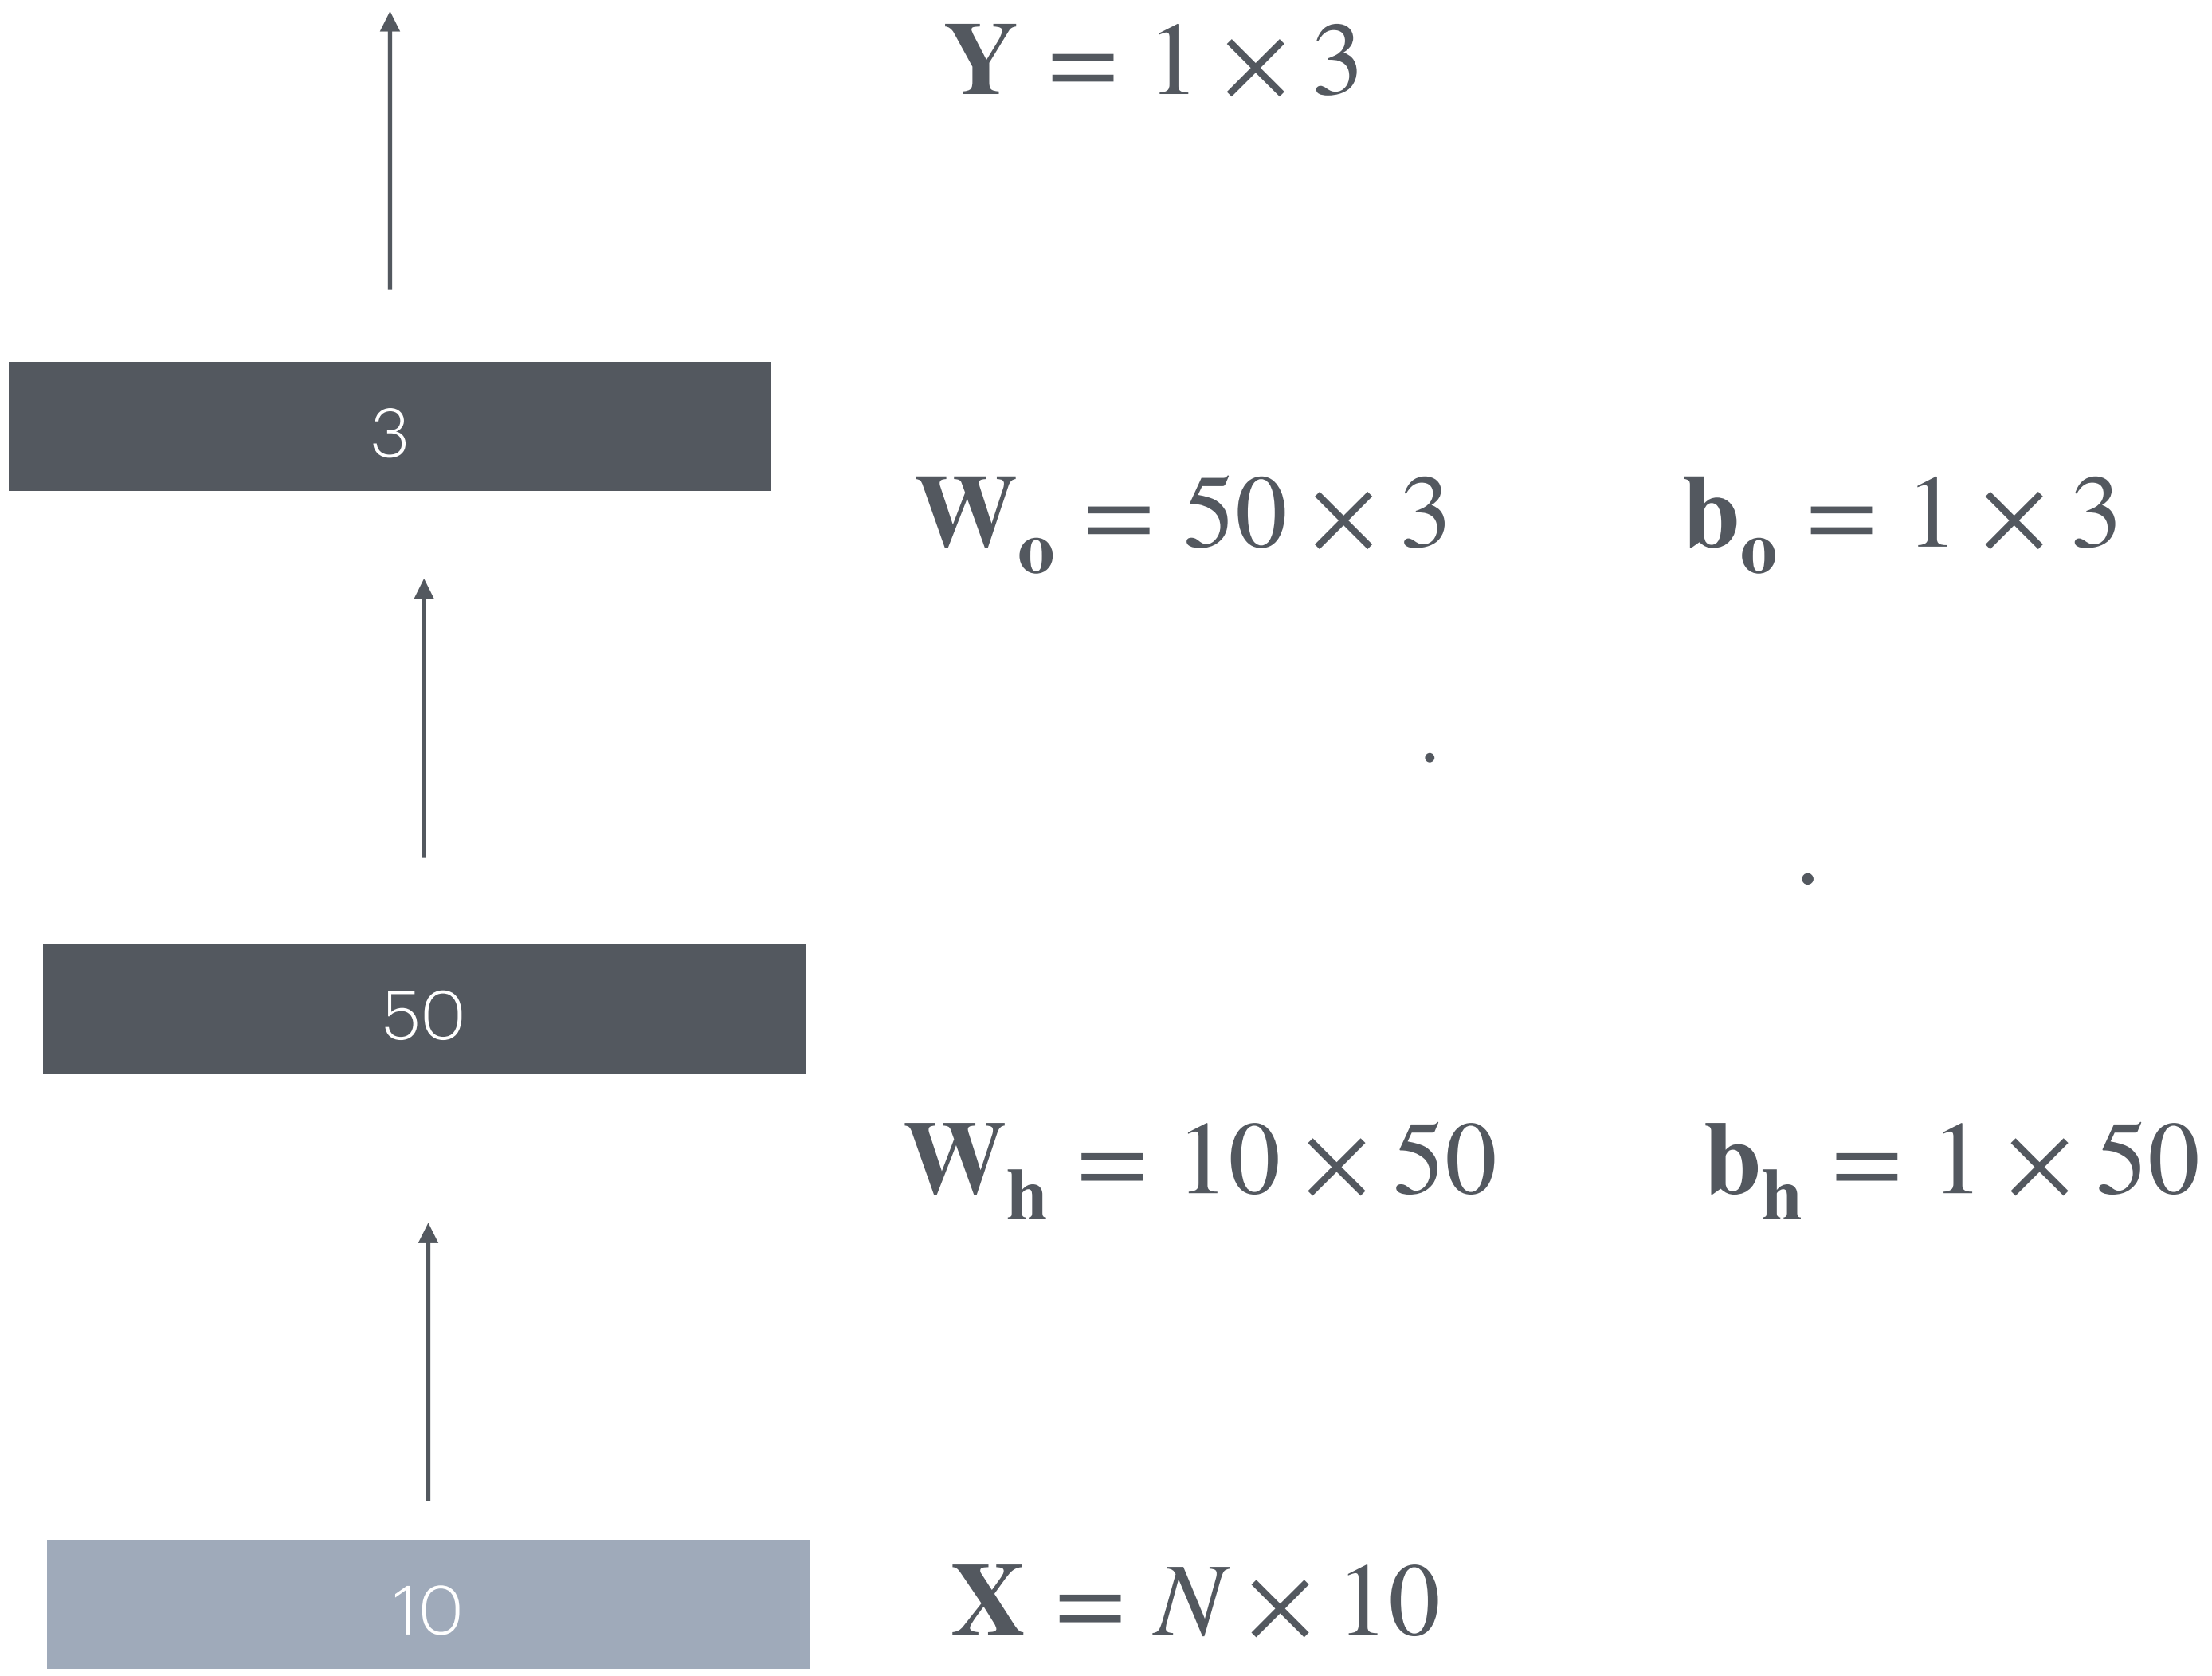
Exercise 2

Draw an ANN using the original artificial neurons that computes $A \oplus B$

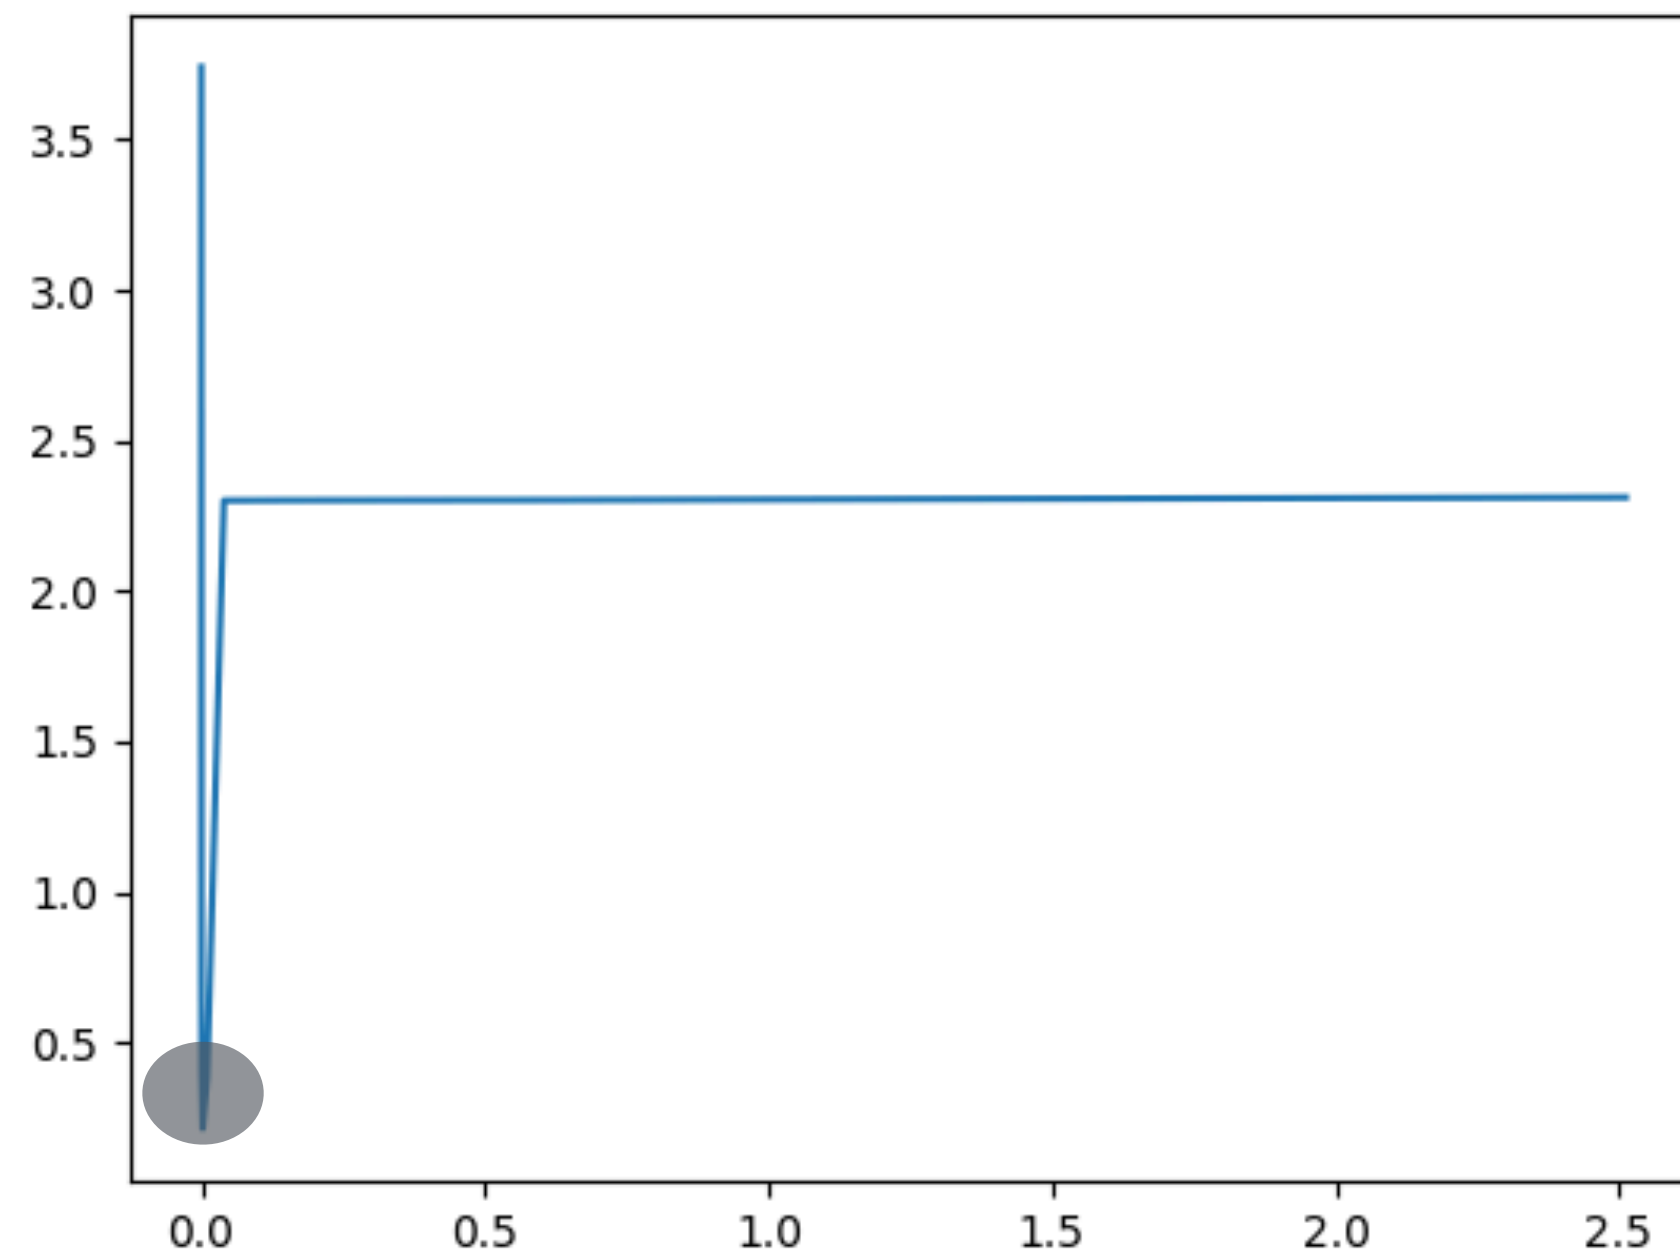
$$(A \cup \neg B) \cup (\neg A \cup B)$$



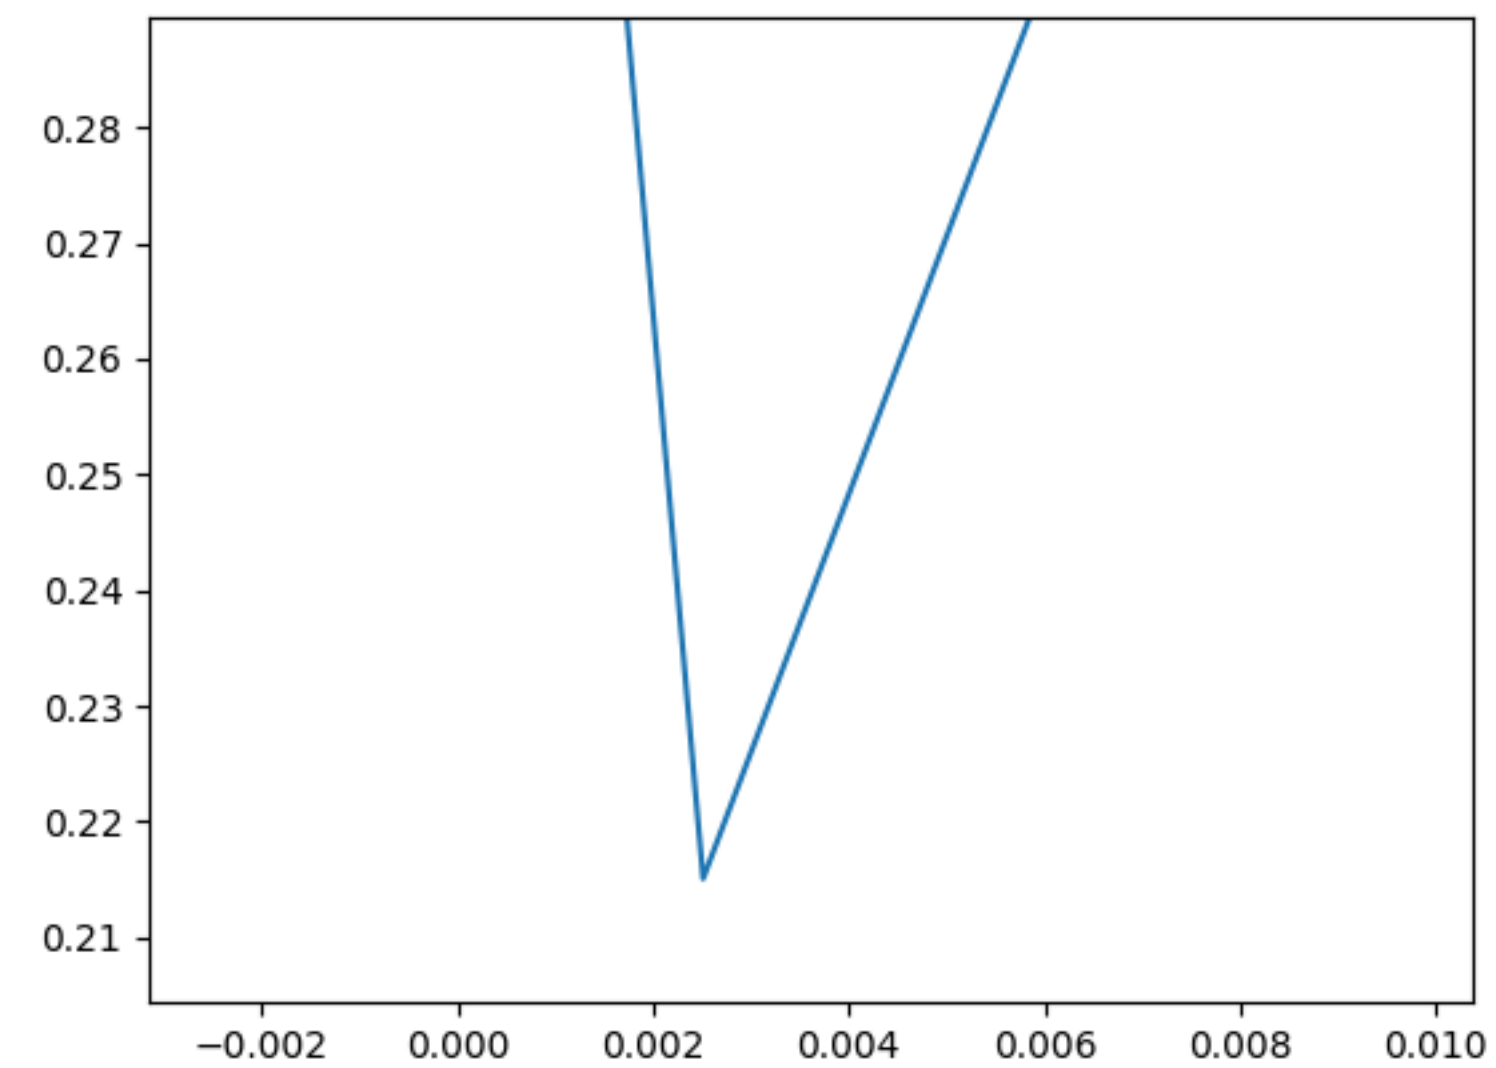
Exercise 3



Find Learning Rate



Gradually increase learning rate



Region where loss decreases and immediately increases is an optimal learning rate